

# Linux Kernel Network Security - Transport Layer Security (TLS)

Deep Hacking

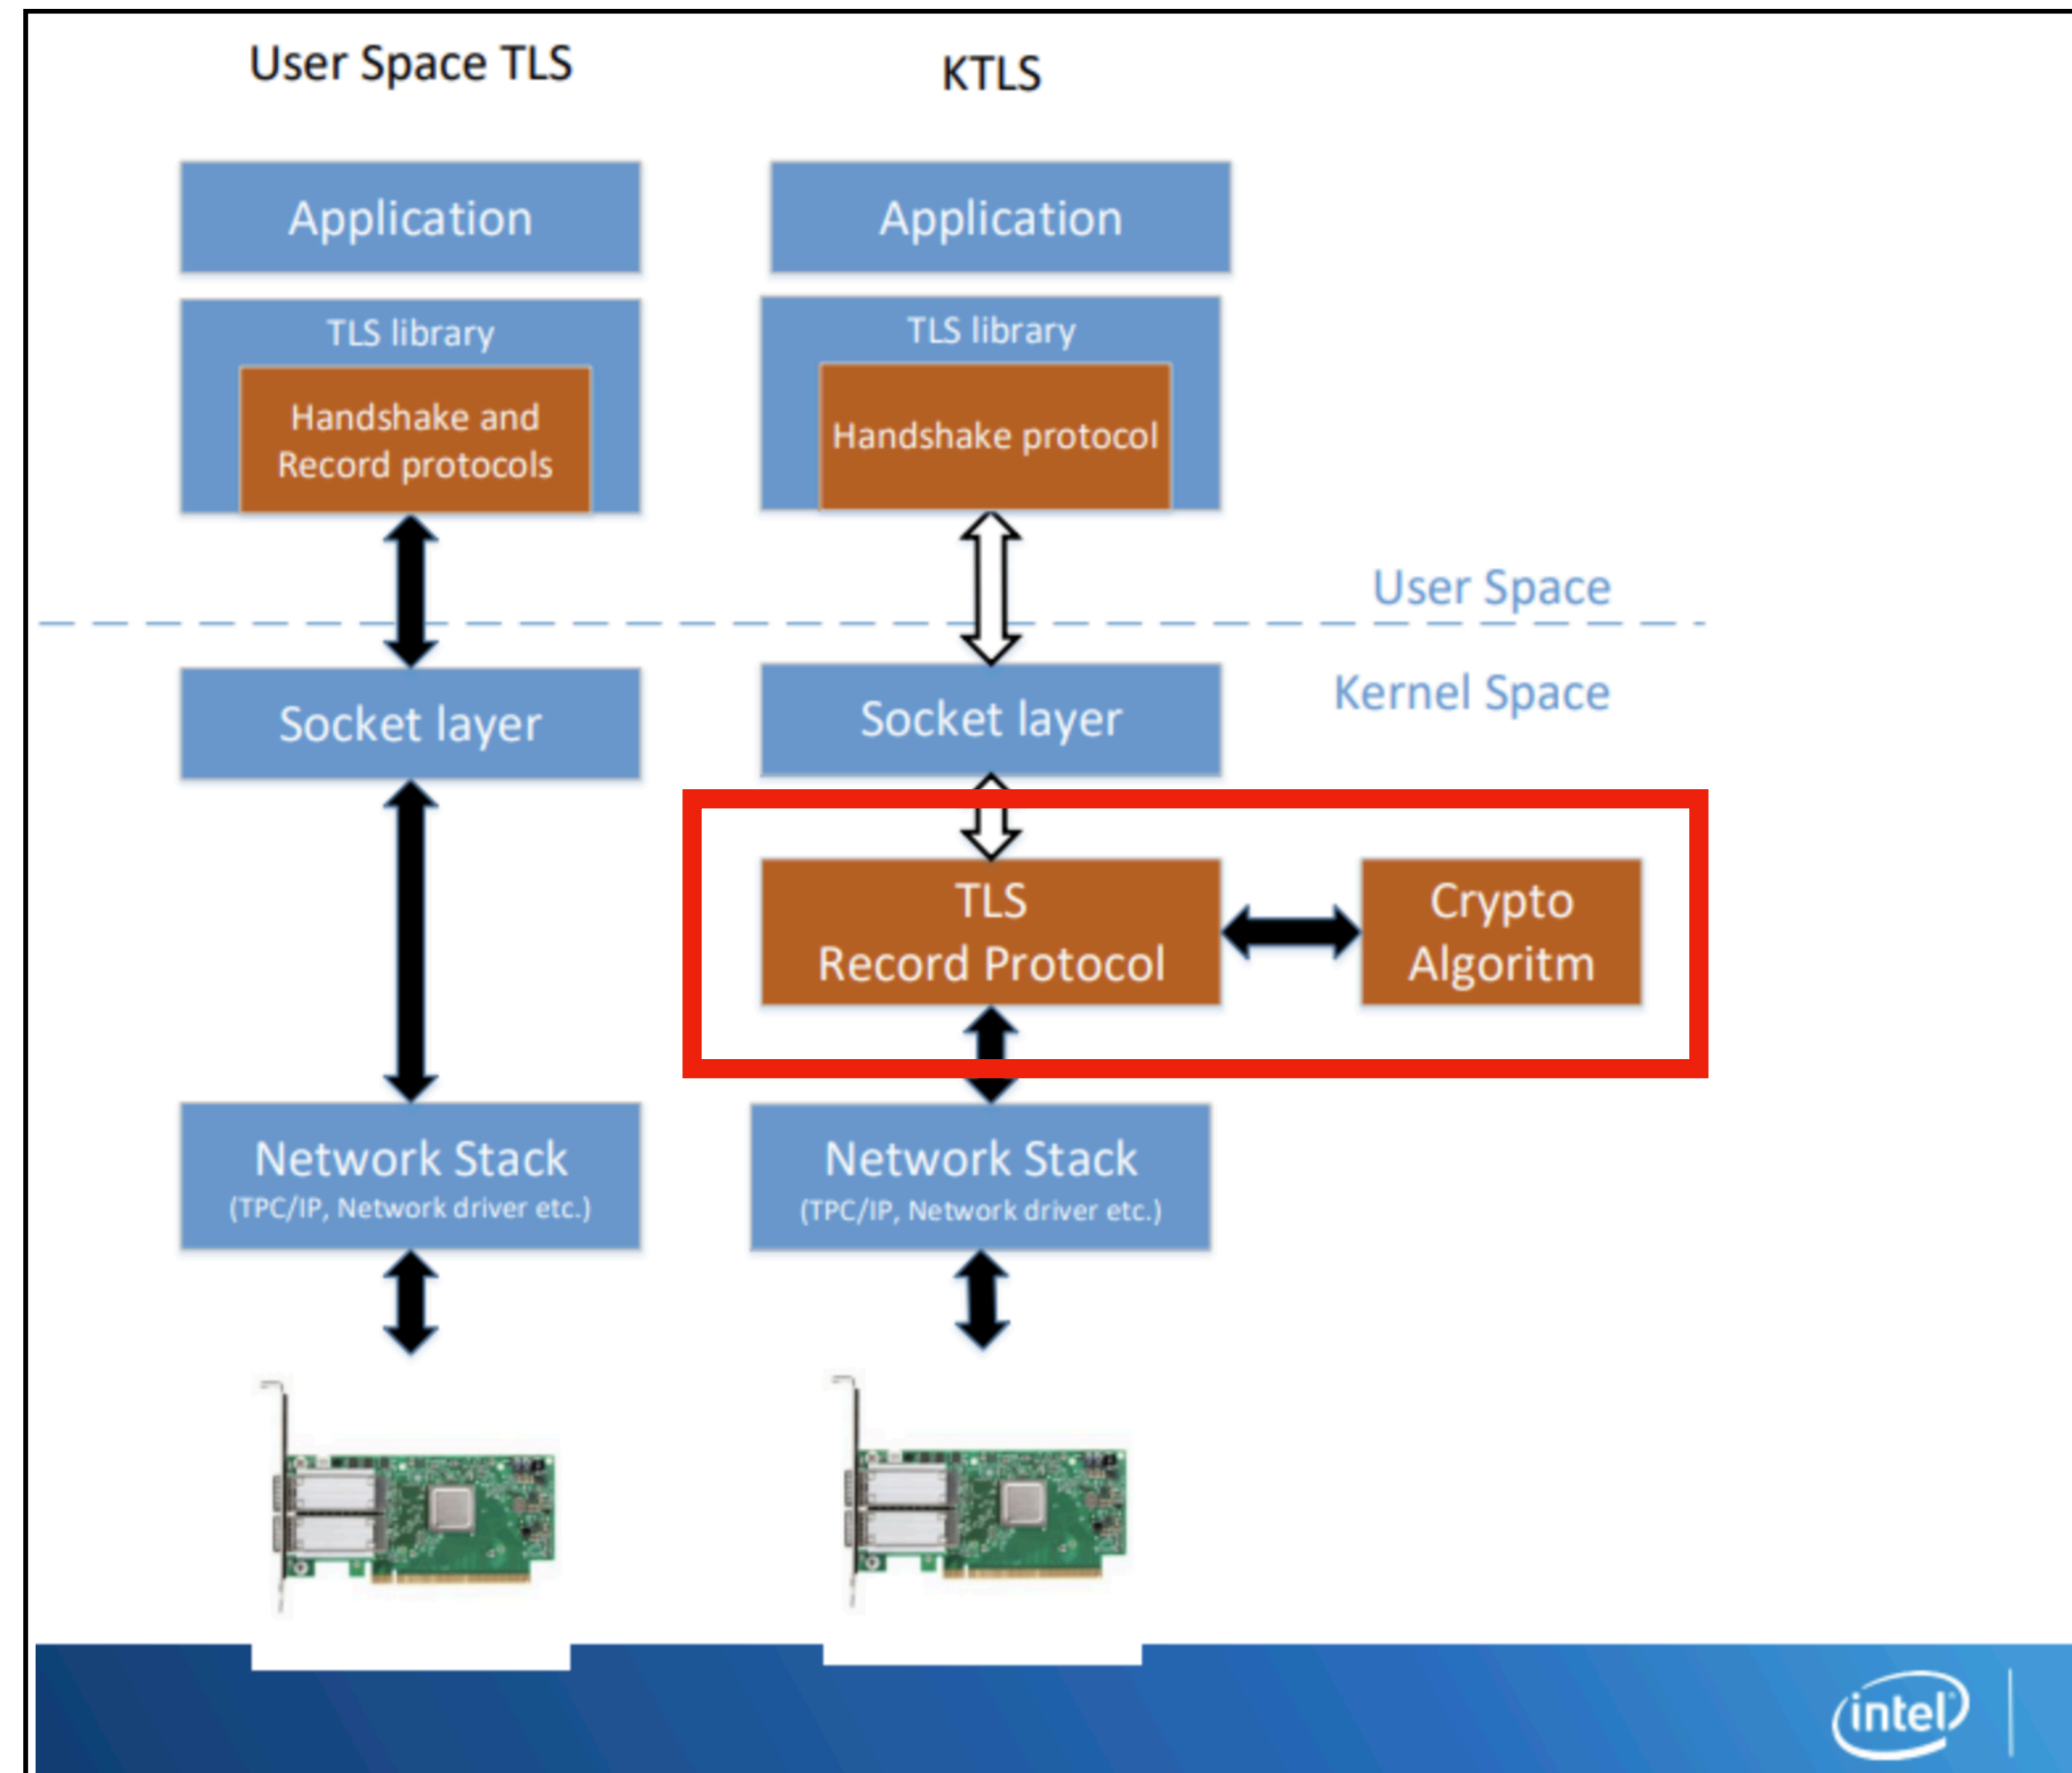
2025/01/19 Pumpkin 🎃

# Outline

- Overview
- Vulnerability

# Overview

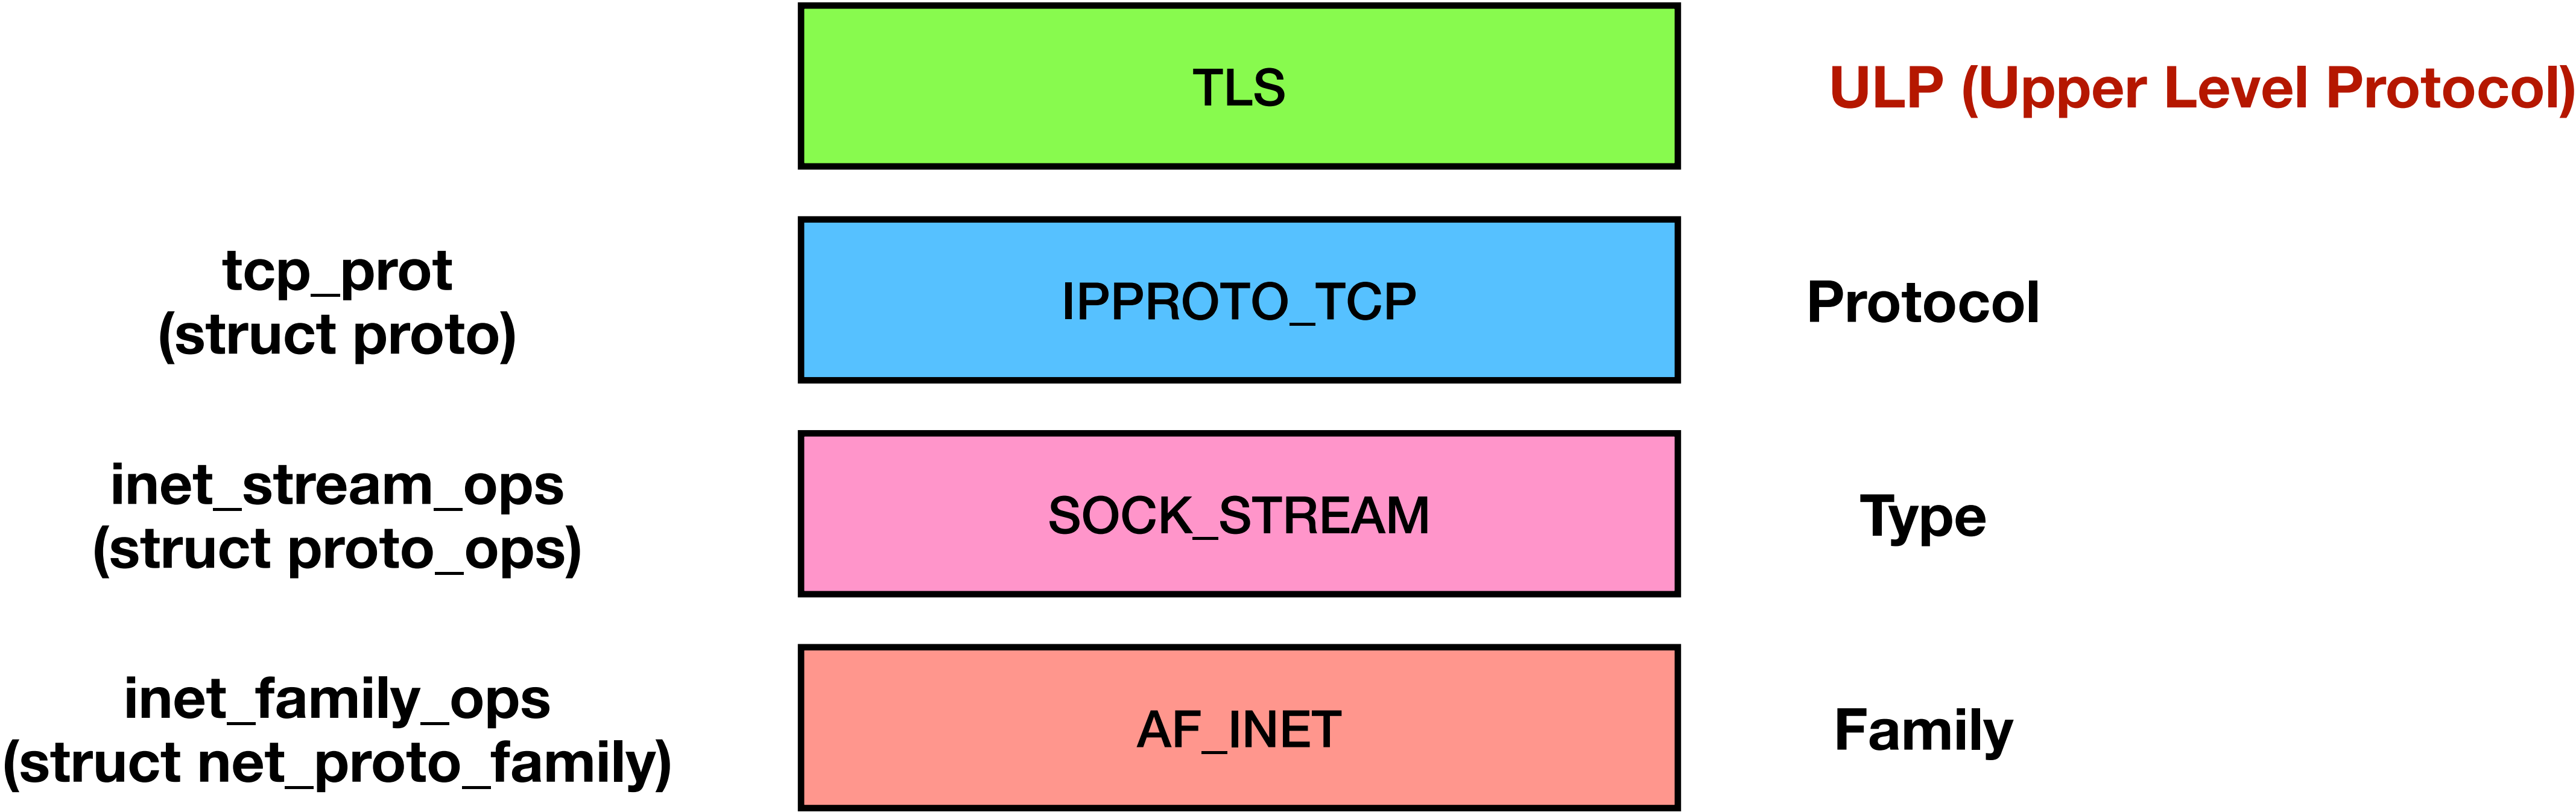
# Overview



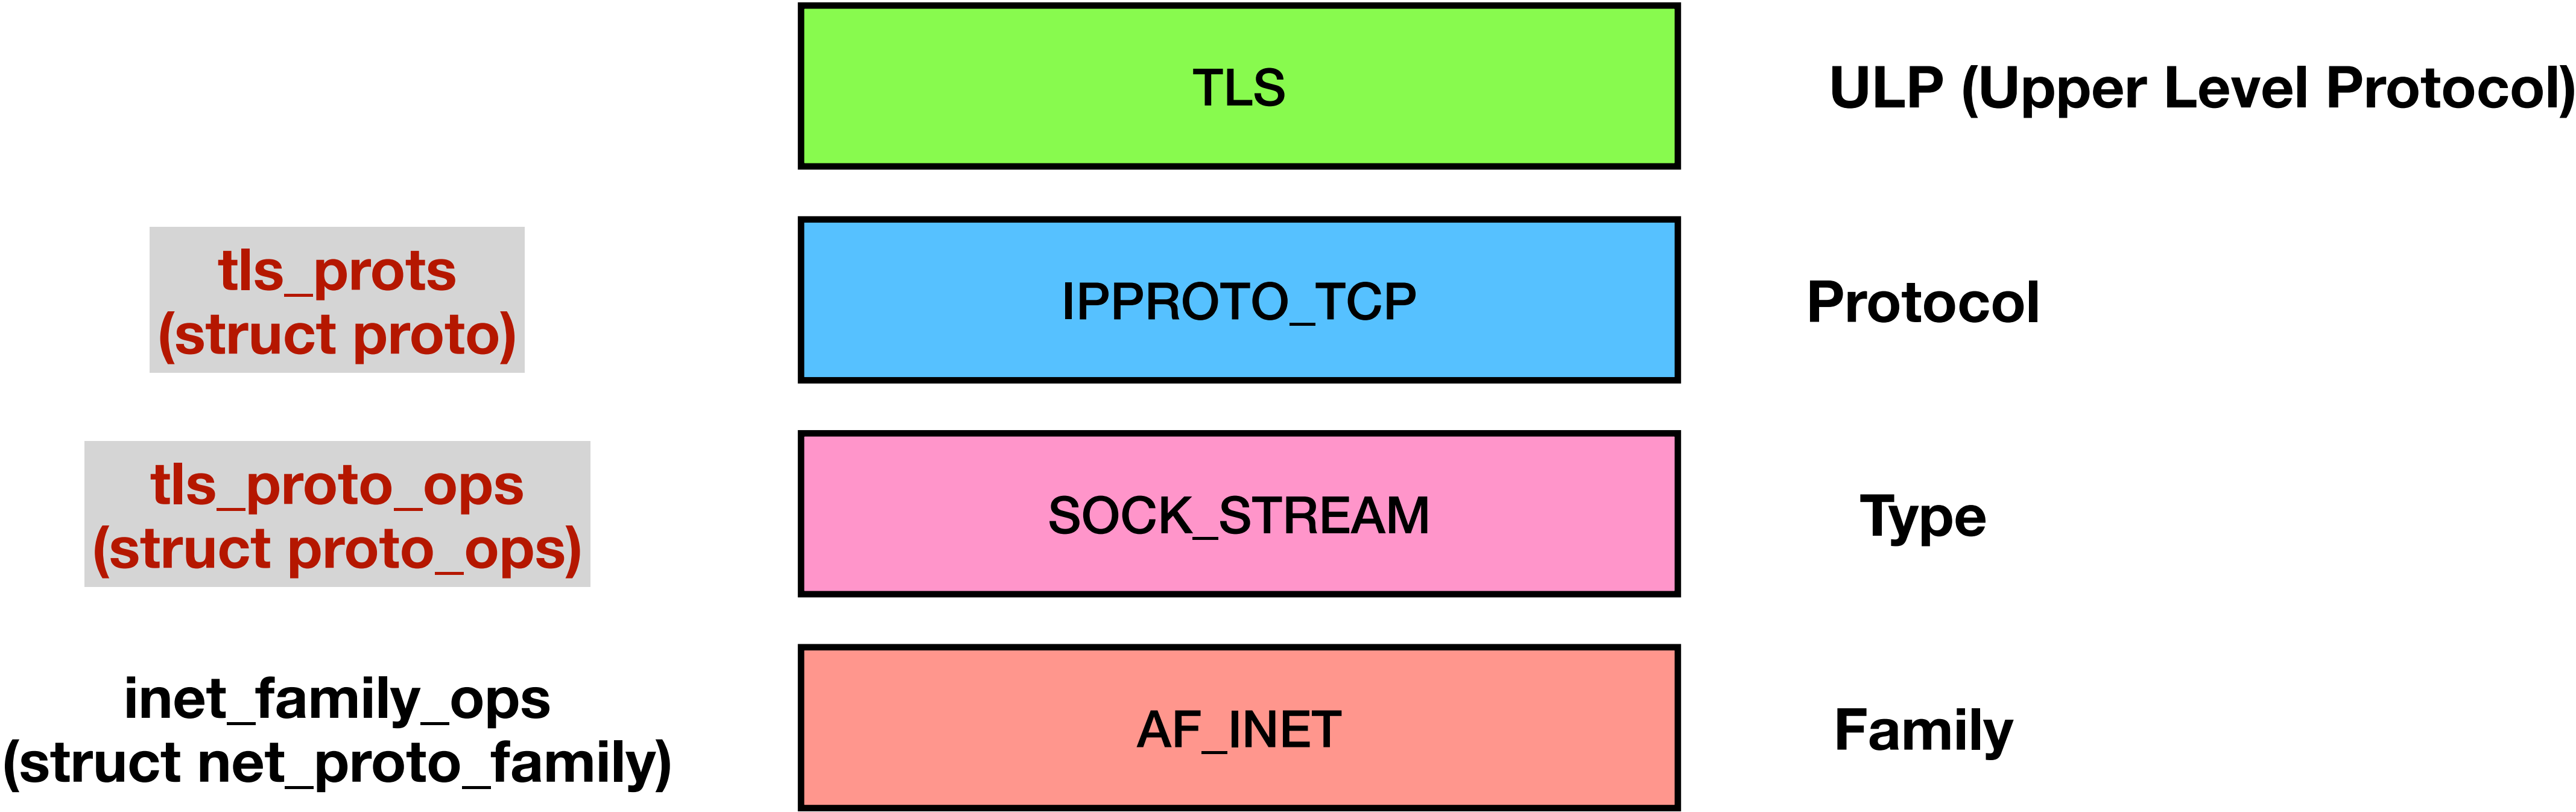
# Overview

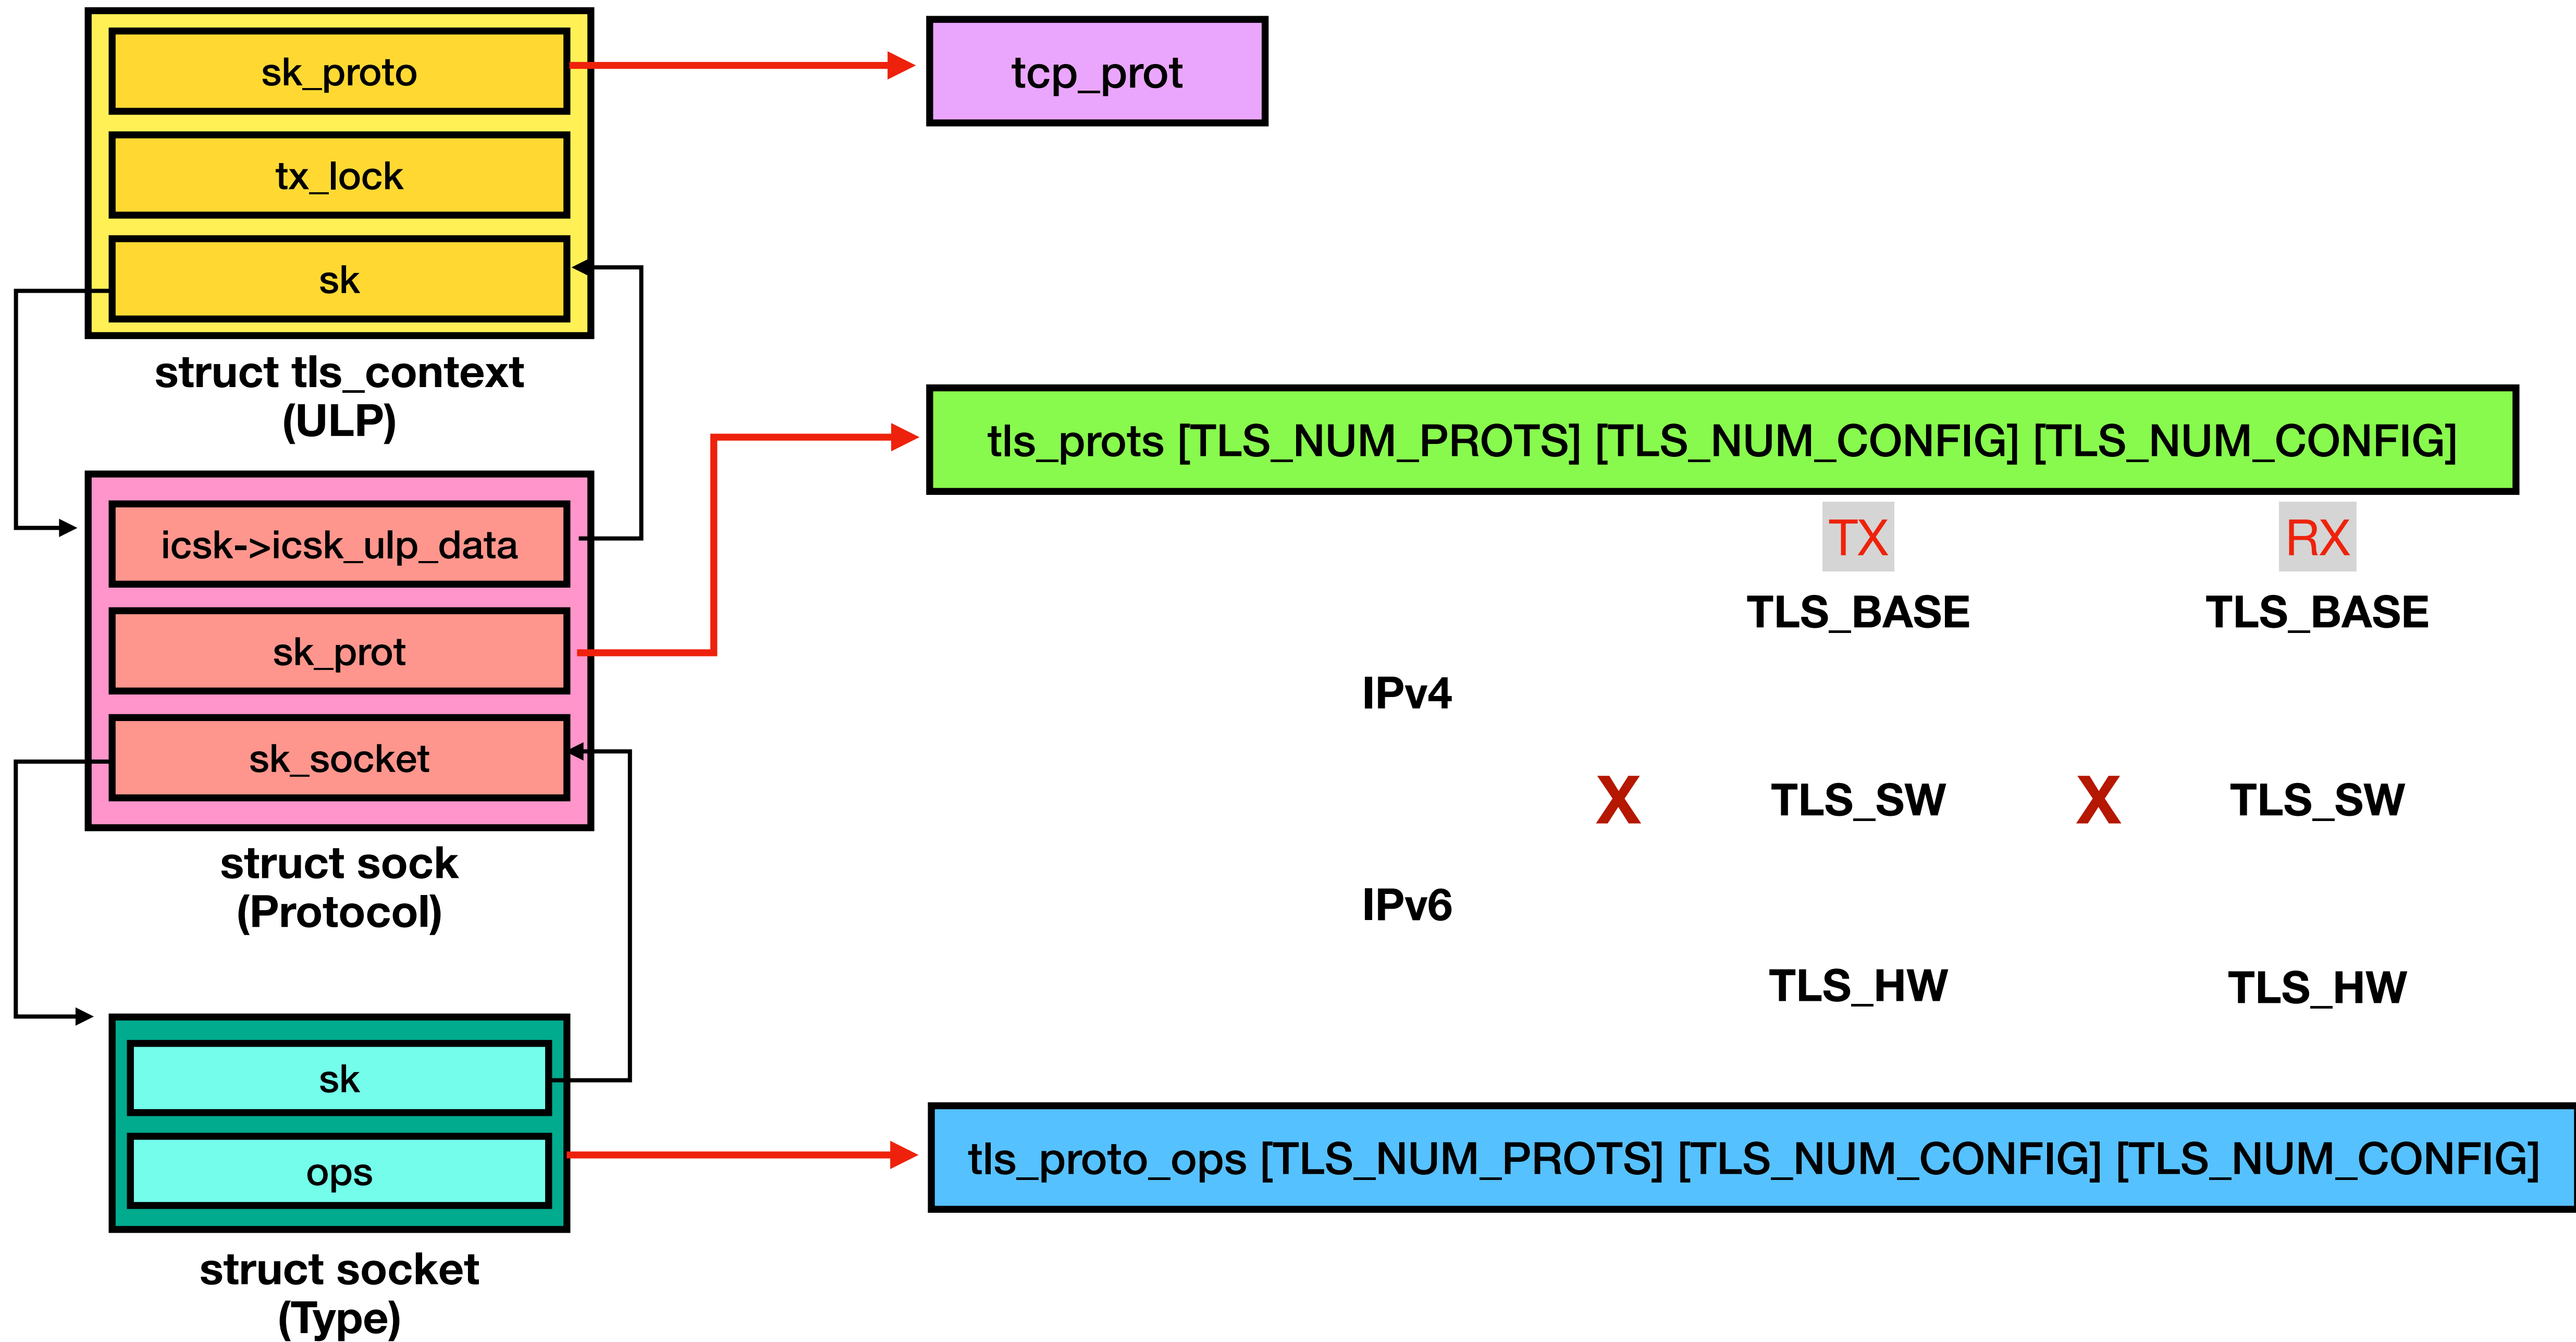
<b>tcp_prot</b> (struct proto)	IPPROTO_TCP	<b>Protocol</b>
<b>inet_stream_ops</b> (struct proto_ops)	SOCK_STREAM	<b>Type</b>
<b>inet_family_ops</b> (struct net_proto_family)	AF_INET	<b>Family</b>

# Overview



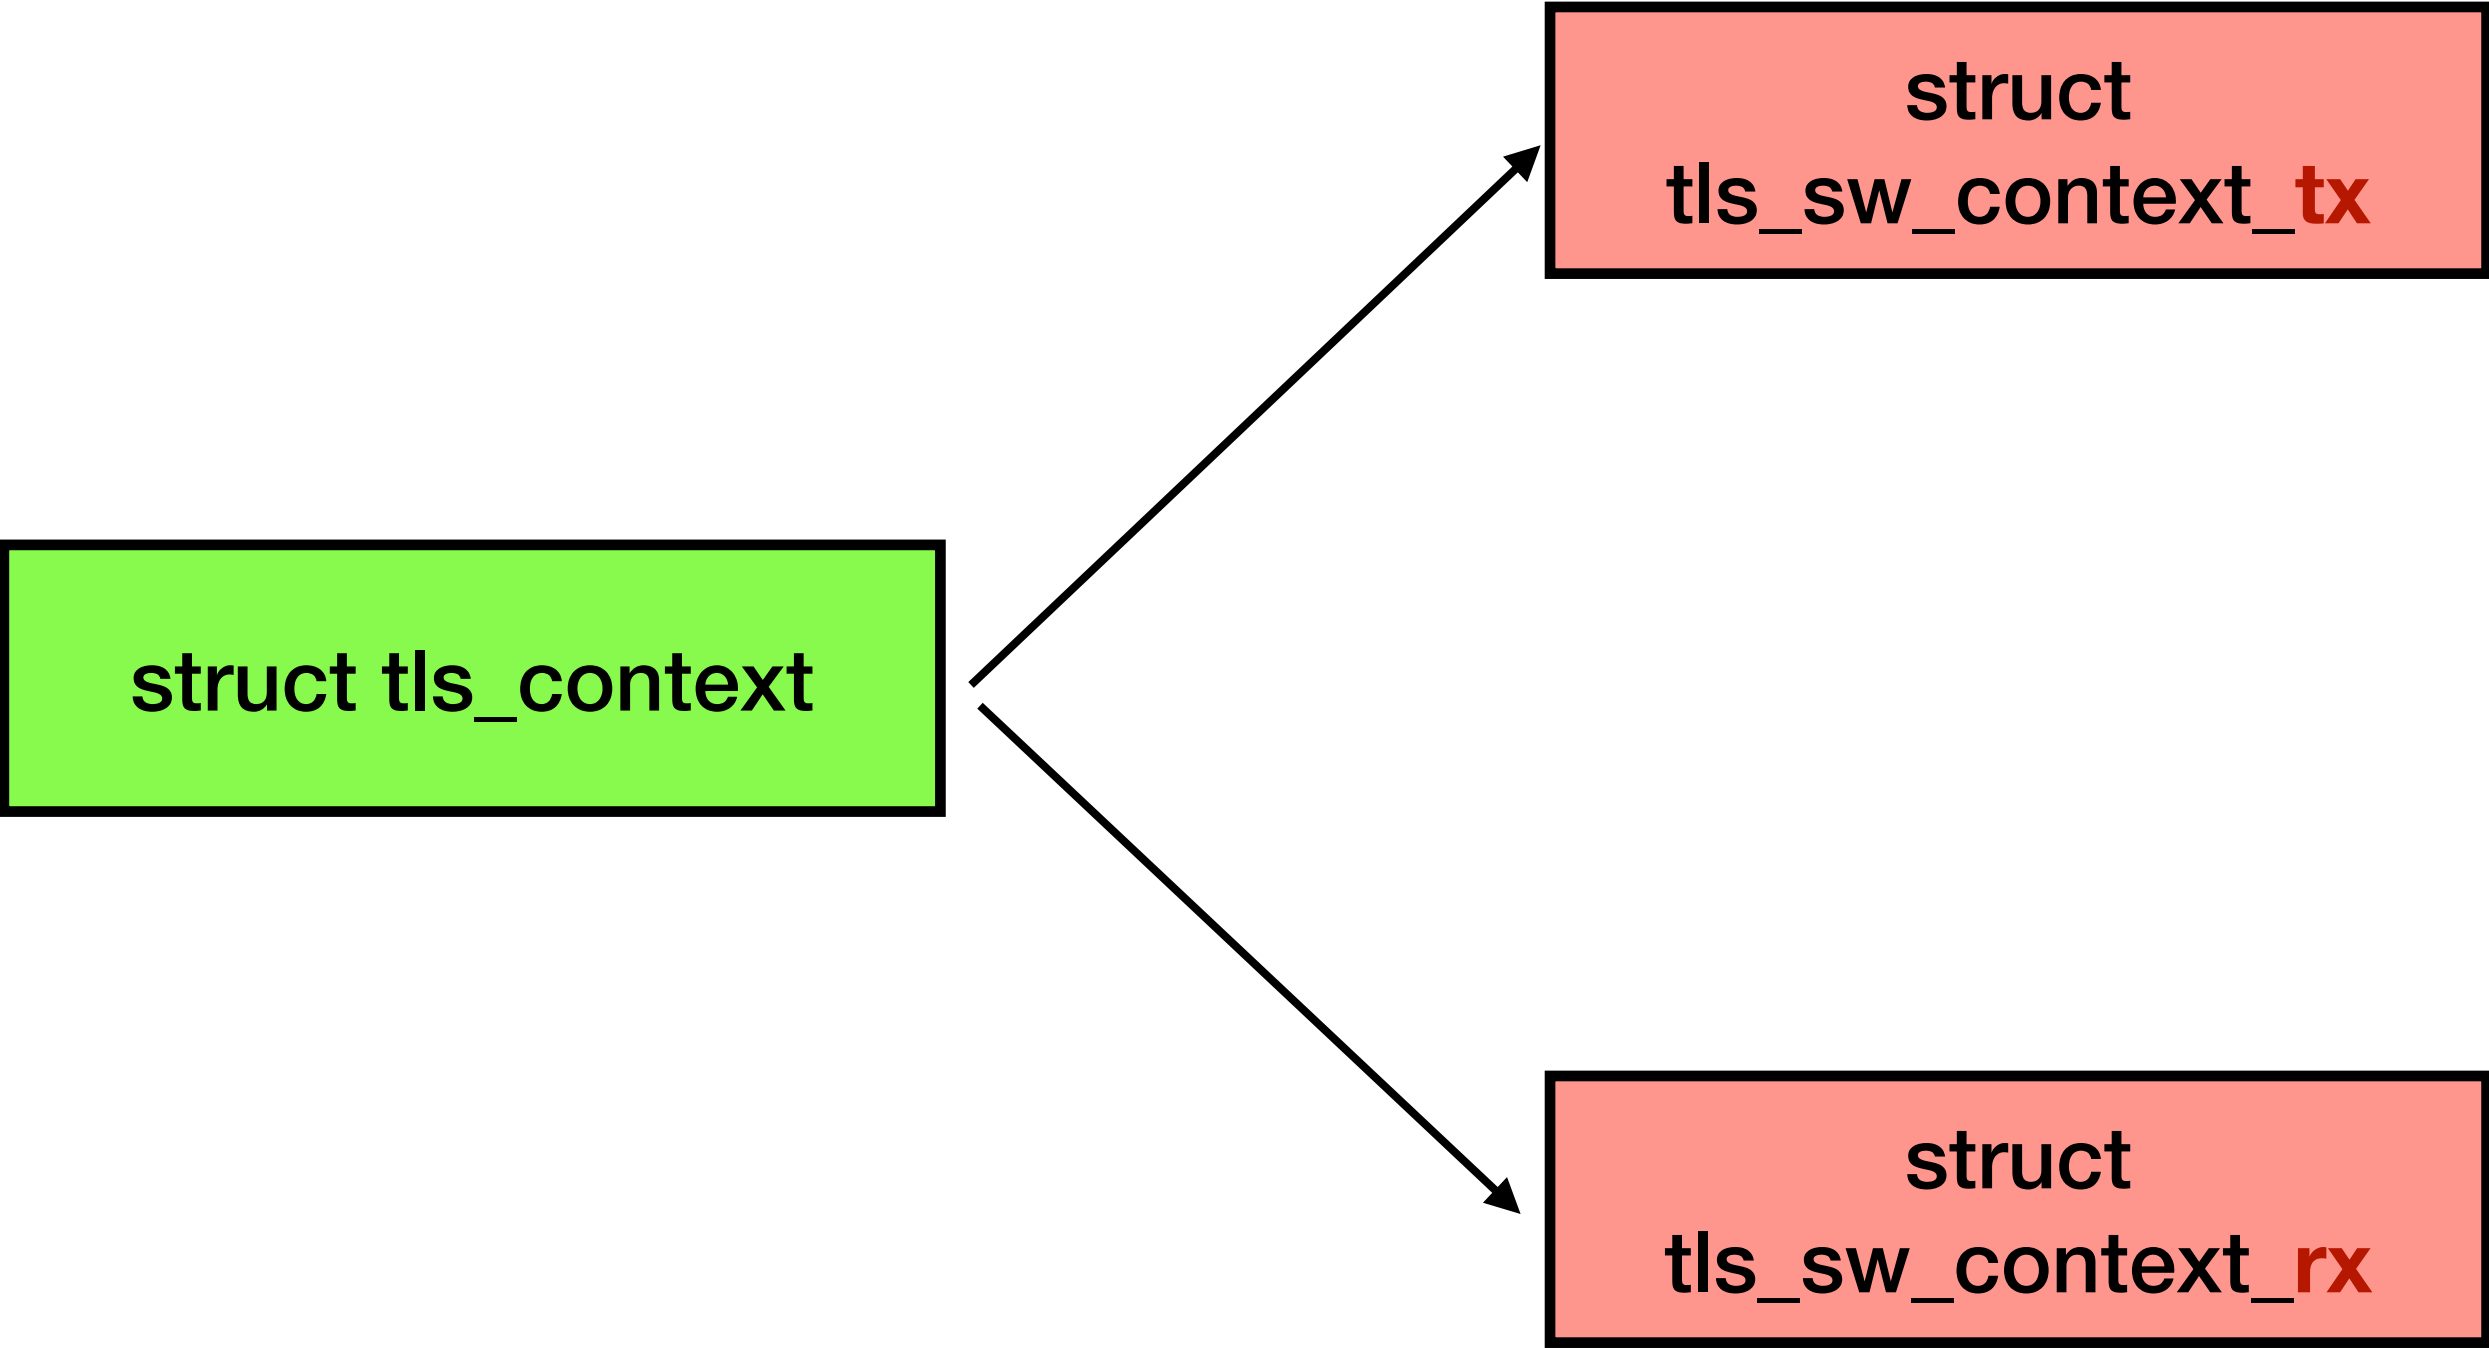
# Overview







# Overview



cipher type	AES_GCM_128
TLS vers.	TLS_1_2
KEY	0123...DEF
IV	12345678
SALT	SALT
...	...

cipher type	...
TLS vers.	...
KEY	...
IV	...
SALT	...
...	...

# Overview

```
// 1. config a TCP socket to TLS
setsockopt(sockfd, SOL_TCP, TCP_ULP, "tls", sizeof("tls"));

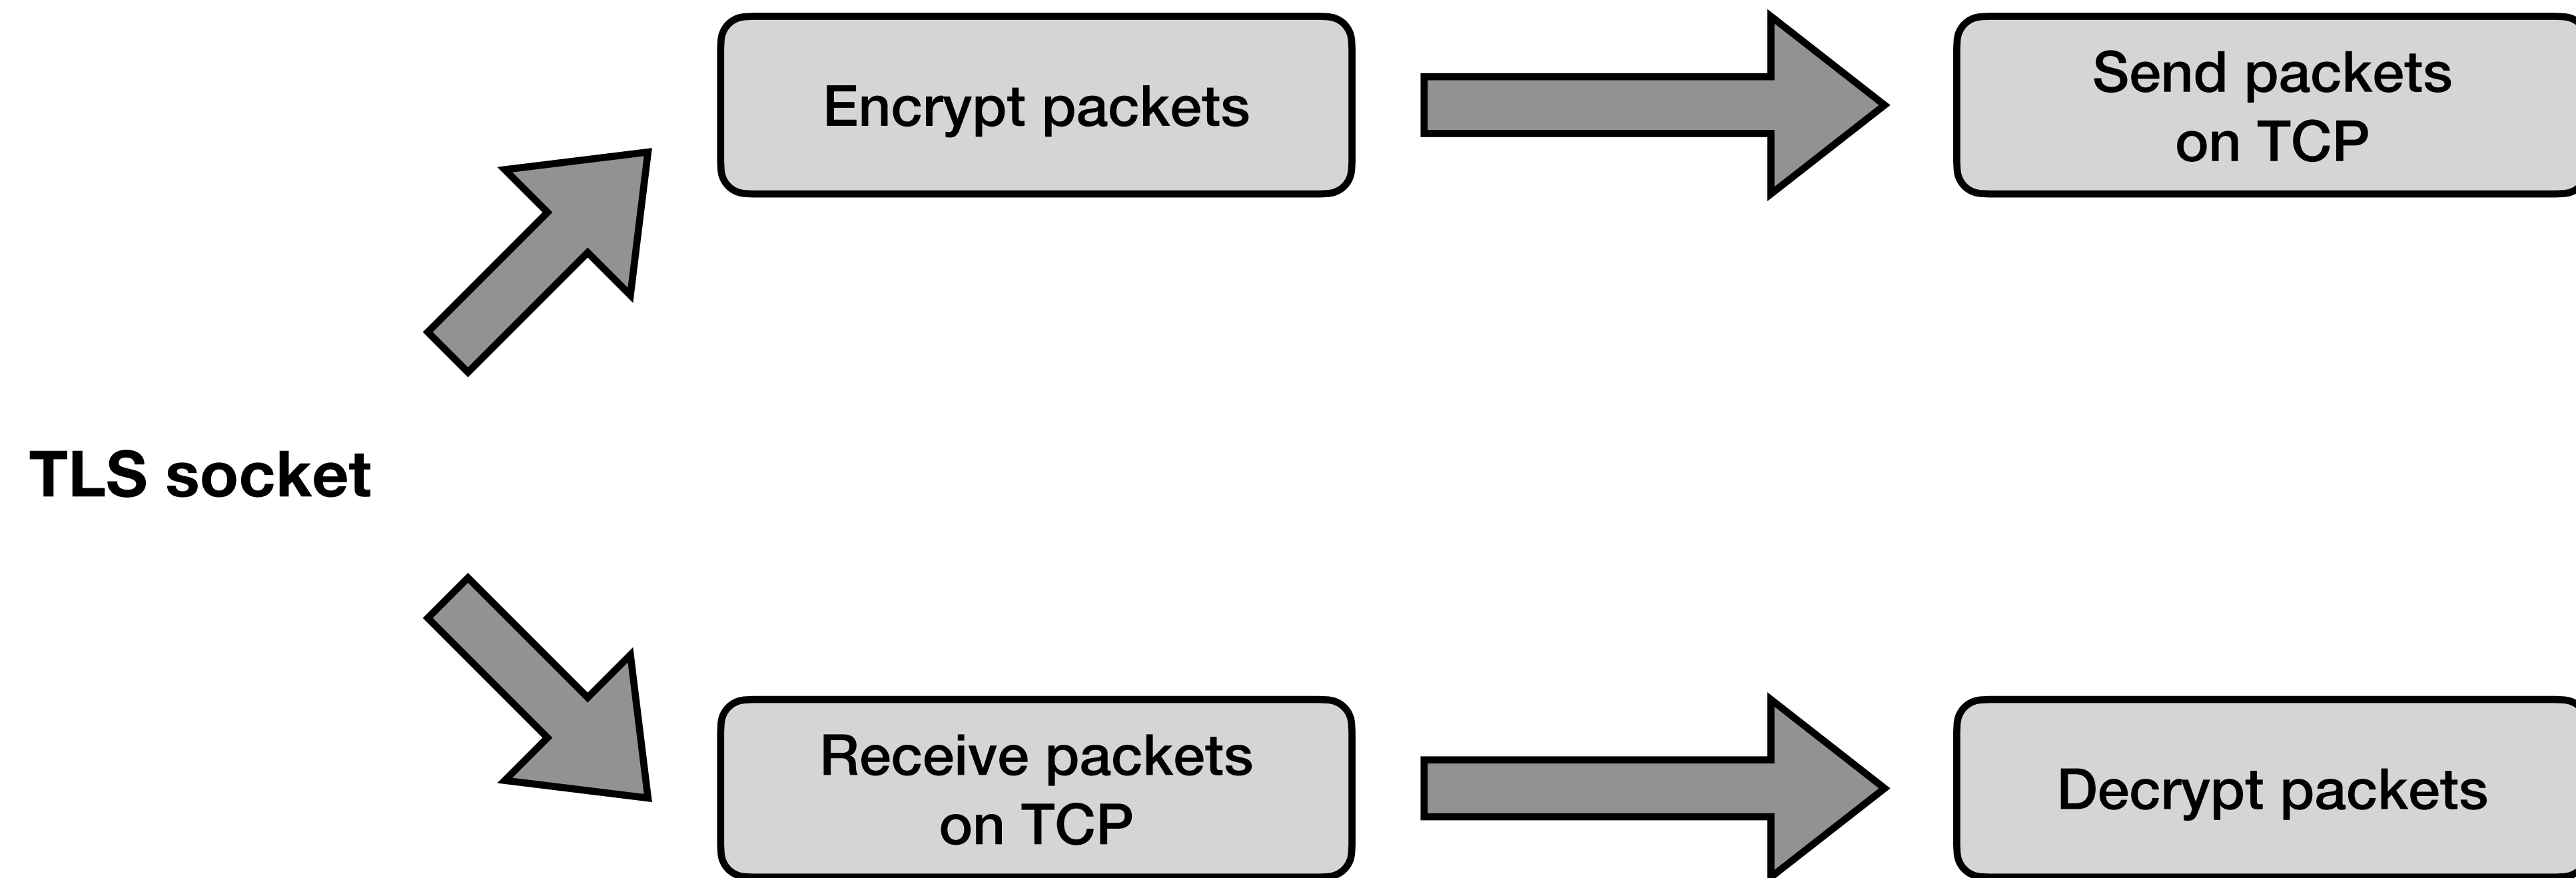
// 2. config TX / RX of the TLS socket
struct tls12_crypto_info_aes_gcm_128 crypto_info = {};

crypto_info.info.version = TLS_1_2_VERSION;
crypto_info.info.cipher_type = TLS_CIPHER_AES_GCM_128;

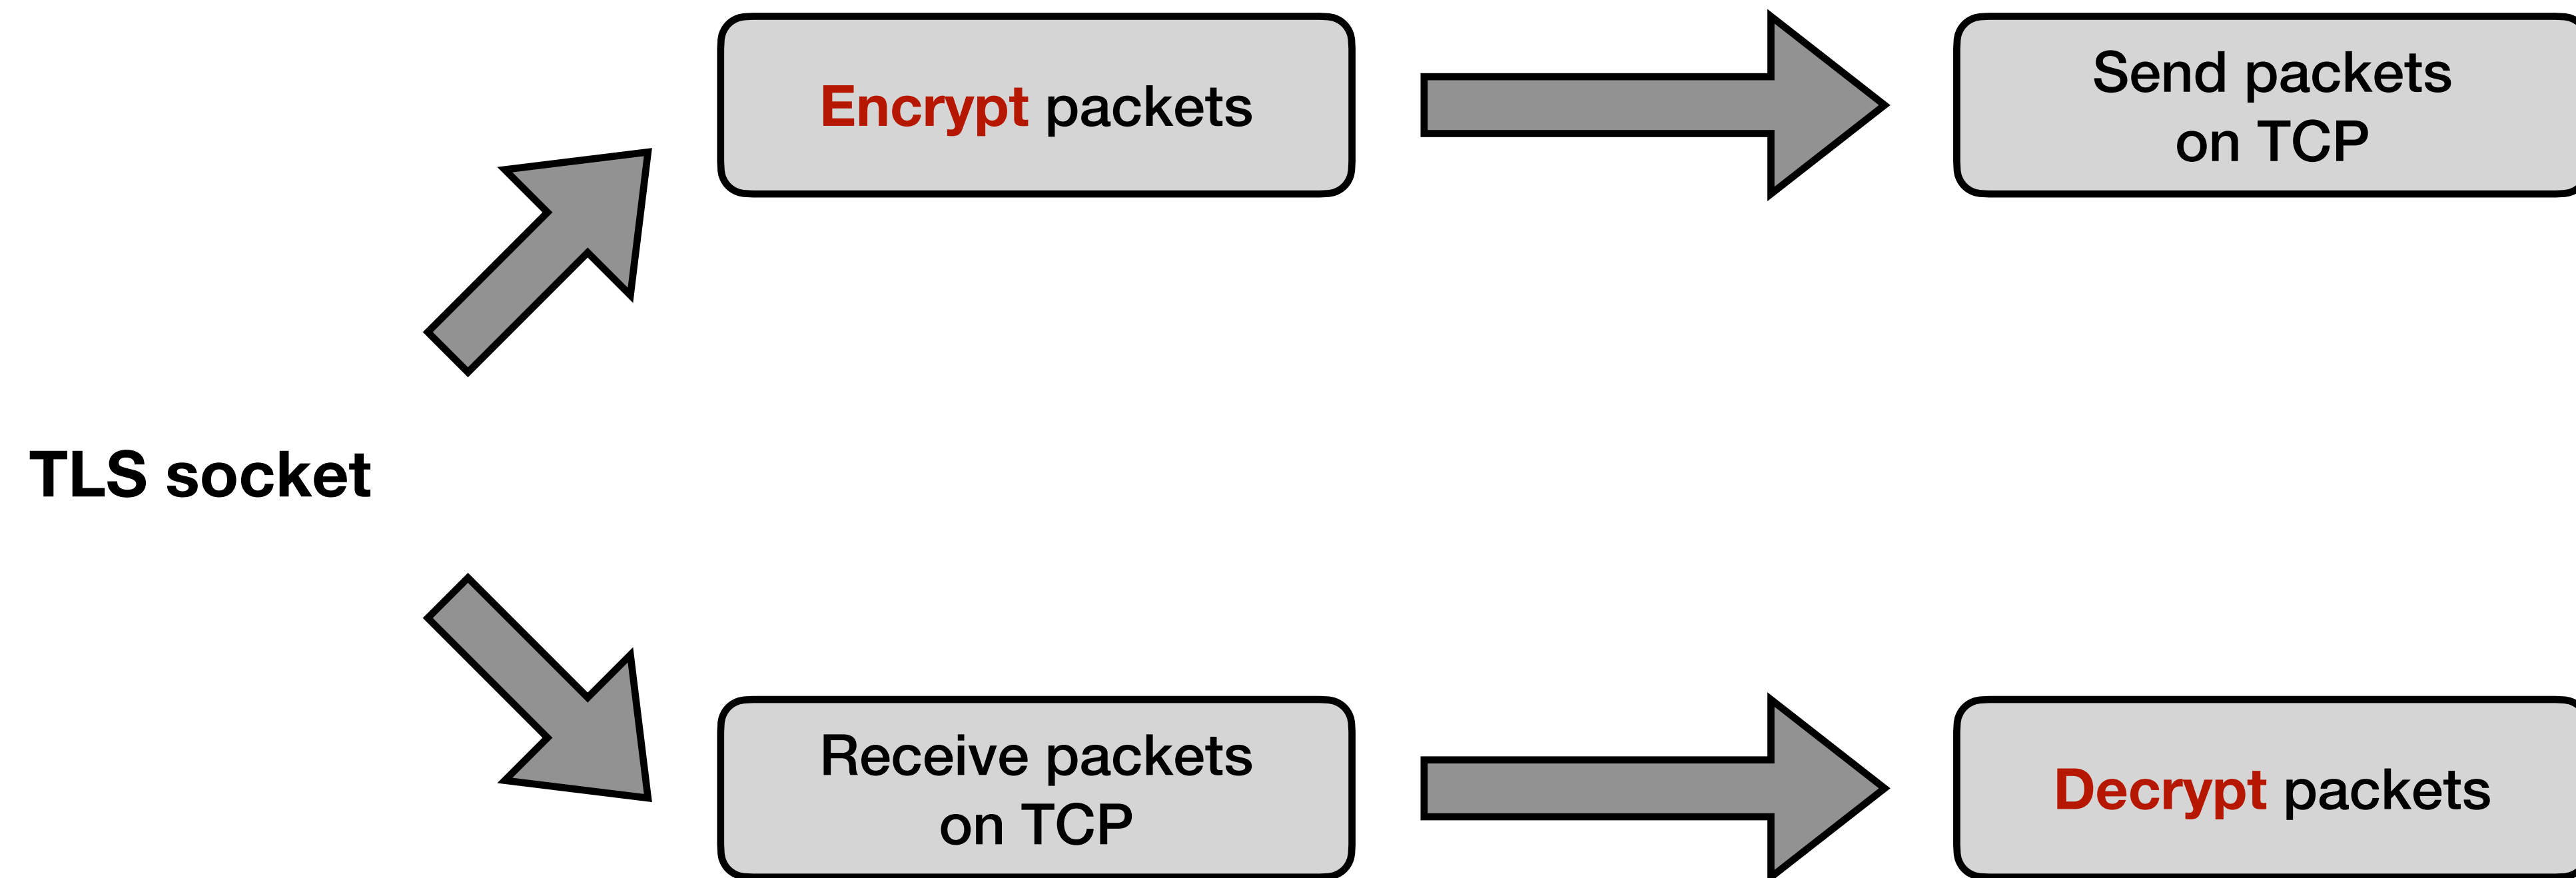
memcpy(crypto_info.key, "0123456789ABCDEF", TLS_CIPHER_AES_GCM_128_KEY_SIZE); // 16
memcpy(crypto_info.iv, "12345678", TLS_CIPHER_AES_GCM_128_IV_SIZE); // 8
memcpy(crypto_info.salt, "SALT", TLS_CIPHER_AES_GCM_128_SALT_SIZE); // 4

setsockopt(sockfd, SOL_TLS, TLS_TX, &crypto_info, sizeof(crypto_info));
setsockopt(sockfd, SOL_TLS, TLS_RX, &crypto_info, sizeof(crypto_info));
```

# Overview



# Overview



# Overview

- Supported TLS algorithms
  - gcm(aes)
  - ccm(aes)
  - gcm(sm4)
  - ...

# Overview

- Supported TLS algorithms
  - **gcm**(aes)
  - **ccm**(aes)
  - **gcm**(sm4)
  - ... **Template name**

# Overview

- Supported TLS algorithms
  - gcm(**aes**)
  - ccm(**aes**)
  - gcm(**sm4**)
  - ...     **Cipher name**

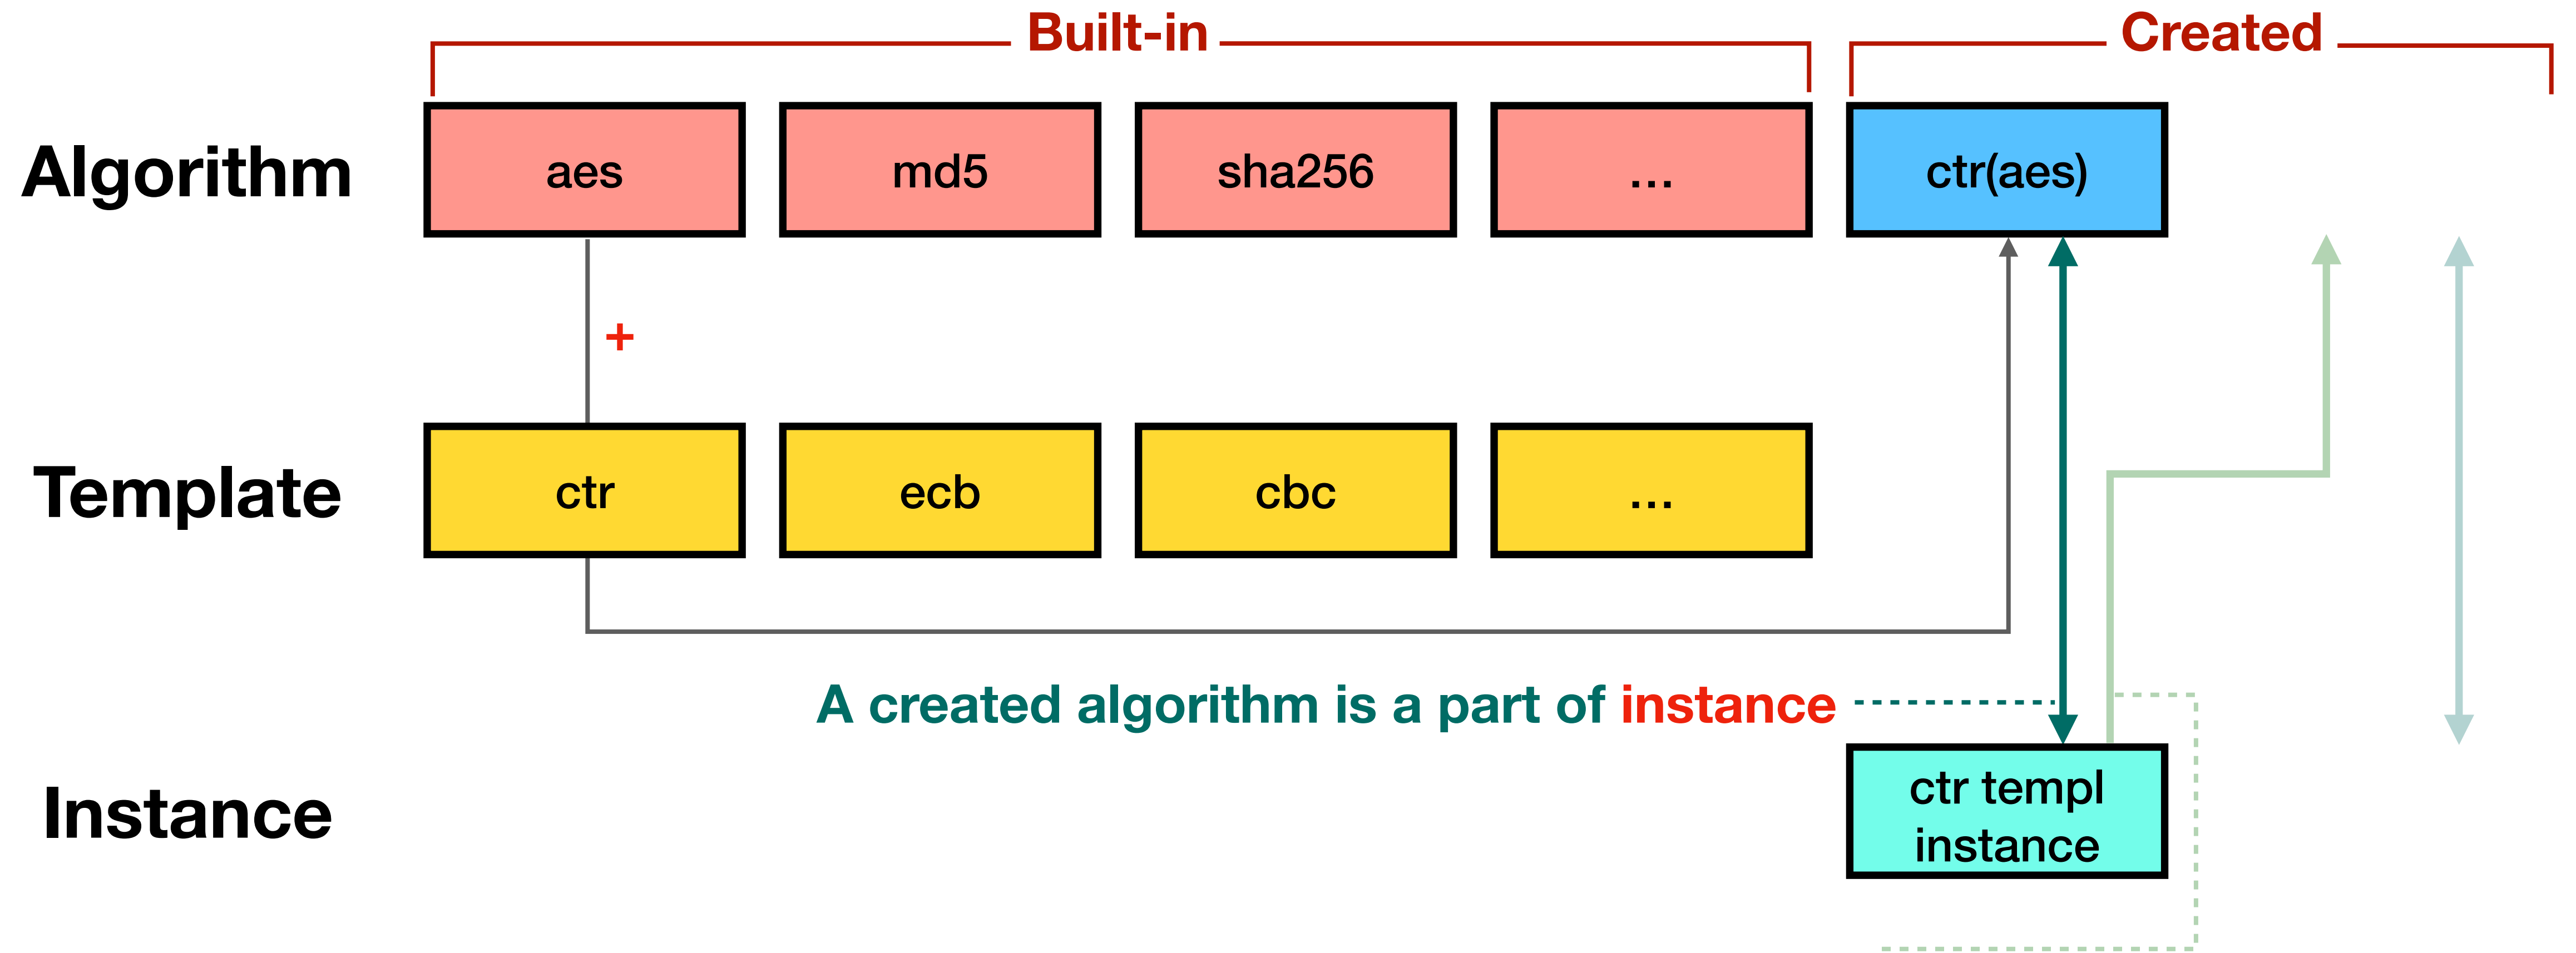
# Overview

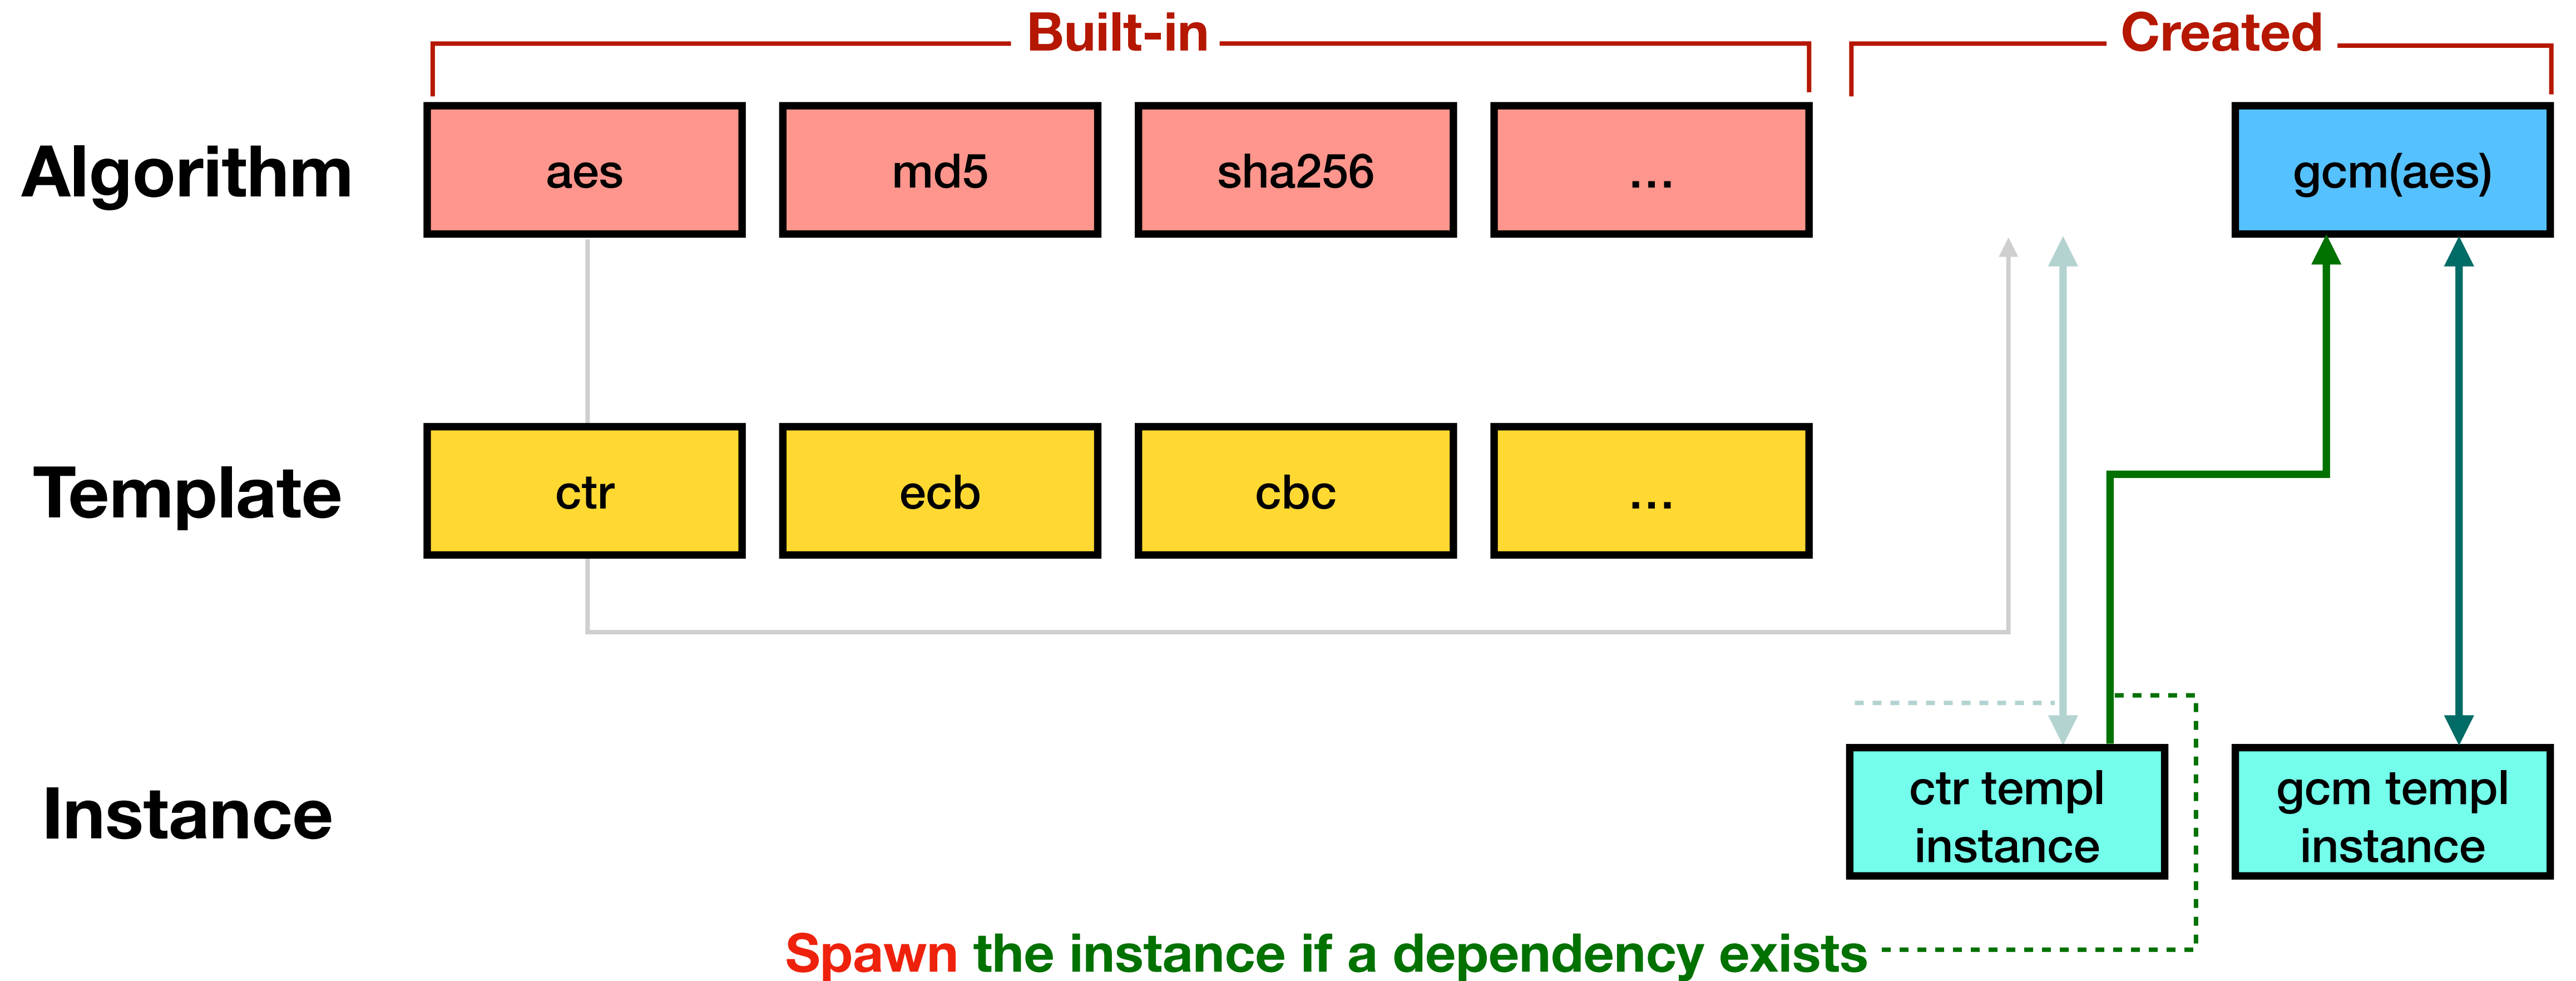
- **Algorithm**
  - Implementation of a specific **cryptographic operation**, such as AES, SHA-256, or HMAC
- **Template**
  - Constructing more complex cryptographic transformations by **combining** or **layering** simpler algorithms

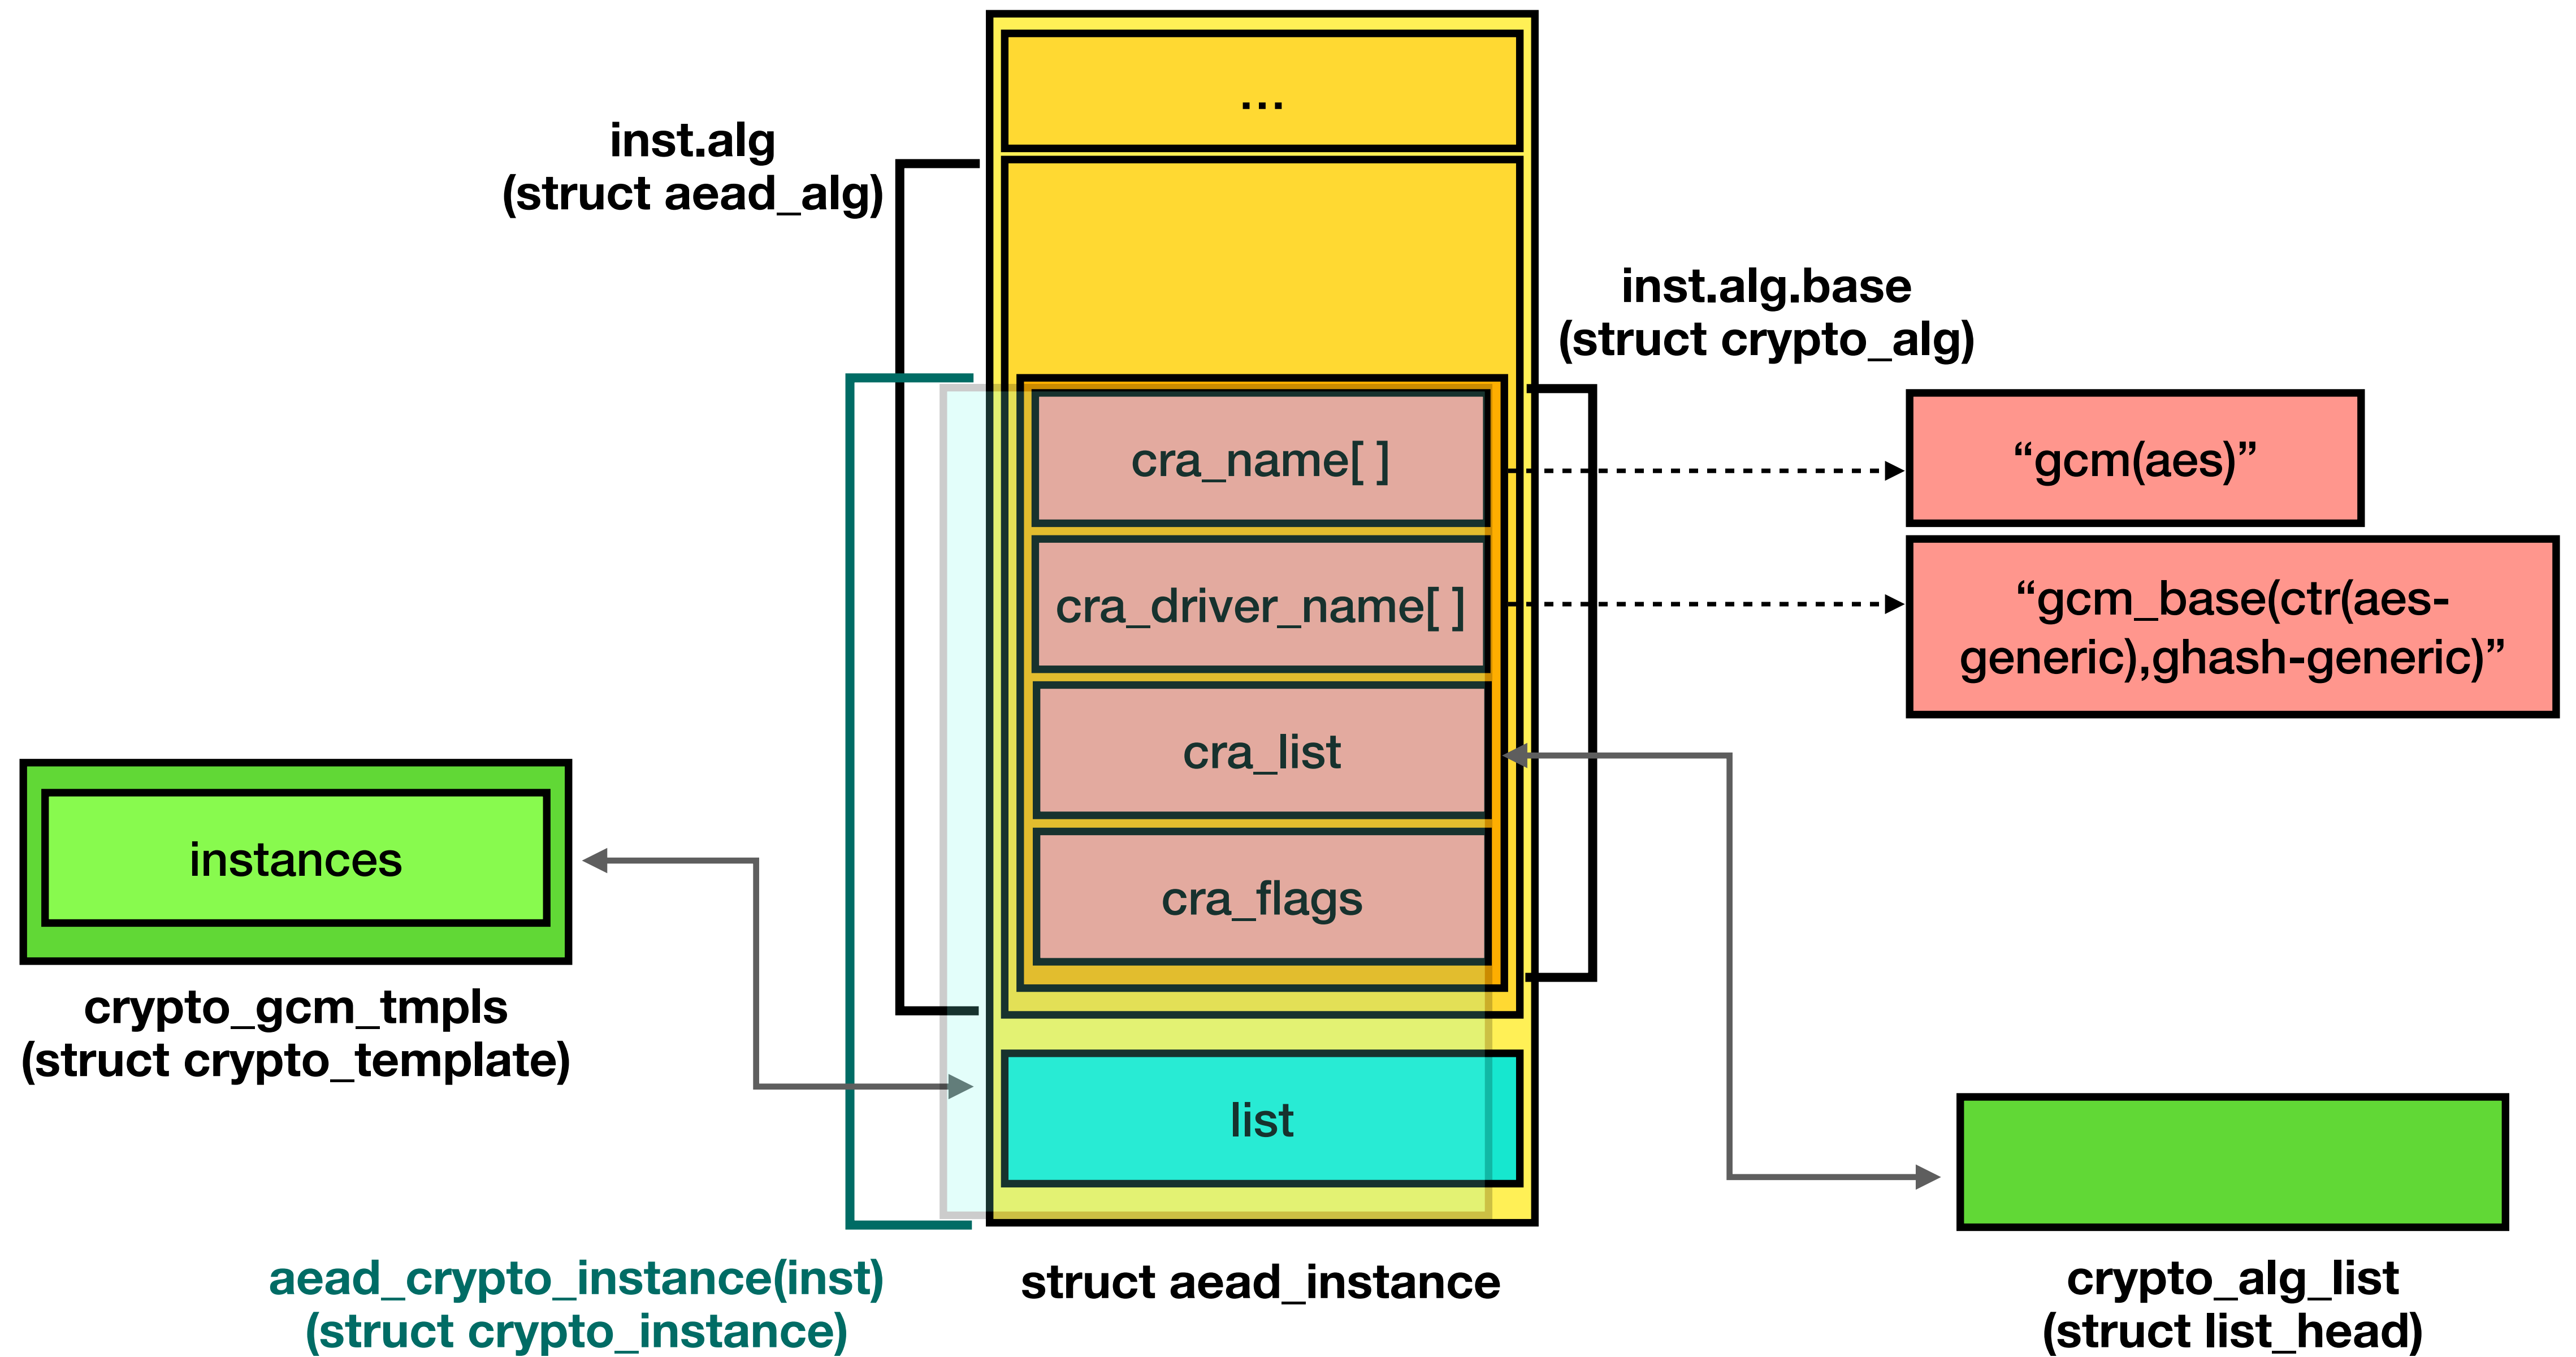


# Overview

- **Instance**
  - Instantiation of a cryptographic template, where specific algorithms and parameters **have been configured**
- **Spawn**
  - Create a linkage or **dependency** between cryptographic instances and algorithms







# Overview

- For example, if we configure “gcm(aes)” as the crypto algorithm of TX...

**Thread-A**

Find “gcm(aes)”

aes

md5

...

**Global variable**

**Thread-B**

Thread-A

Thread-B

Not found

Find “gcm(aes)”

aes

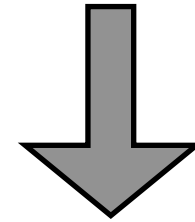
md5

...

Global variable

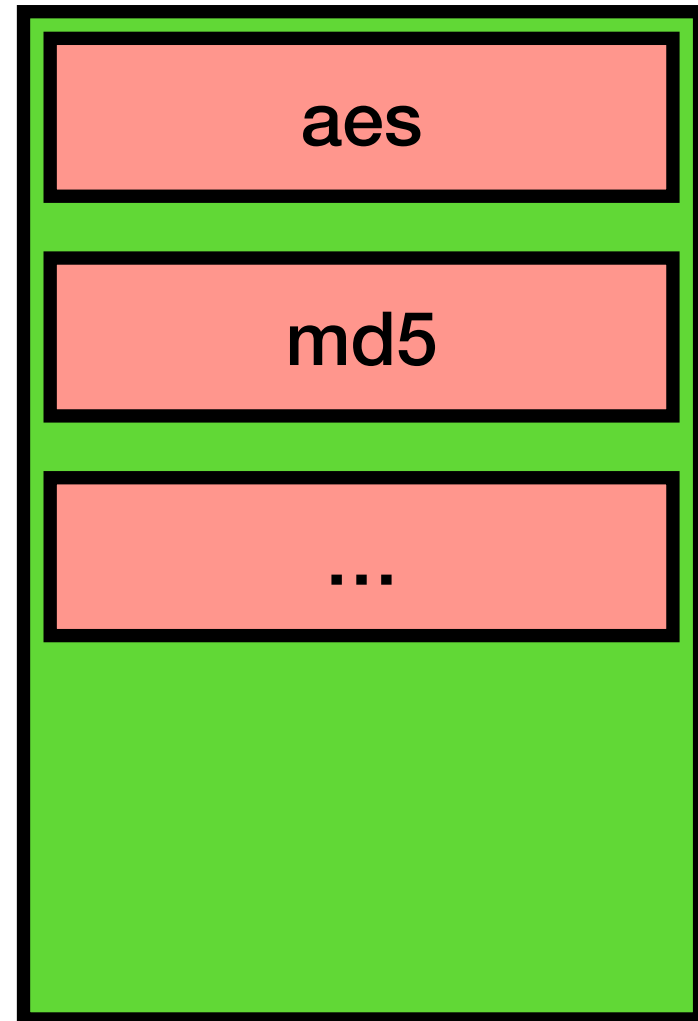
## Thread-A

Find “gcm(aes)”



Setup probe

## Thread-B

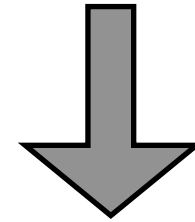


Global variable

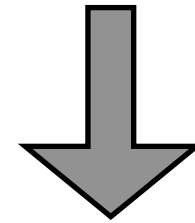


## Thread-A

Find “gcm(aes)”

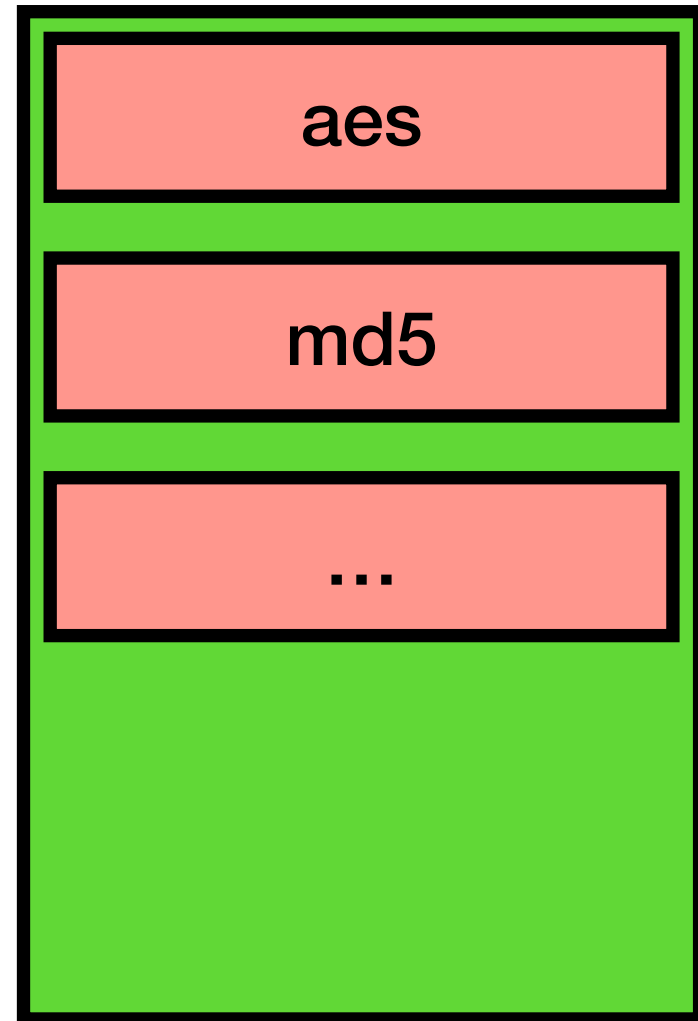


Setup probe



Dispatch probing

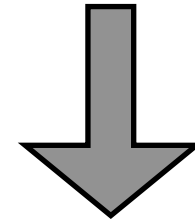
## Thread-B



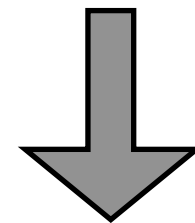
Global variable

## Thread-A

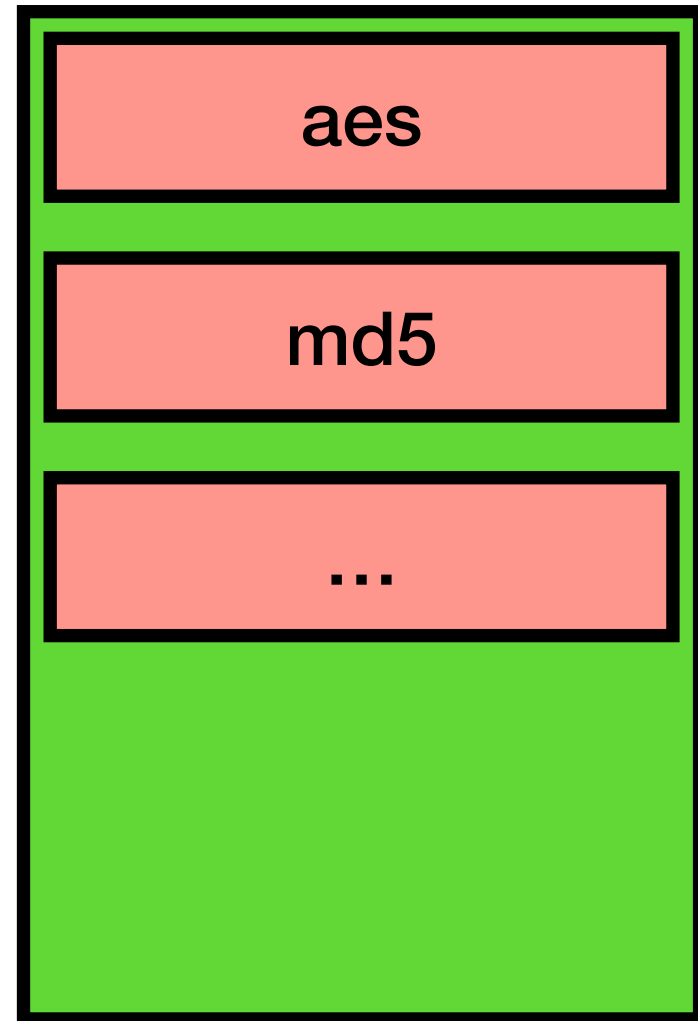
Find “gcm(aes)”



Setup probe



Dispatch probing



Global variable

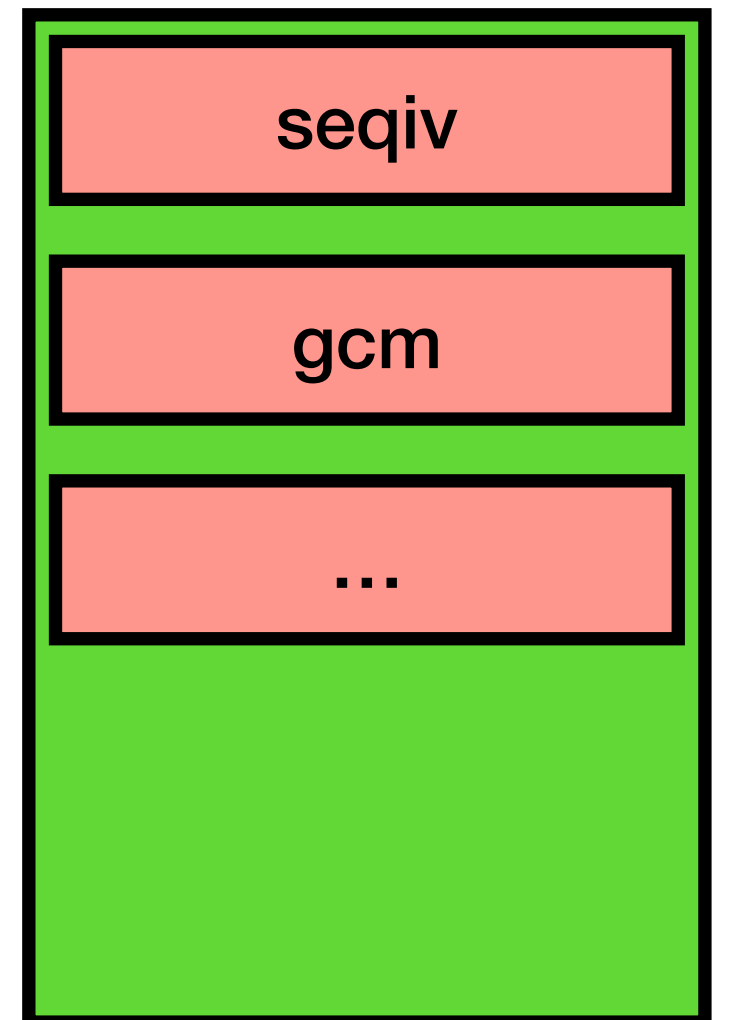
1. Template name = “gcm”
2. Cipher name = “aes”



## Thread-B

“cryptomgr\_probe”

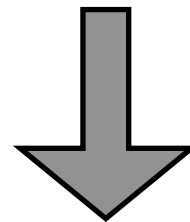
Find template



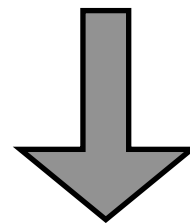
Global variable

## Thread-A

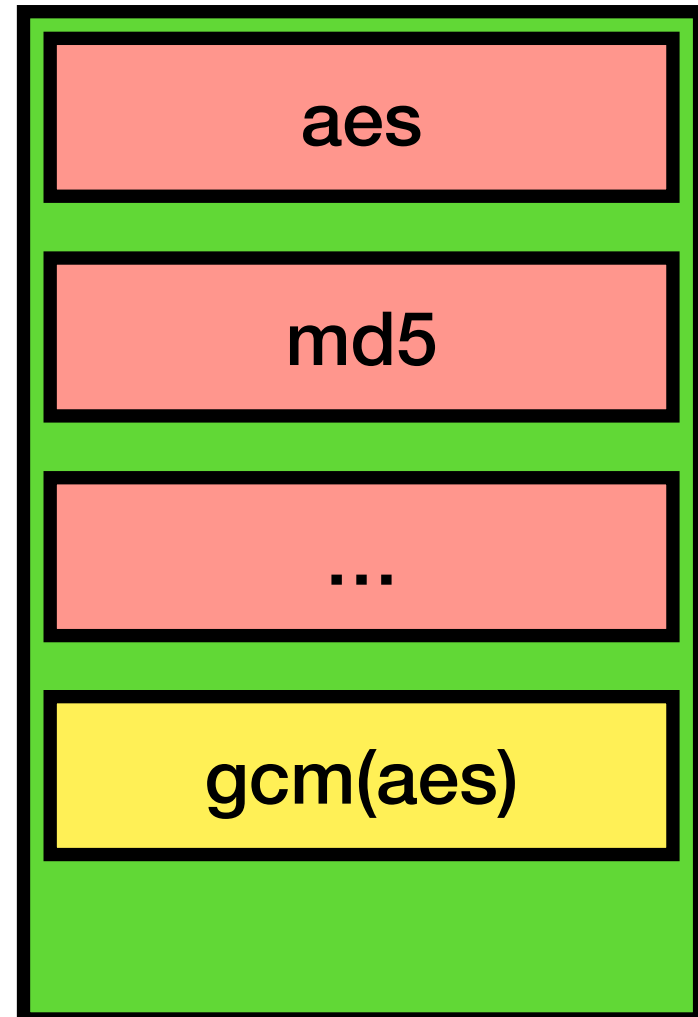
Find “gcm(aes)”



Setup probe



Dispatch probing



Global variable

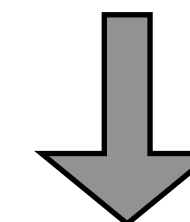
1. Template name = “gcm”
2. Cipher name = “aes”



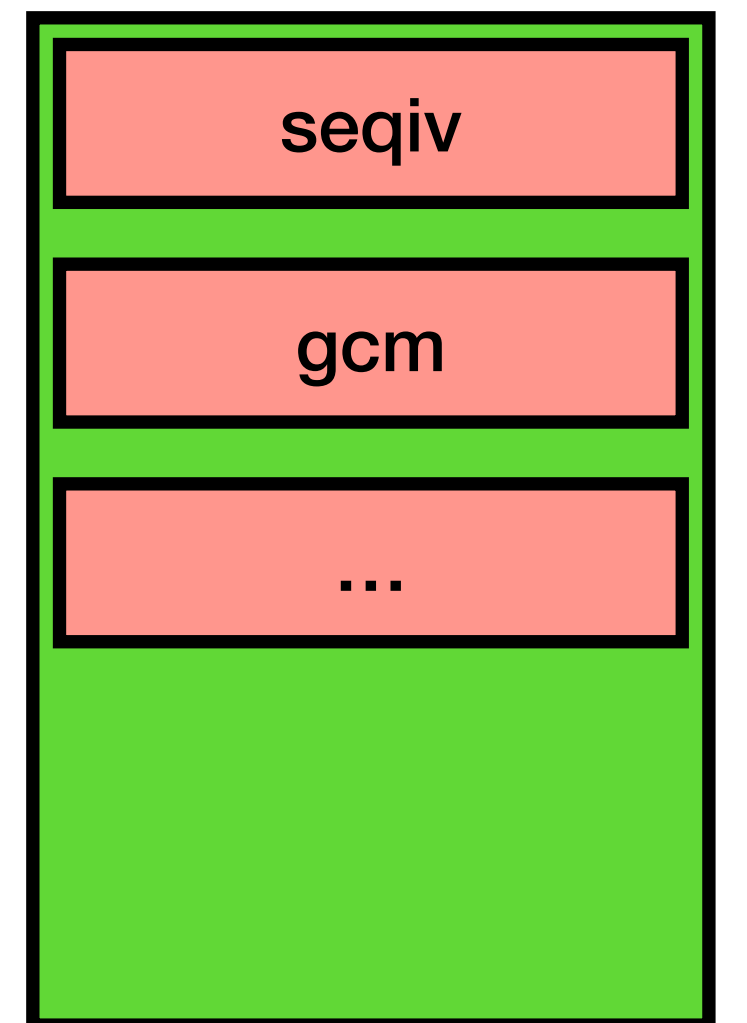
## Thread-B

“cryptomgr\_probe”

Find template

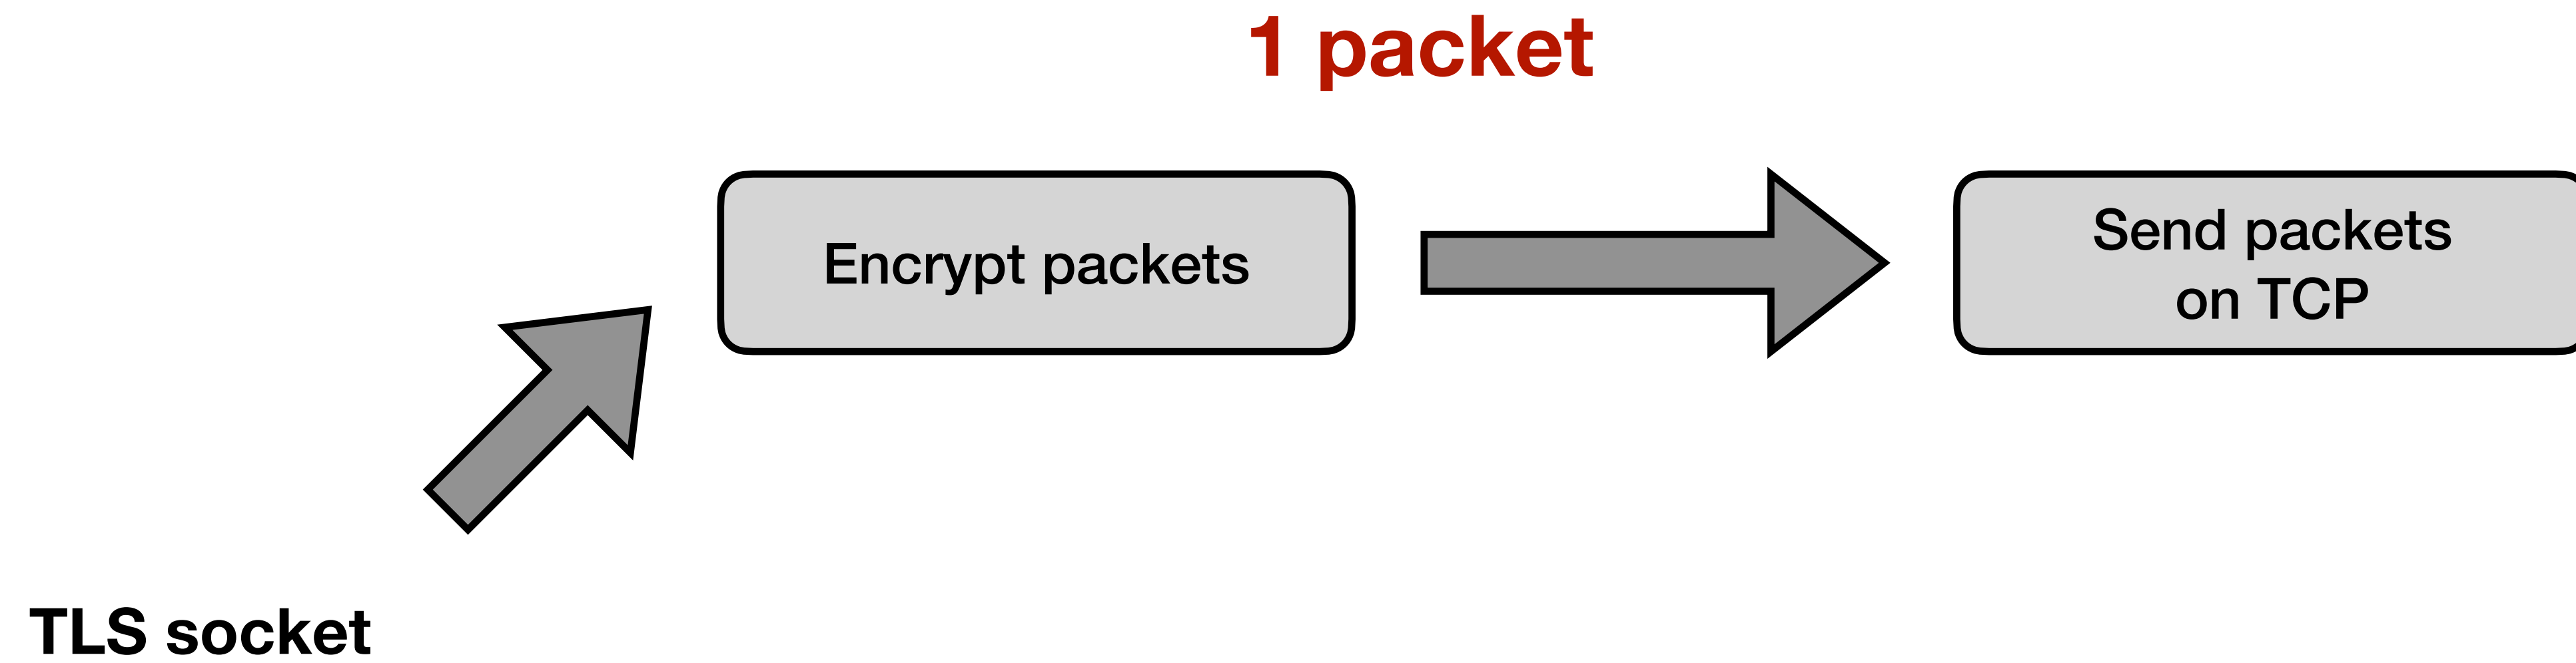


Create & initialize instance

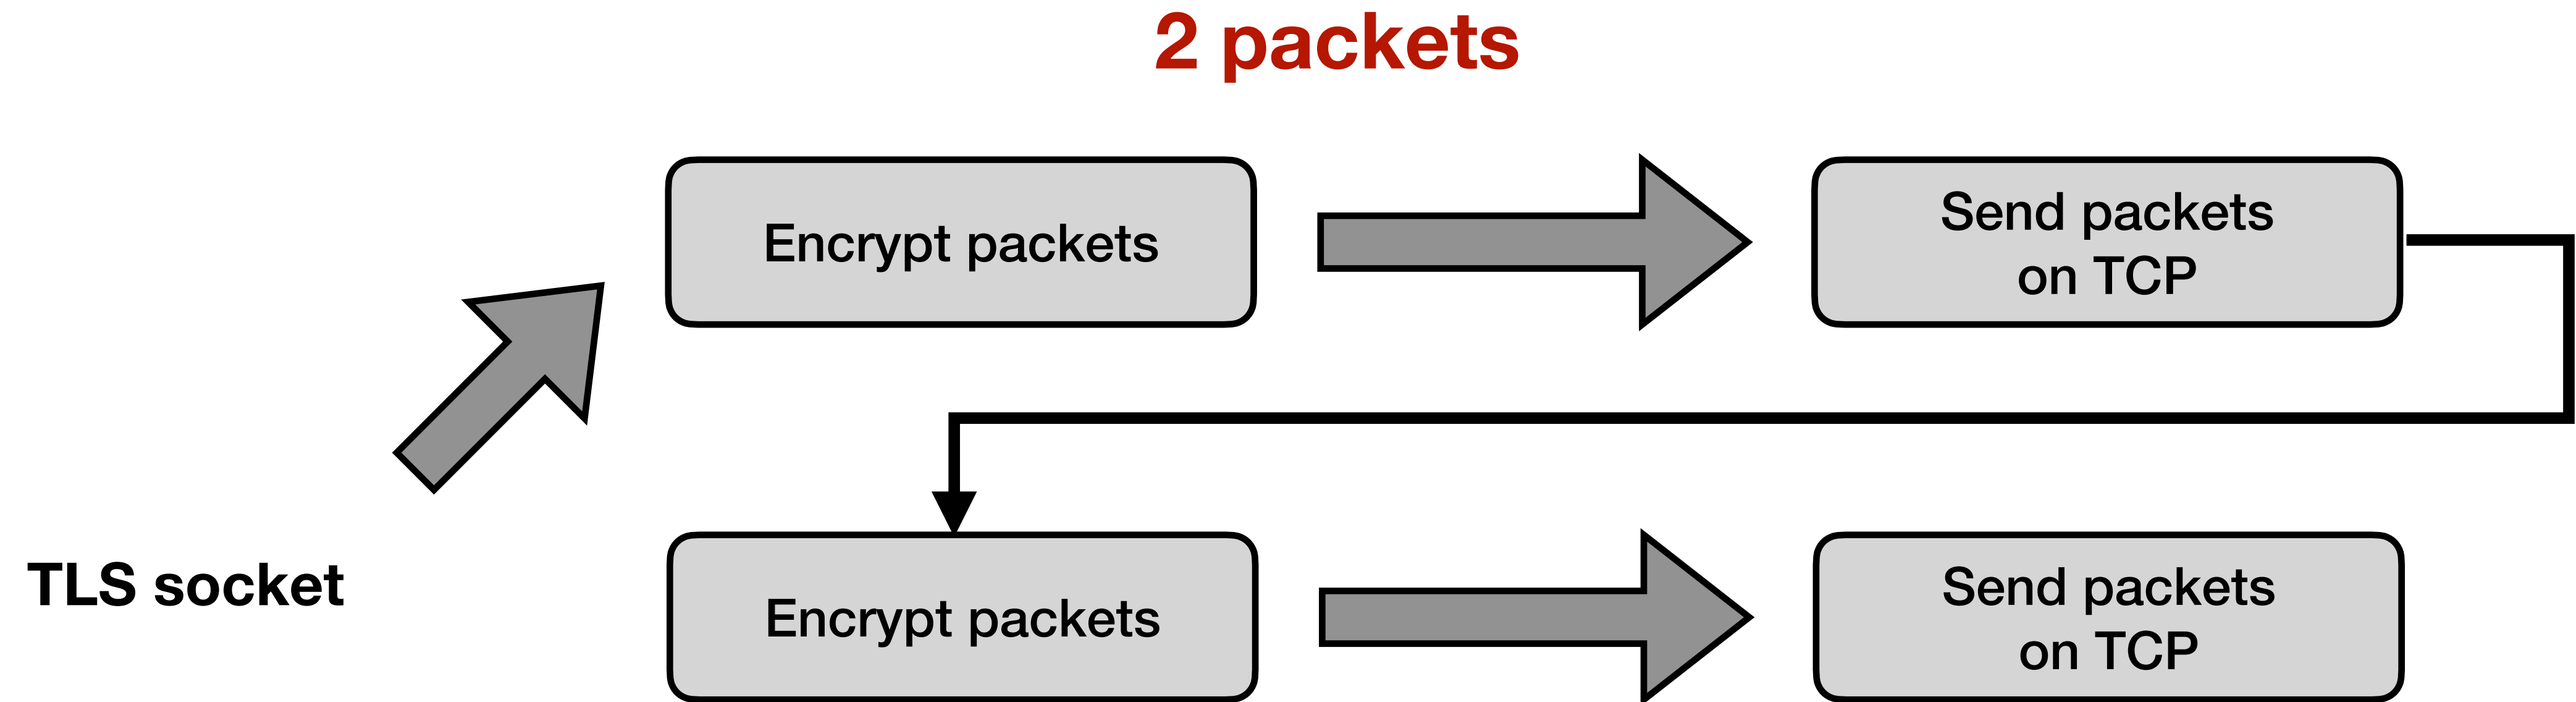


Global variable

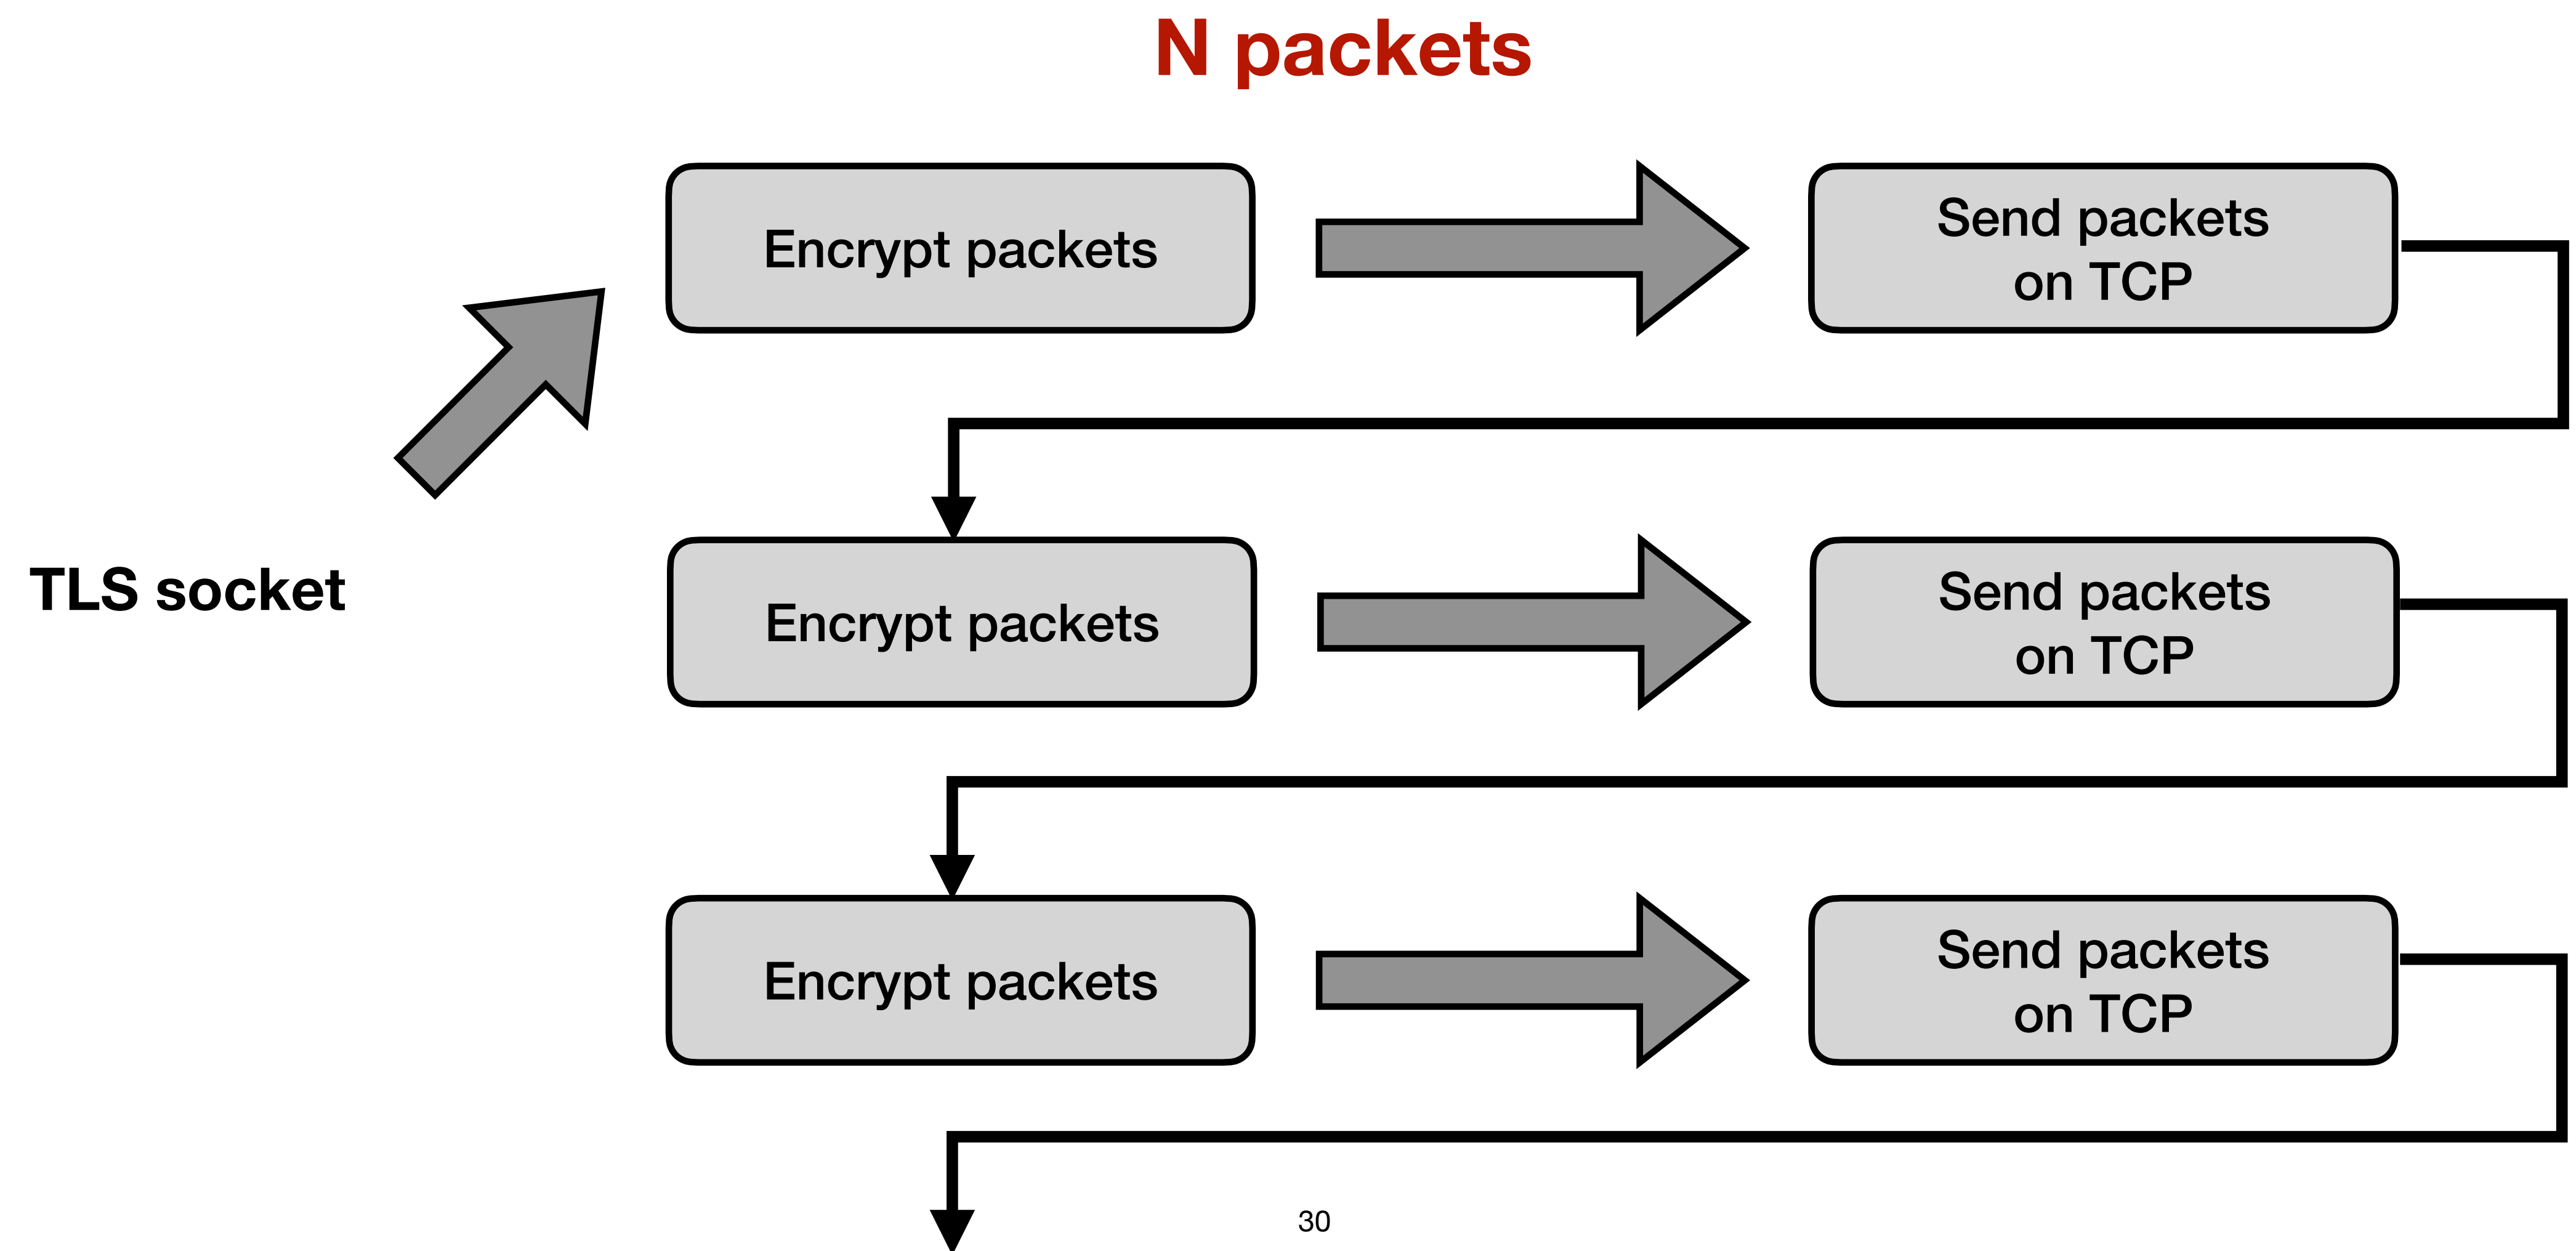
# Overview



# Overview



# Overview

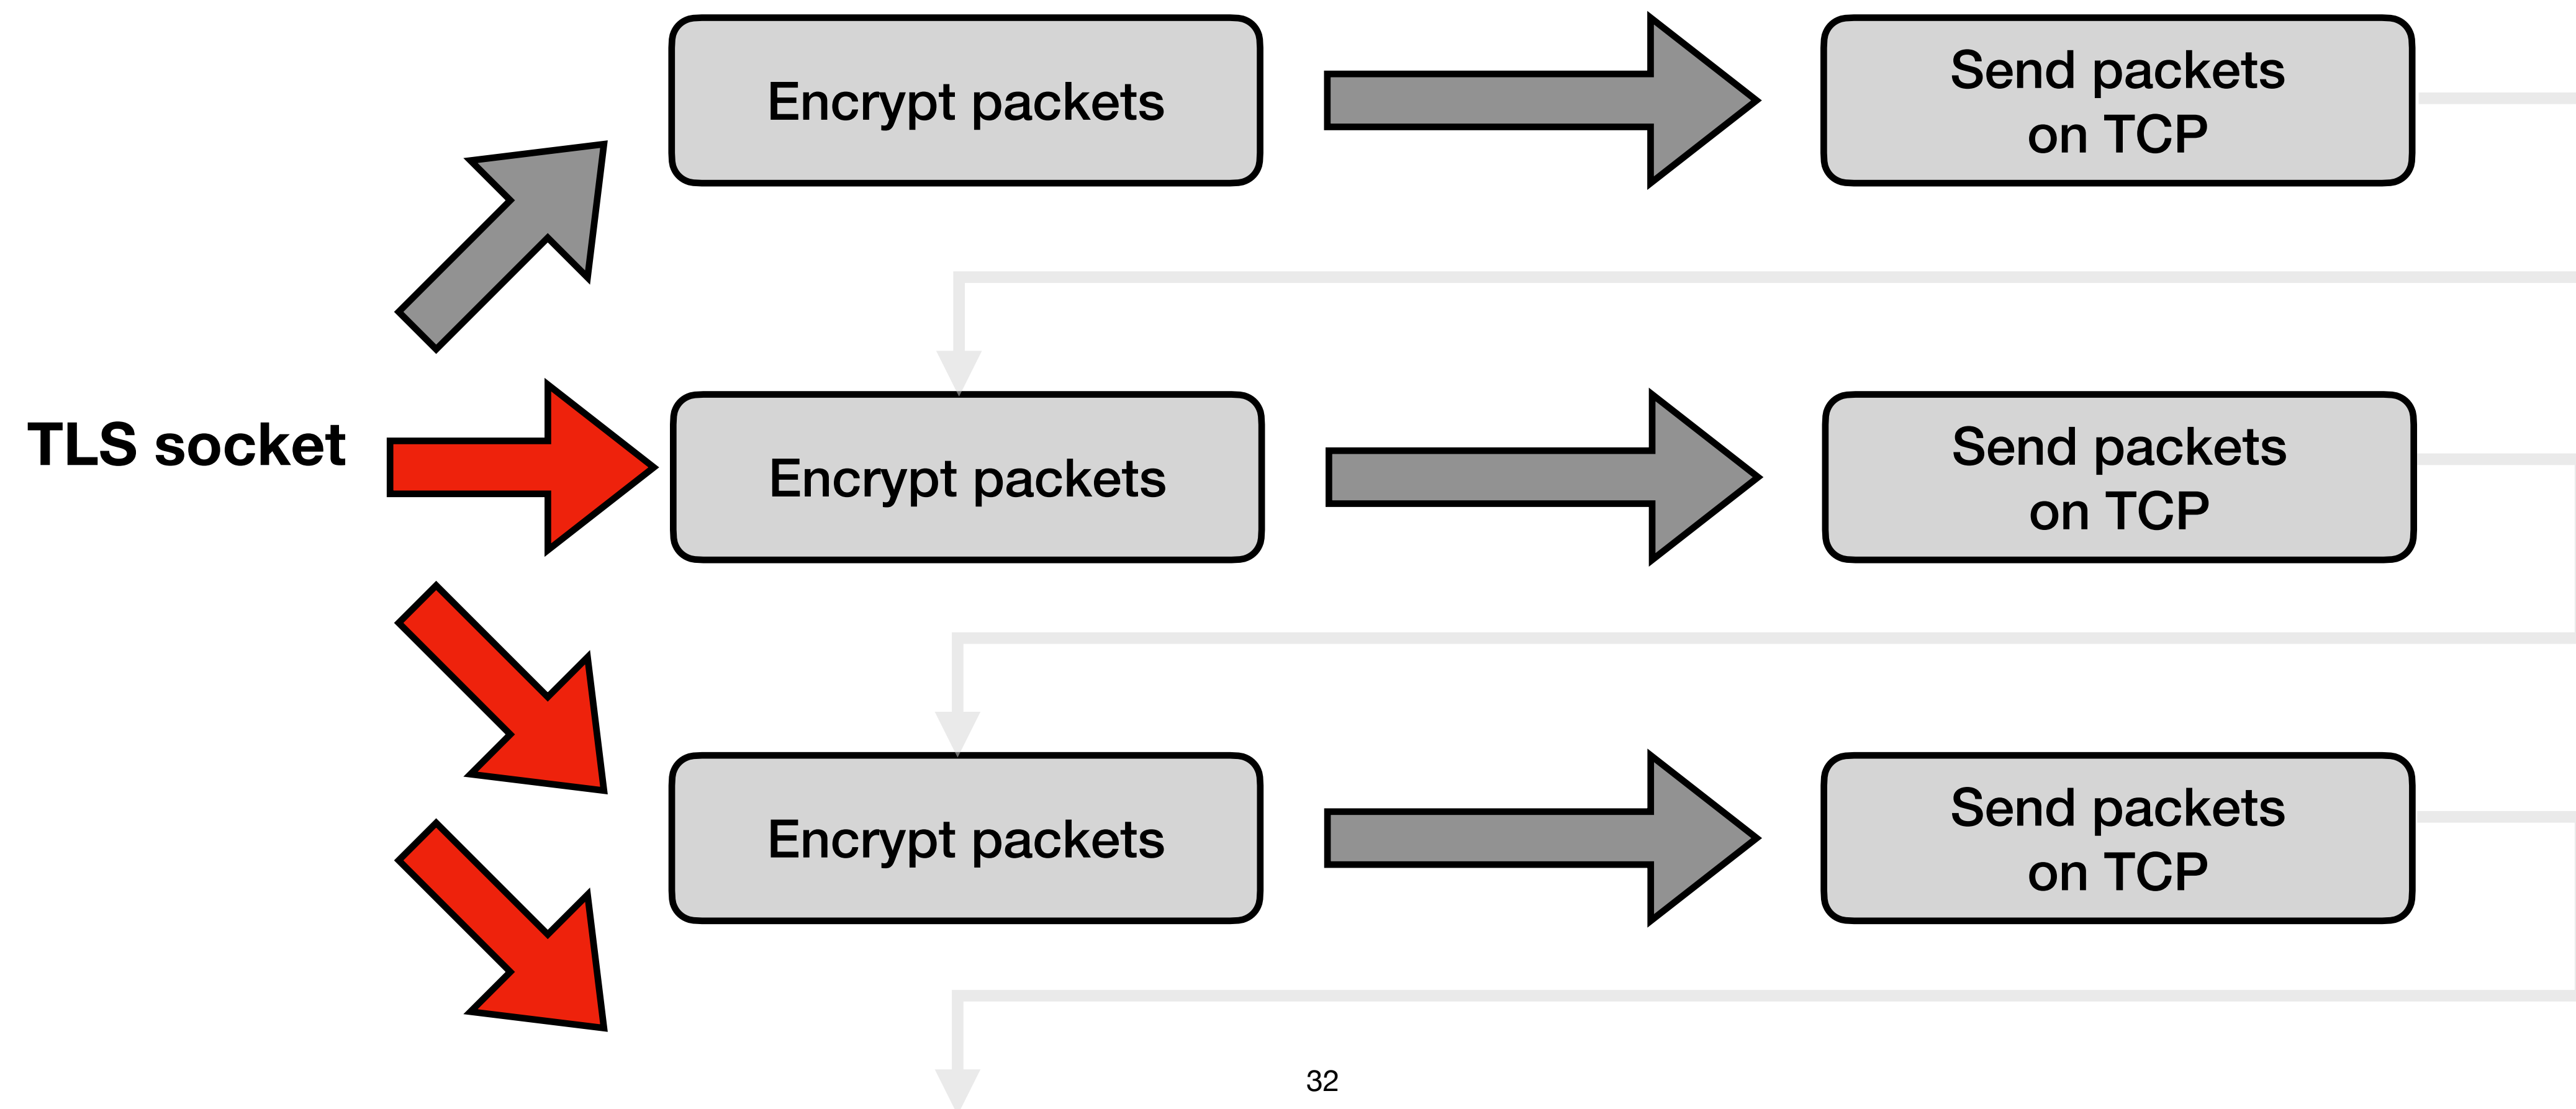


# Overview



# Overview

**N packets (asynchronous mode)**





# Overview

## Vendor specific drivers

```
✓ C aes_cbc.c drivers/crypto/vmx 1  
  CRYPTO_ALG_ASYNC);  
✓ C aes_ctr.c drivers/crypto/vmx 1  
  CRYPTO_ALG_ASYNC);  
✓ C aes_xts.c drivers/crypto/vmx 1  
  CRYPTO_ALG_ASYNC);  
✓ C zynqmp-aes-gcm.c drivers/crypto/xilinx 1  
  CRYPTO_ALG_ASYNC |  
✓ C dm-verity-target.c drivers/md 1  
  v->use_tasklet ? CRYPTO_ALG_ASYNC : 0);  
✓ C ppp_mppe.c drivers/net/ppp 1  
  ...crypto_has_ahash("sha1", 0, CRYPTO_ALG_ASYNC))  
✓ C tcp.c drivers/nvme/host 1  
  ...crypto_alloc_ahash("crc32c", 0, CRYPTO_ALG_ASYNC);  
✓ C tcp.c drivers/nvme/target 1  
  ...crypto_alloc_ahash("crc32c", 0, CRYPTO_ALG_ASYNC);  
✓ C iscsi_tcp.c drivers/scsi 1  
  ...crypto_alloc_ahash("crc32c", 0, CRYPTO_ALG_ASYNC);
```



# Overview

- **Cryptd**
  - Enabled when the CONFIG\_CRYPTD\_CRYPTD compile option is set
  - A **crypto daemon** which converts an arbitrary synchronous crypto algorithm into an **asynchronous algorithm** that runs in a **kthread**





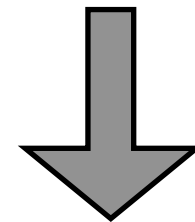
# Overview

- **Cryptd**
  - Enabled when the CONFIG\_CRYPTO\_CRYPTD compile option is set
  - A **crypto daemon** which converts an arbitrary synchronous crypto algorithm into an **asynchronous algorithm** that runs in a **kthread**
  - Used as a **template**

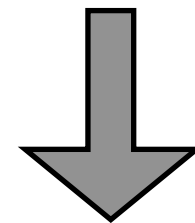


## Thread-A

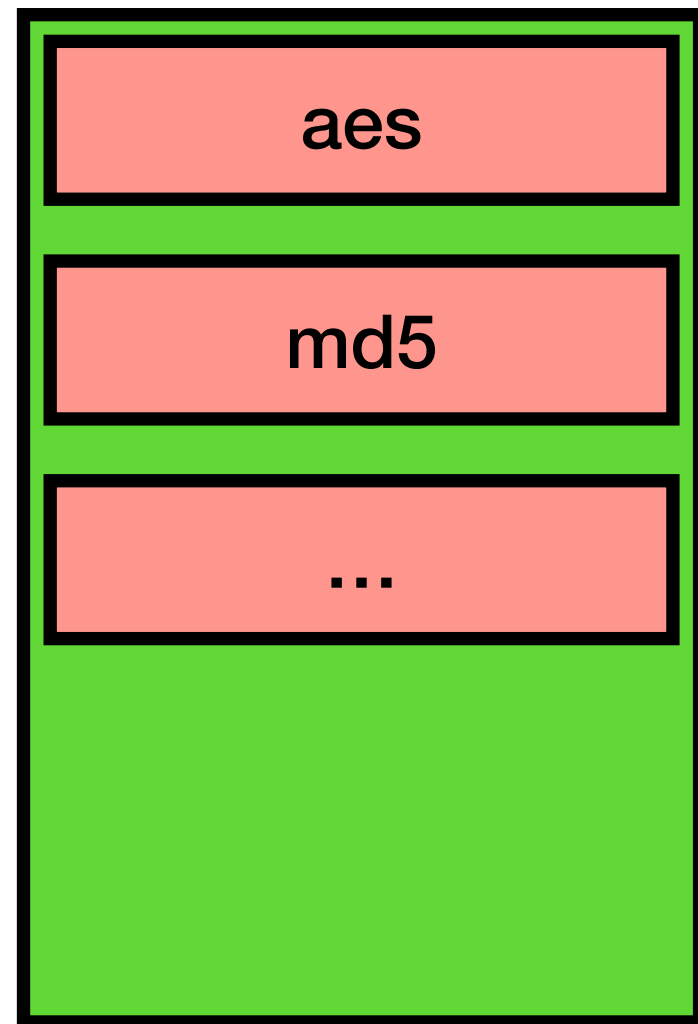
Find “cryptd(XXX)”



Setup probe



Dispatch probing



Global variable

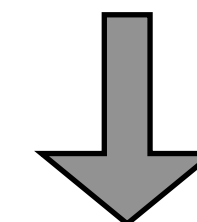
1. Template name = “**cryptd**”
2. Cipher name = “**XXX**”



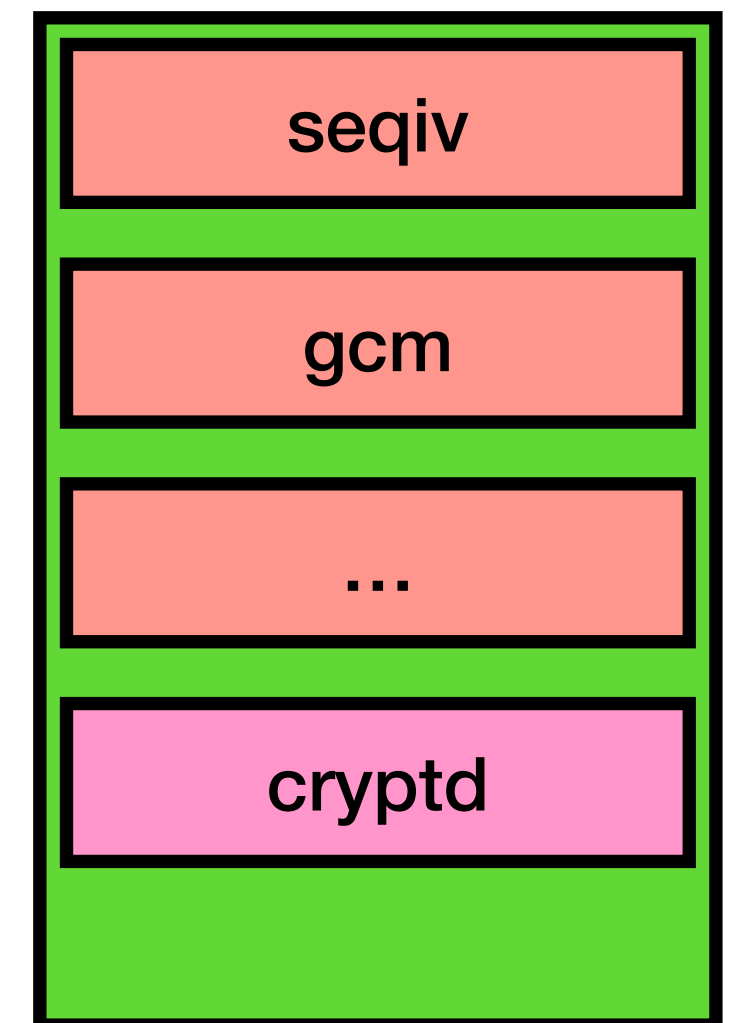
## Thread-B

“**cryptomgr\_probe**”

Find template



Create & initialize instance



Global variable

## Thread-A

```
static int cryptd_create_aead(struct crypto_template *tmpl,
                             struct rtattr **tb,
                             struct crypto_attr_type *algt,
                             struct cryptd_queue *queue)
{
    struct aead_instance_ctx *ctx;
    struct aead_instance *inst;

    // [...]
    inst = kzalloc(sizeof(*inst) + sizeof(*ctx), GFP_KERNEL);
    ctx = aead_instance_ctx(inst);

    // [...]
    inst->alg.base.cra_flags |= CRYPTO_ALG_ASYNC
        (alg->base.cra_flags & CRYPTO_ALG_INTERNAL);

    // [...]
    err = aead_register_instance(tmpl, inst);

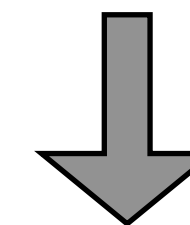
    // [...]
}
```

Global

## Thread-B

“cryptomgr\_probe”

Find template



Create & initialize  
instance

seqiv

gcm

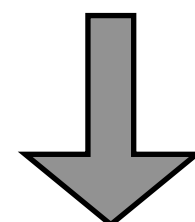
...

cryptd

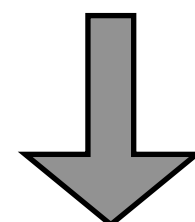
Global variable

## Thread-A

Find “cryptd(XXX)”



Setup probe



Dispatch probing

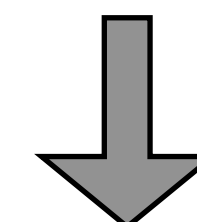
1. Template name = “**cryptd**”
2. Cipher name = “**XXX**”



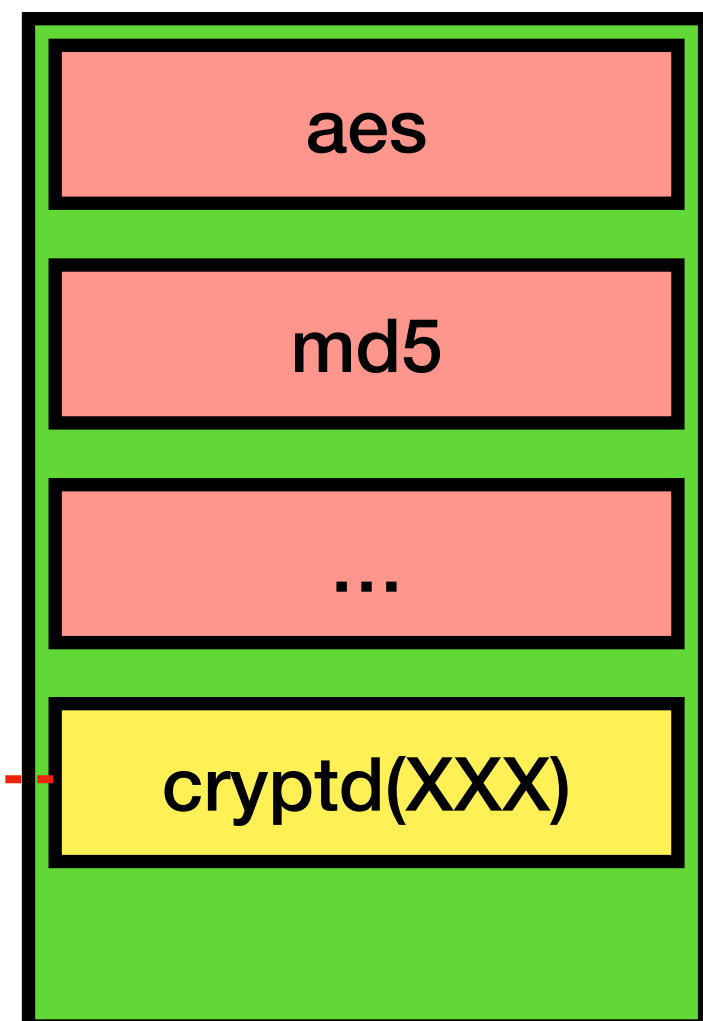
## Thread-B

“**cryptomgr\_probe**”

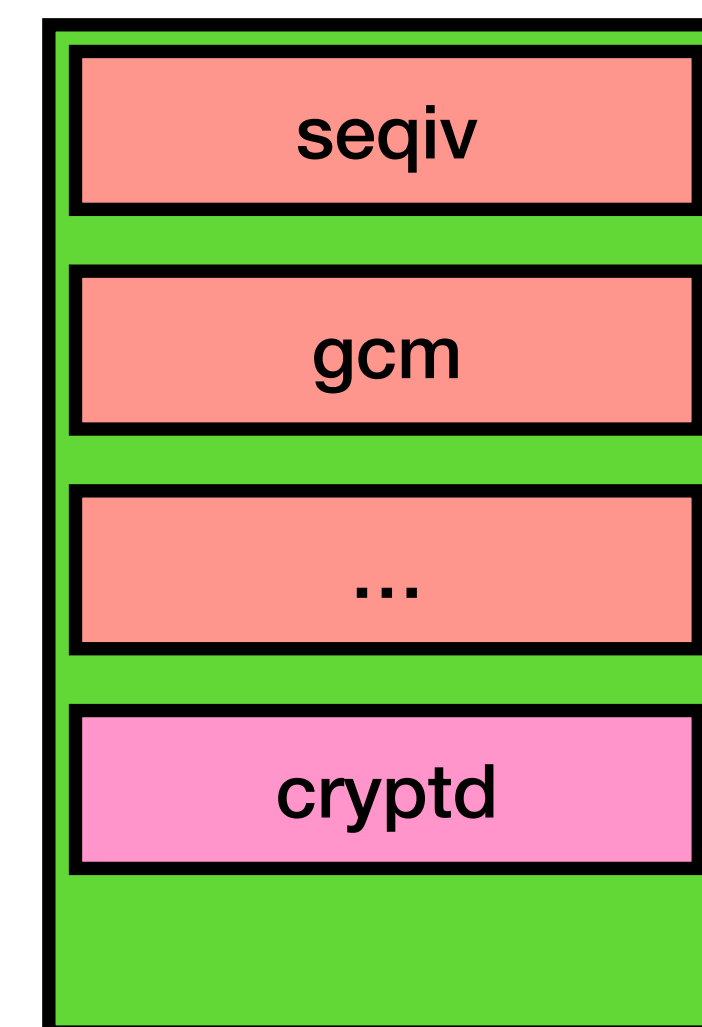
Find template



Create & initialize  
instance



Global variable



Global variable

... Same as the original algorithm (XXX in here) but with asynchronous mode enabled

# Overview

- But how?

```
const struct tls_cipher_desc tls_cipher_desc[TLS_CIPHER_MAX + 1 - TLS_CIPHER_MIN] = {  
    TLS_CIPHER_AES_GCM_128, ..., "gcm(aes)"  
    TLS_CIPHER_AES_GCM_256, ..., "gcm(aes)"  
    TLS_CIPHER_AES_CCM_128, ..., "ccm(aes)"  
    TLS_CIPHER_CHACHA20_POLY1305, ..., "rfc7539(chacha20,poly1305)"  
    TLS_CIPHER_SM4_GCM, ..., "gcm(sm4)"  
    TLS_CIPHER_SM4_CCM, ..., "ccm(sm4)"  
    TLS_CIPHER_ARIA_GCM_128, ..., "gcm(aria)"  
    TLS_CIPHER_ARIA_GCM_256, ..., "gcm(aria)"  
};
```



# Overview

- **AF\_ALG**
  - Interface to kernel crypto API
  - Algorithm probing with user-provided algorithm name

```
#include <linux/if_alg.h>

int sock = socket(AF_ALG, SOCK_SEQPACKET, 0);
struct sockaddr_alg sa = {
    .salg_family = AF_ALG,
    .salg_type = "aead",
    .salg_name = "cryptd(gcm(aes))",
};
bind(sock, (struct sockaddr *)&sa, sizeof(sa));
```



**Thread-A**

Encrypt a packet

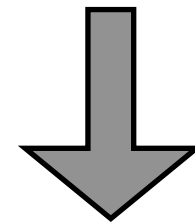
**Thread-B**  
**(cryptd\_queue\_worker)**

**Pending queue**



**Thread-A**

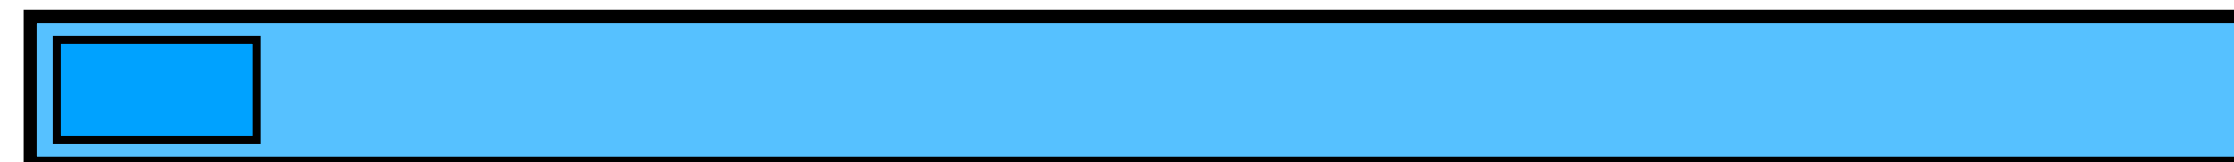
Encrypt a packet



Enqueue request  
(Cryptd)



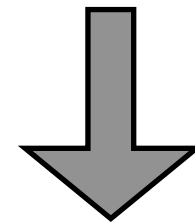
**Pending queue**



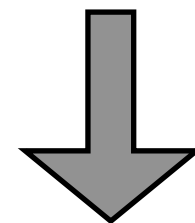
**Thread-B**  
(cryptd\_queue\_worker)

**Thread-A**

Encrypt a packet



Encqueue request  
(Cryptd)



Wakeup worker



**Thread-B**  
(cryptd\_queue\_worker)

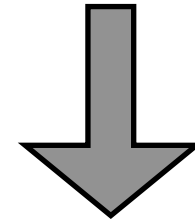


**Pending queue**

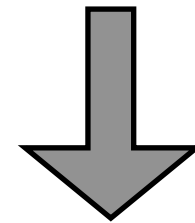


## Thread-A

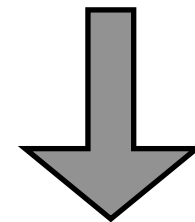
Encrypt a packet



Encqueue request  
(Cryptd)



Wakeup worker



Return

## Thread-B (cryptd\_queue\_worker)

Dequeue request

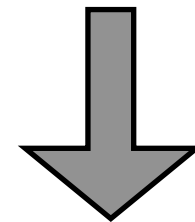


Pending queue

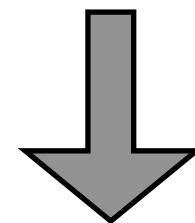


## Thread-A

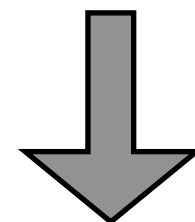
Encrypt a packet



Encqueue request  
(Cryptd)



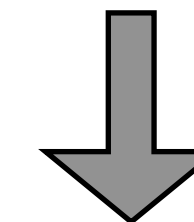
Wakeup worker



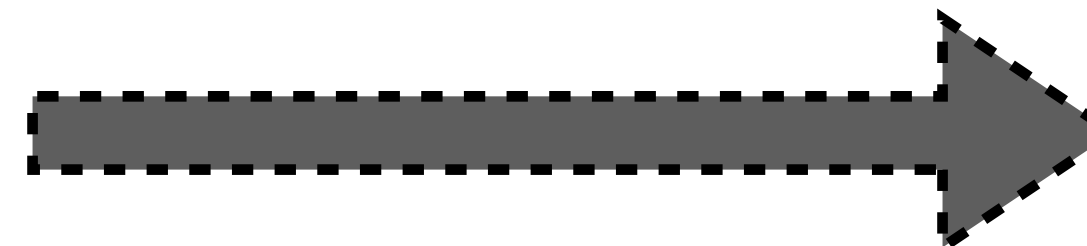
Return

## Thread-B (cryptd\_queue\_worker)

Dequeue request



Handle request




Pending queue



# Vulnerability

# Vulnerability



<b>CVE-2024-26800</b>	2024-02-29	tls: fix use-after-free on failed backlog decryption
	2024-02-10	Merge branch 'tls-fixes'
	2024-02-10	net: tls: fix returned read length with async decrypt
	2024-02-10	selftests: tls: use exact comparison in recv_partial
<b>CVE-2024-26582</b>	2024-02-10	net: tls: fix use-after-free with partial reads and async decrypt
<b>CVE-2024-26584</b>	2024-02-10	net: tls: handle backlogging of crypto requests
<b>CVE-2024-26583</b>	2024-02-10	tls: fix race between tx work scheduling and socket close
<b>CVE-2024-26585</b>	2024-02-10	tls: fix race between async notify and socket close
	2024-02-10	net: tls: factor out tls_*crypt_async_wait()

# Vulnerability

2024-02-29    tls: fix use-after-free on failed backlog decryption

2024-02-10    Merge branch 'tls-fixes'

2024-02-10    net: tls: fix returned read length with async decrypt

2024-02-10    selftests: tls: use exact comparison in recv\_partial

2024-02-10    net: tls: fix use-after-free with partial reads and async decrypt

2024-02-10    net: tls: handle backlogging of crypto requests

2024-02-10    tls: fix race between tx work scheduling and socket close

**CVE-2024-26585**

2024-02-10    tls: fix race between async notify and socket close

2024-02-10    net: tls: factor out tls\_\*crypt\_async\_wait()



# Vulnerability

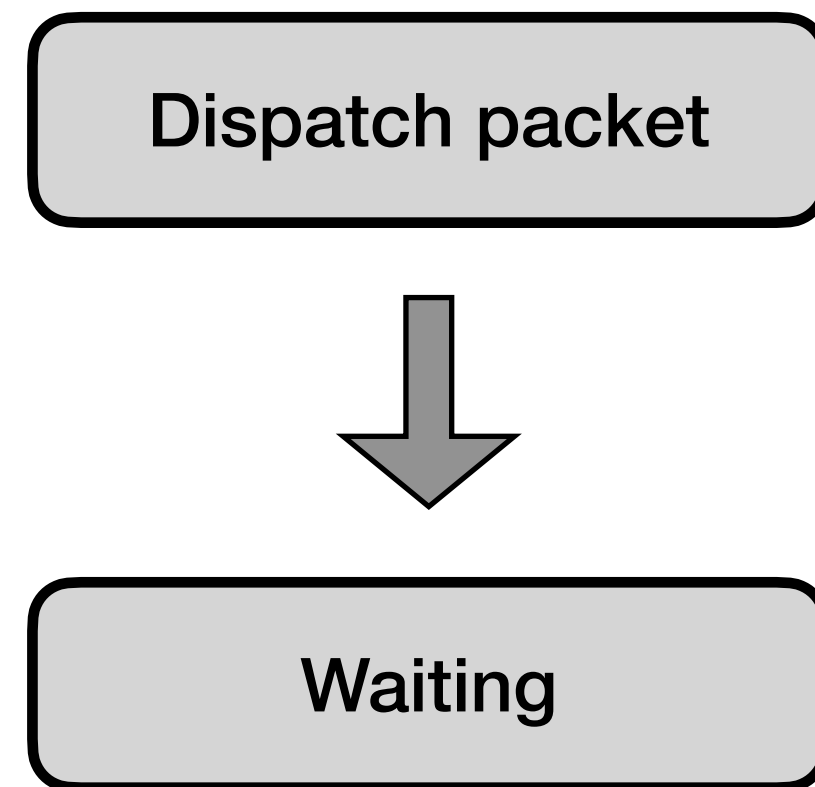
**CVE-2024-26585**

## **tls: fix race between async notify and socket close**

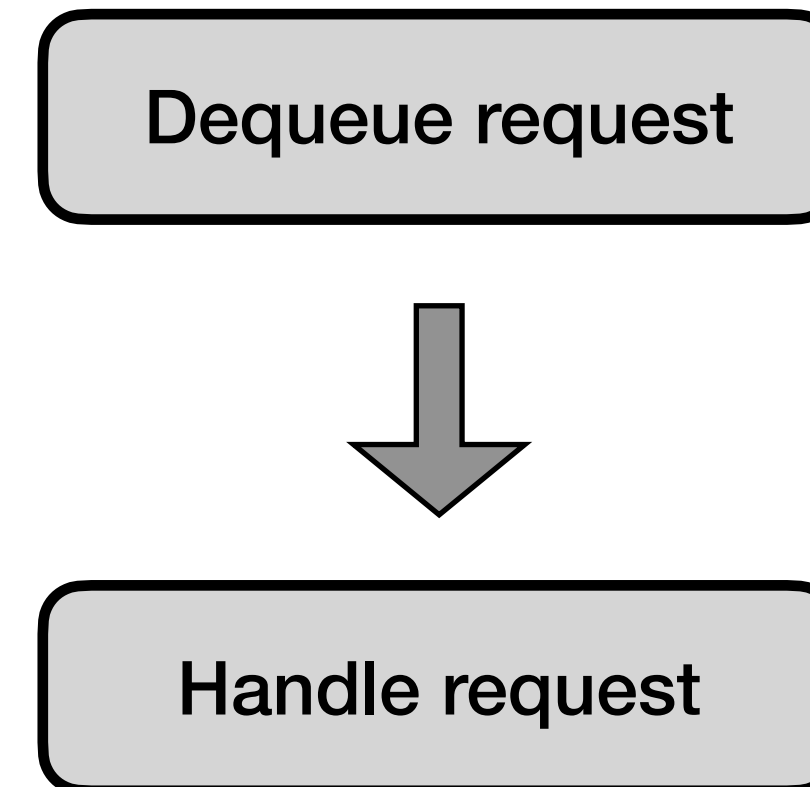
The submitting thread (one which called `recvmsg/sendmsg`) may exit as soon as the async crypto handler calls `complete()` so any code past that point risks touching already freed data.

Try to avoid the locking and extra flags altogether. Have the main thread hold an extra reference, this way we can depend solely on the atomic ref counter for synchronization.

## Thread-A

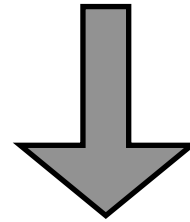


## Thread-B (cryptd\_queue\_worker)



## Thread-A

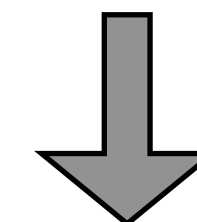
Dispatch packet



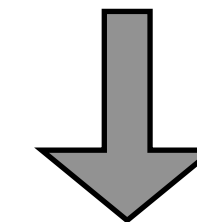
Waiting

## Thread-B (cryptd\_queue\_worker)

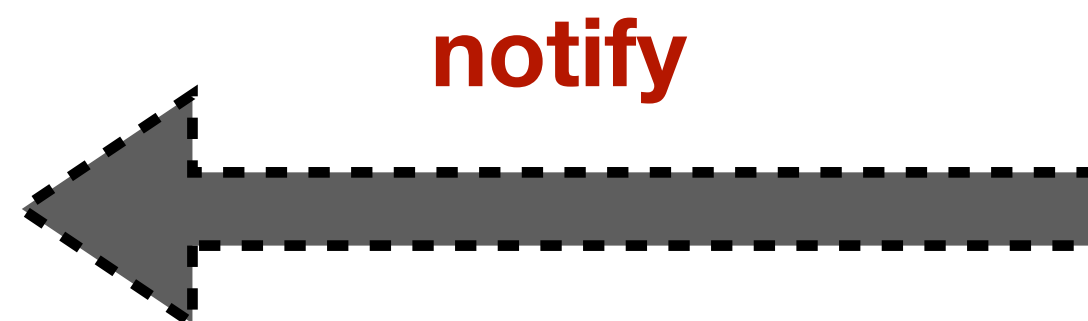
Dequeue request



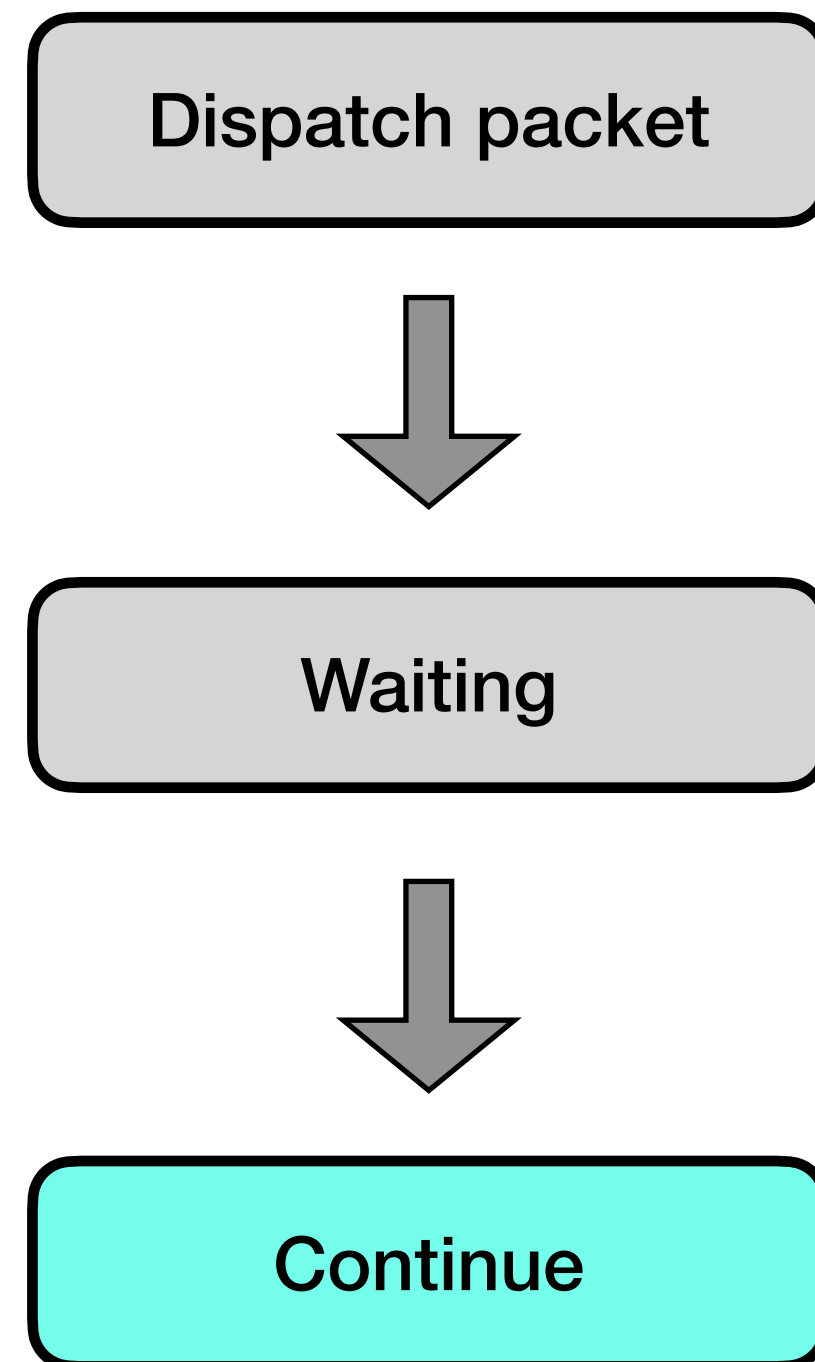
Handle request



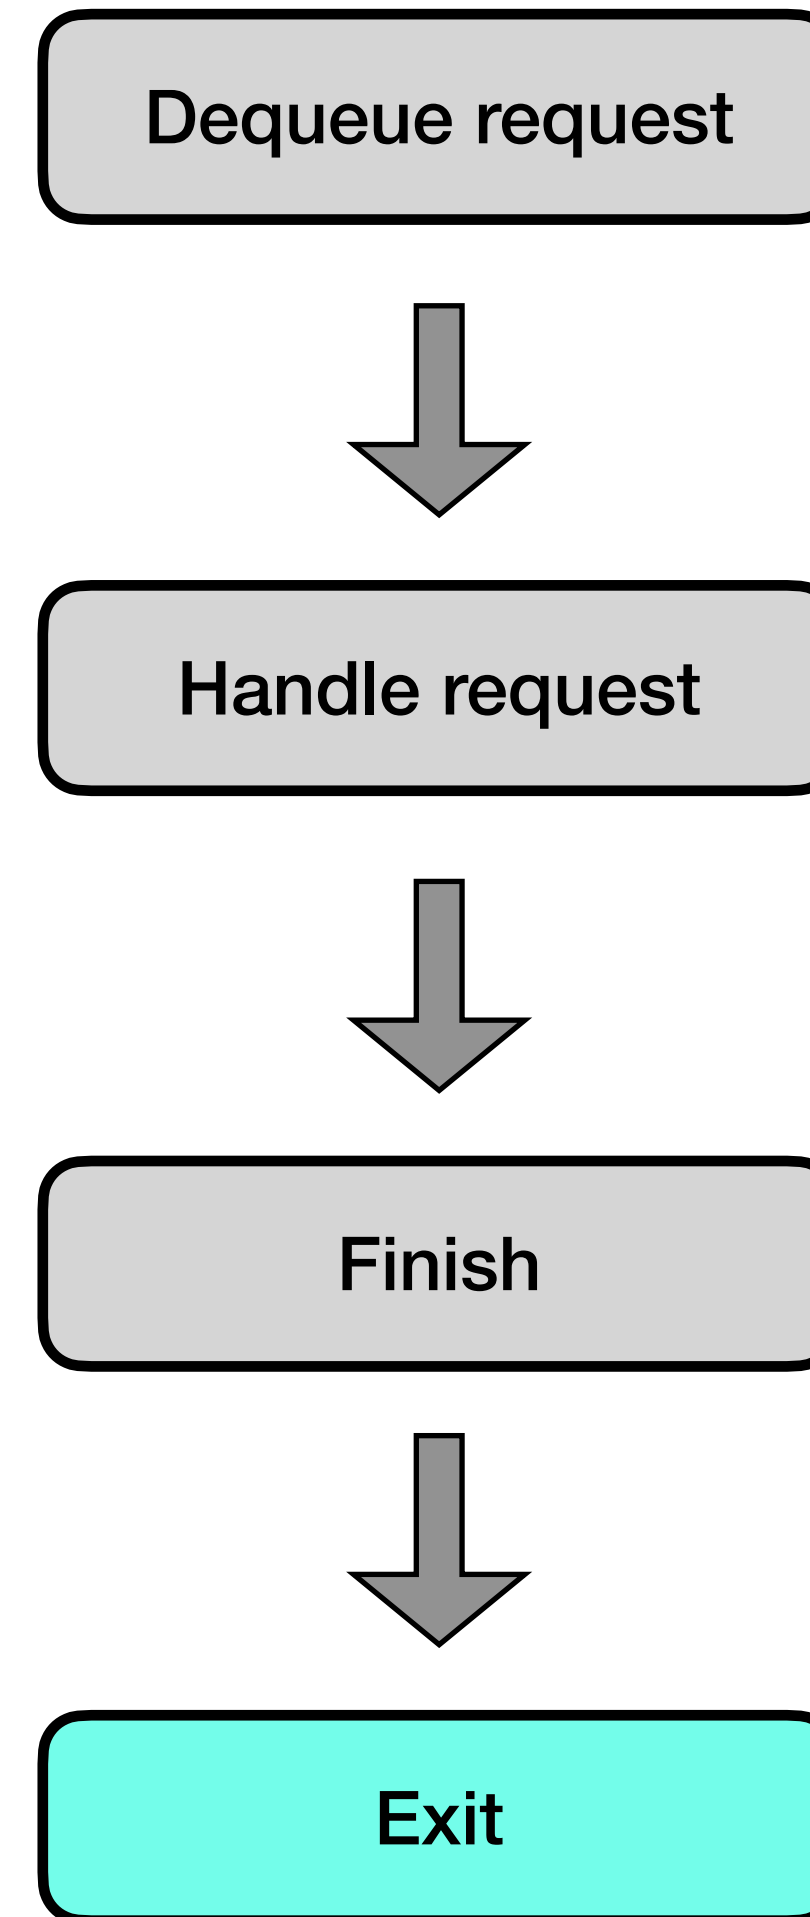
Finish



## Thread-A

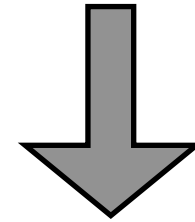


## Thread-B (cryptd\_queue\_worker)

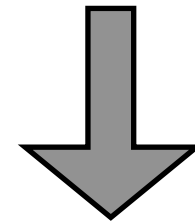


## Thread-A

Dispatch packet



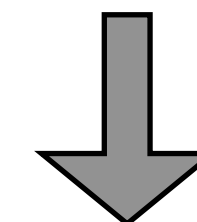
Waiting



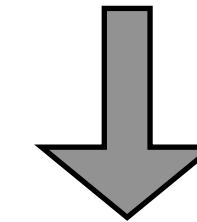
**wait\_for\_completion()**

## Thread-B (cryptd\_queue\_worker)

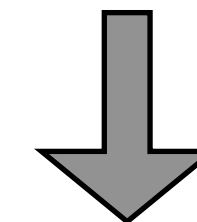
Dequeue request



Handle request

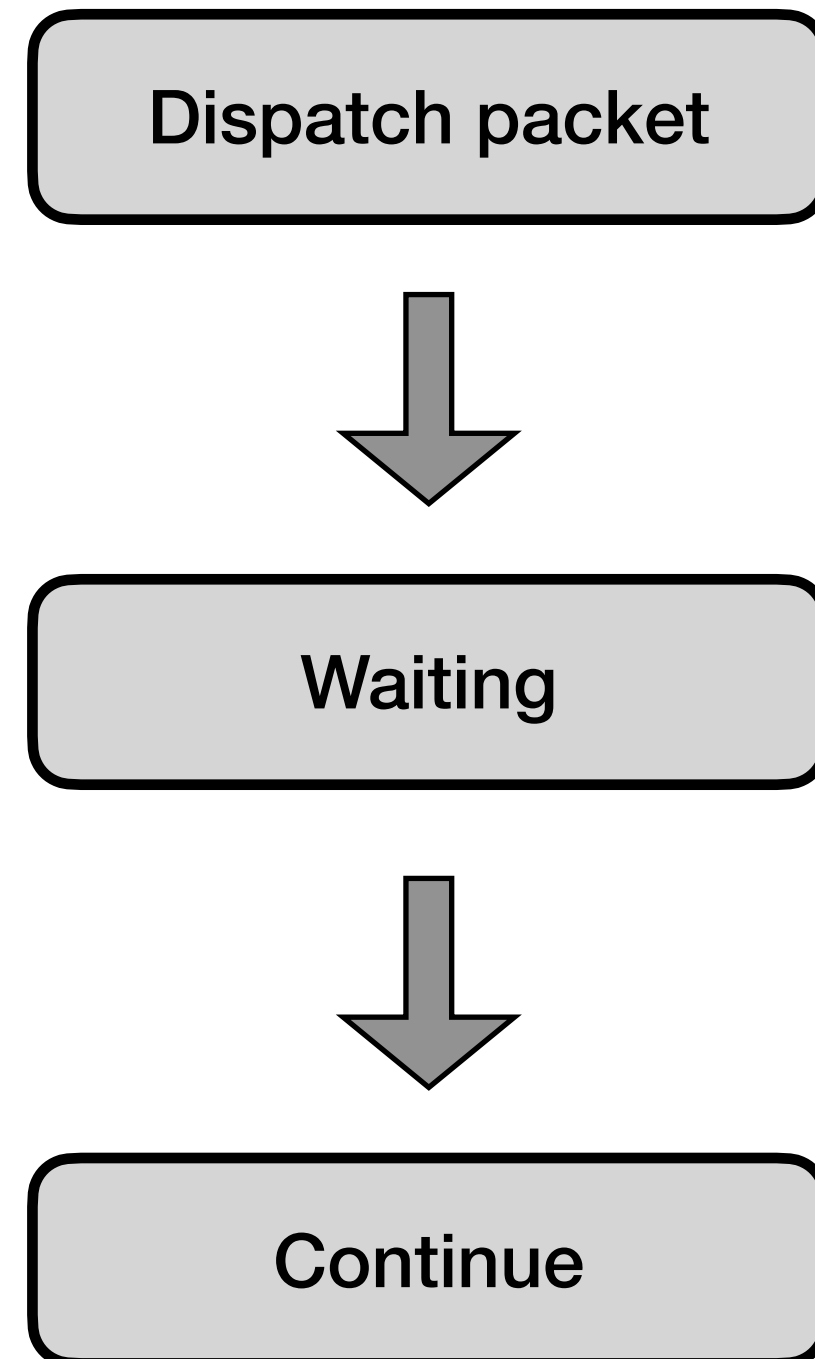


**complete()**

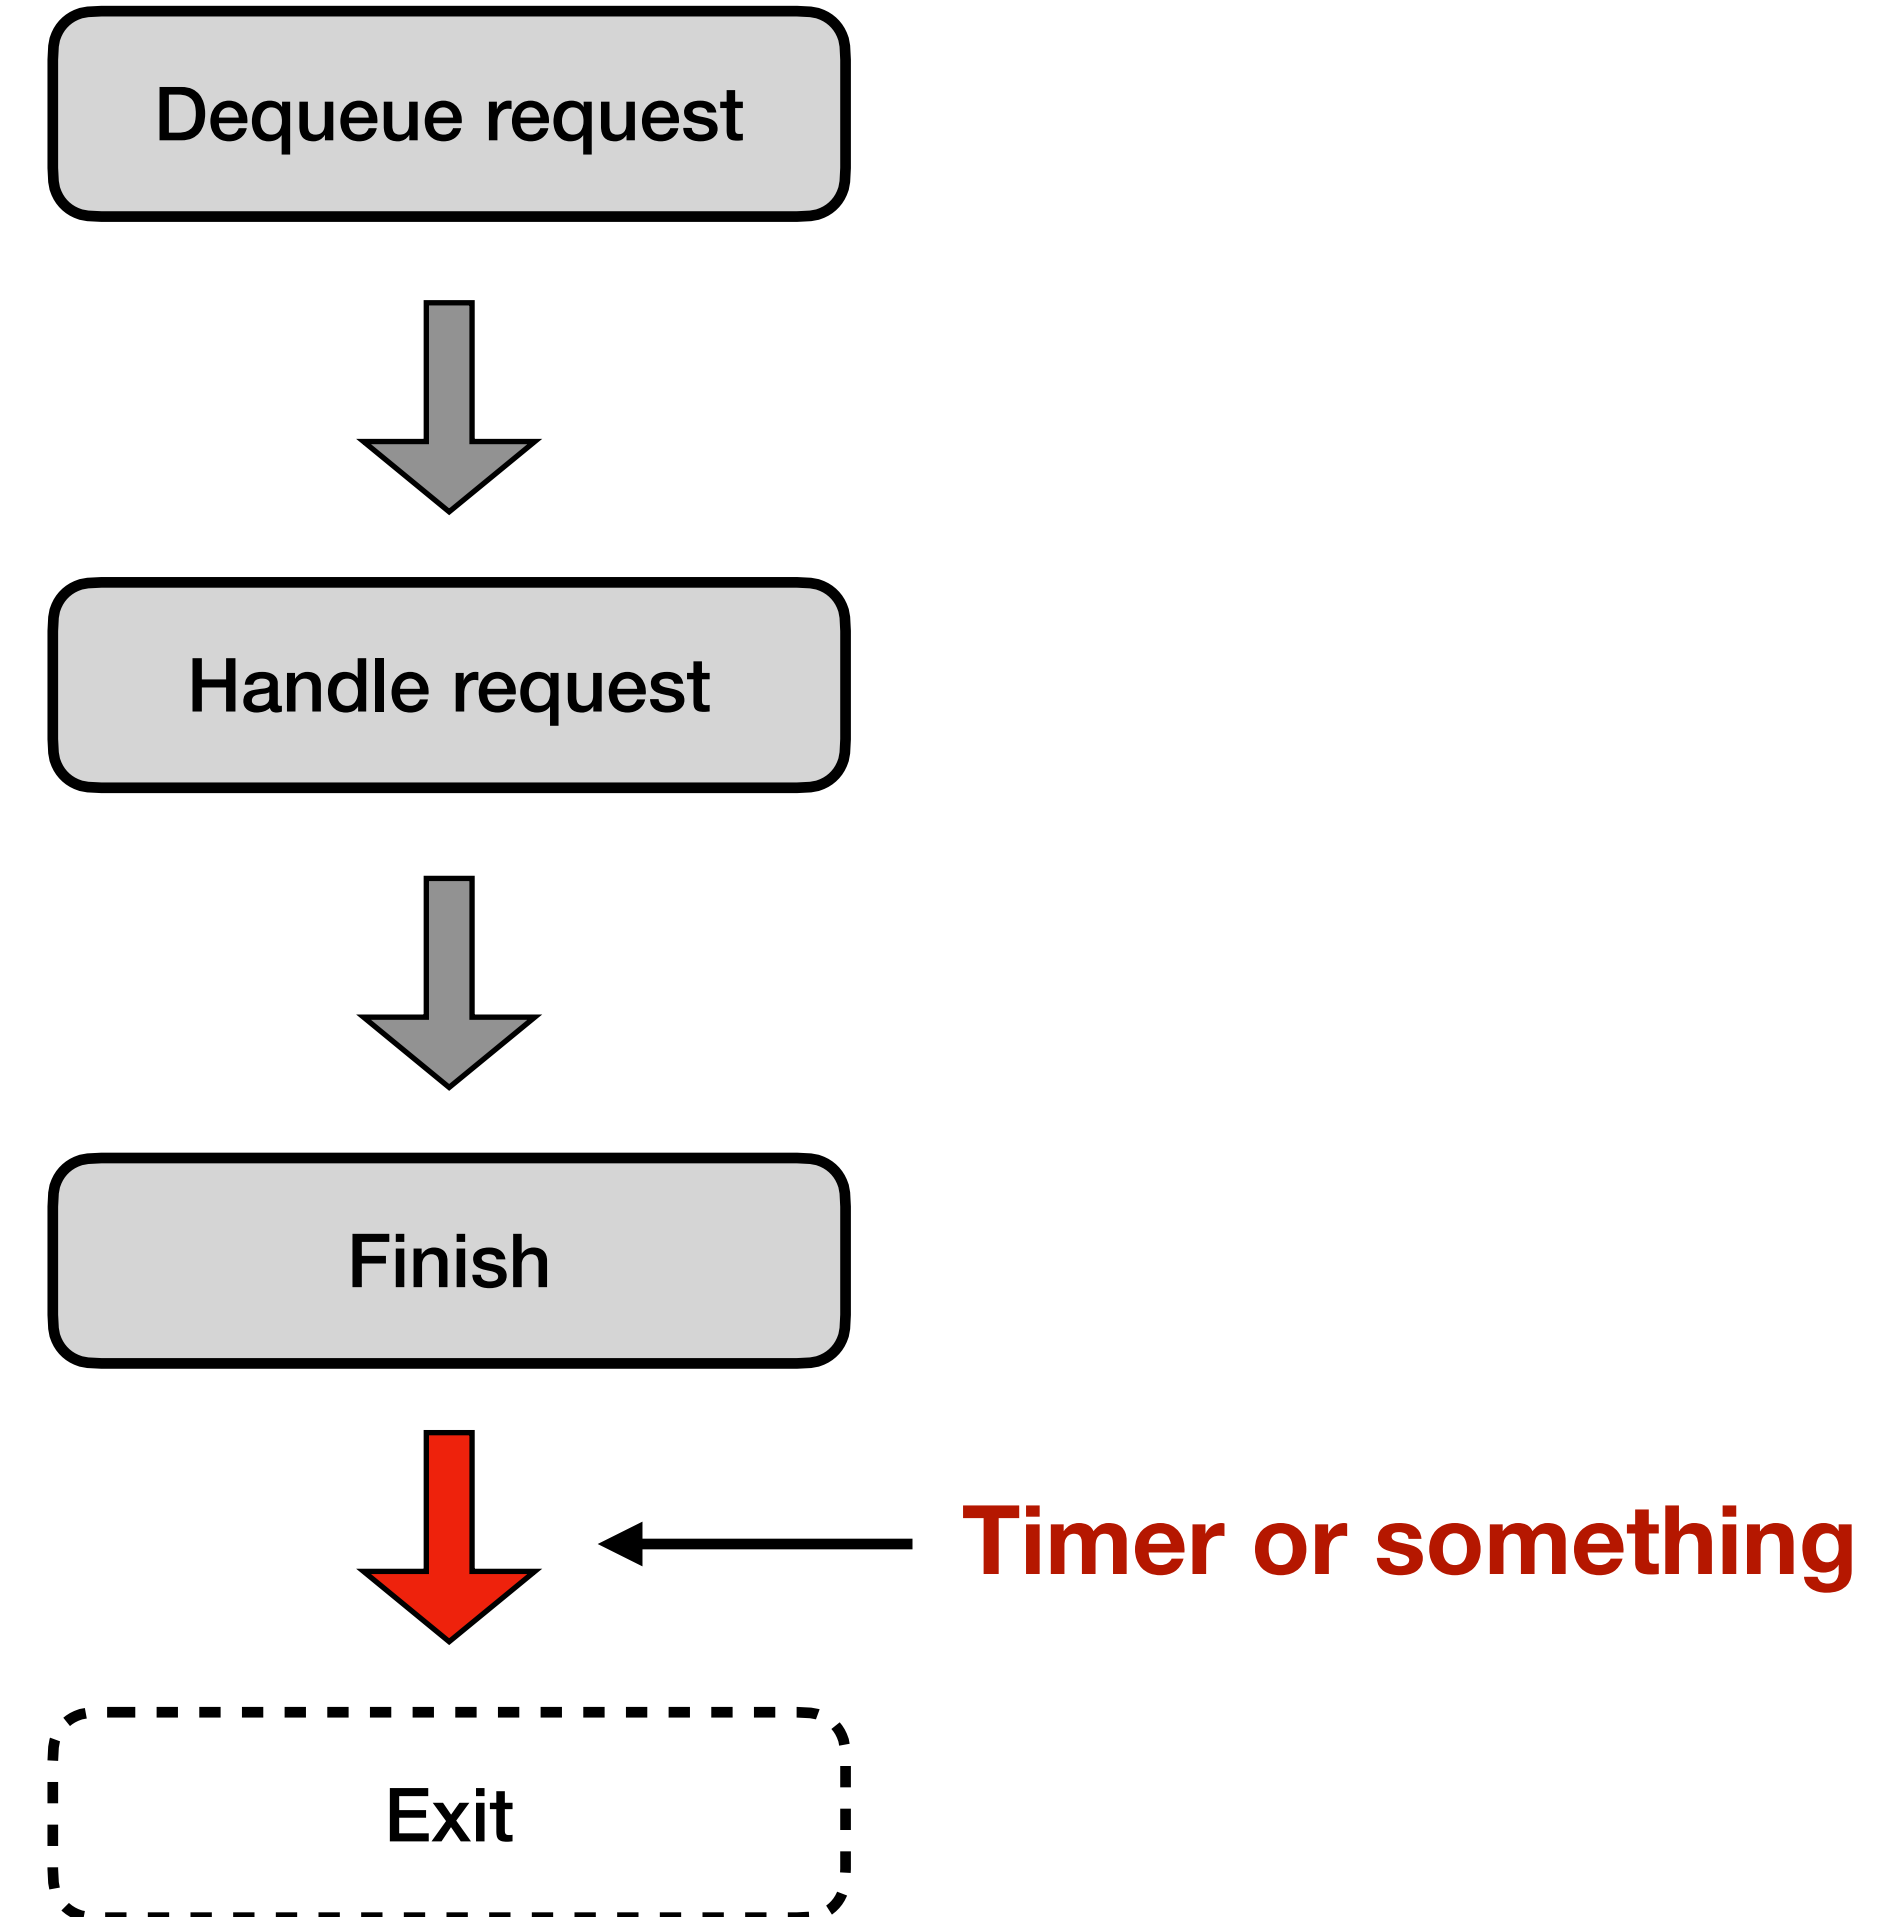


Exit

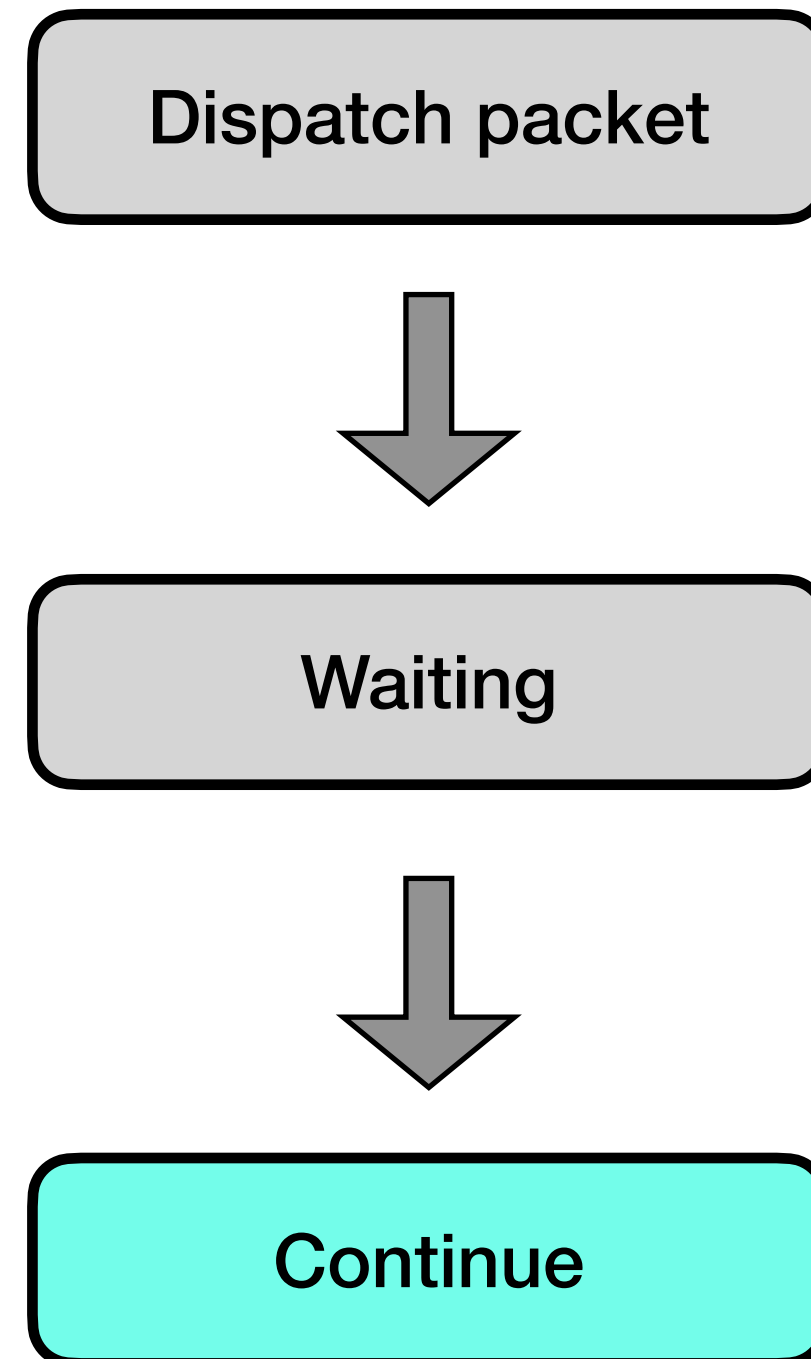
## Thread-A



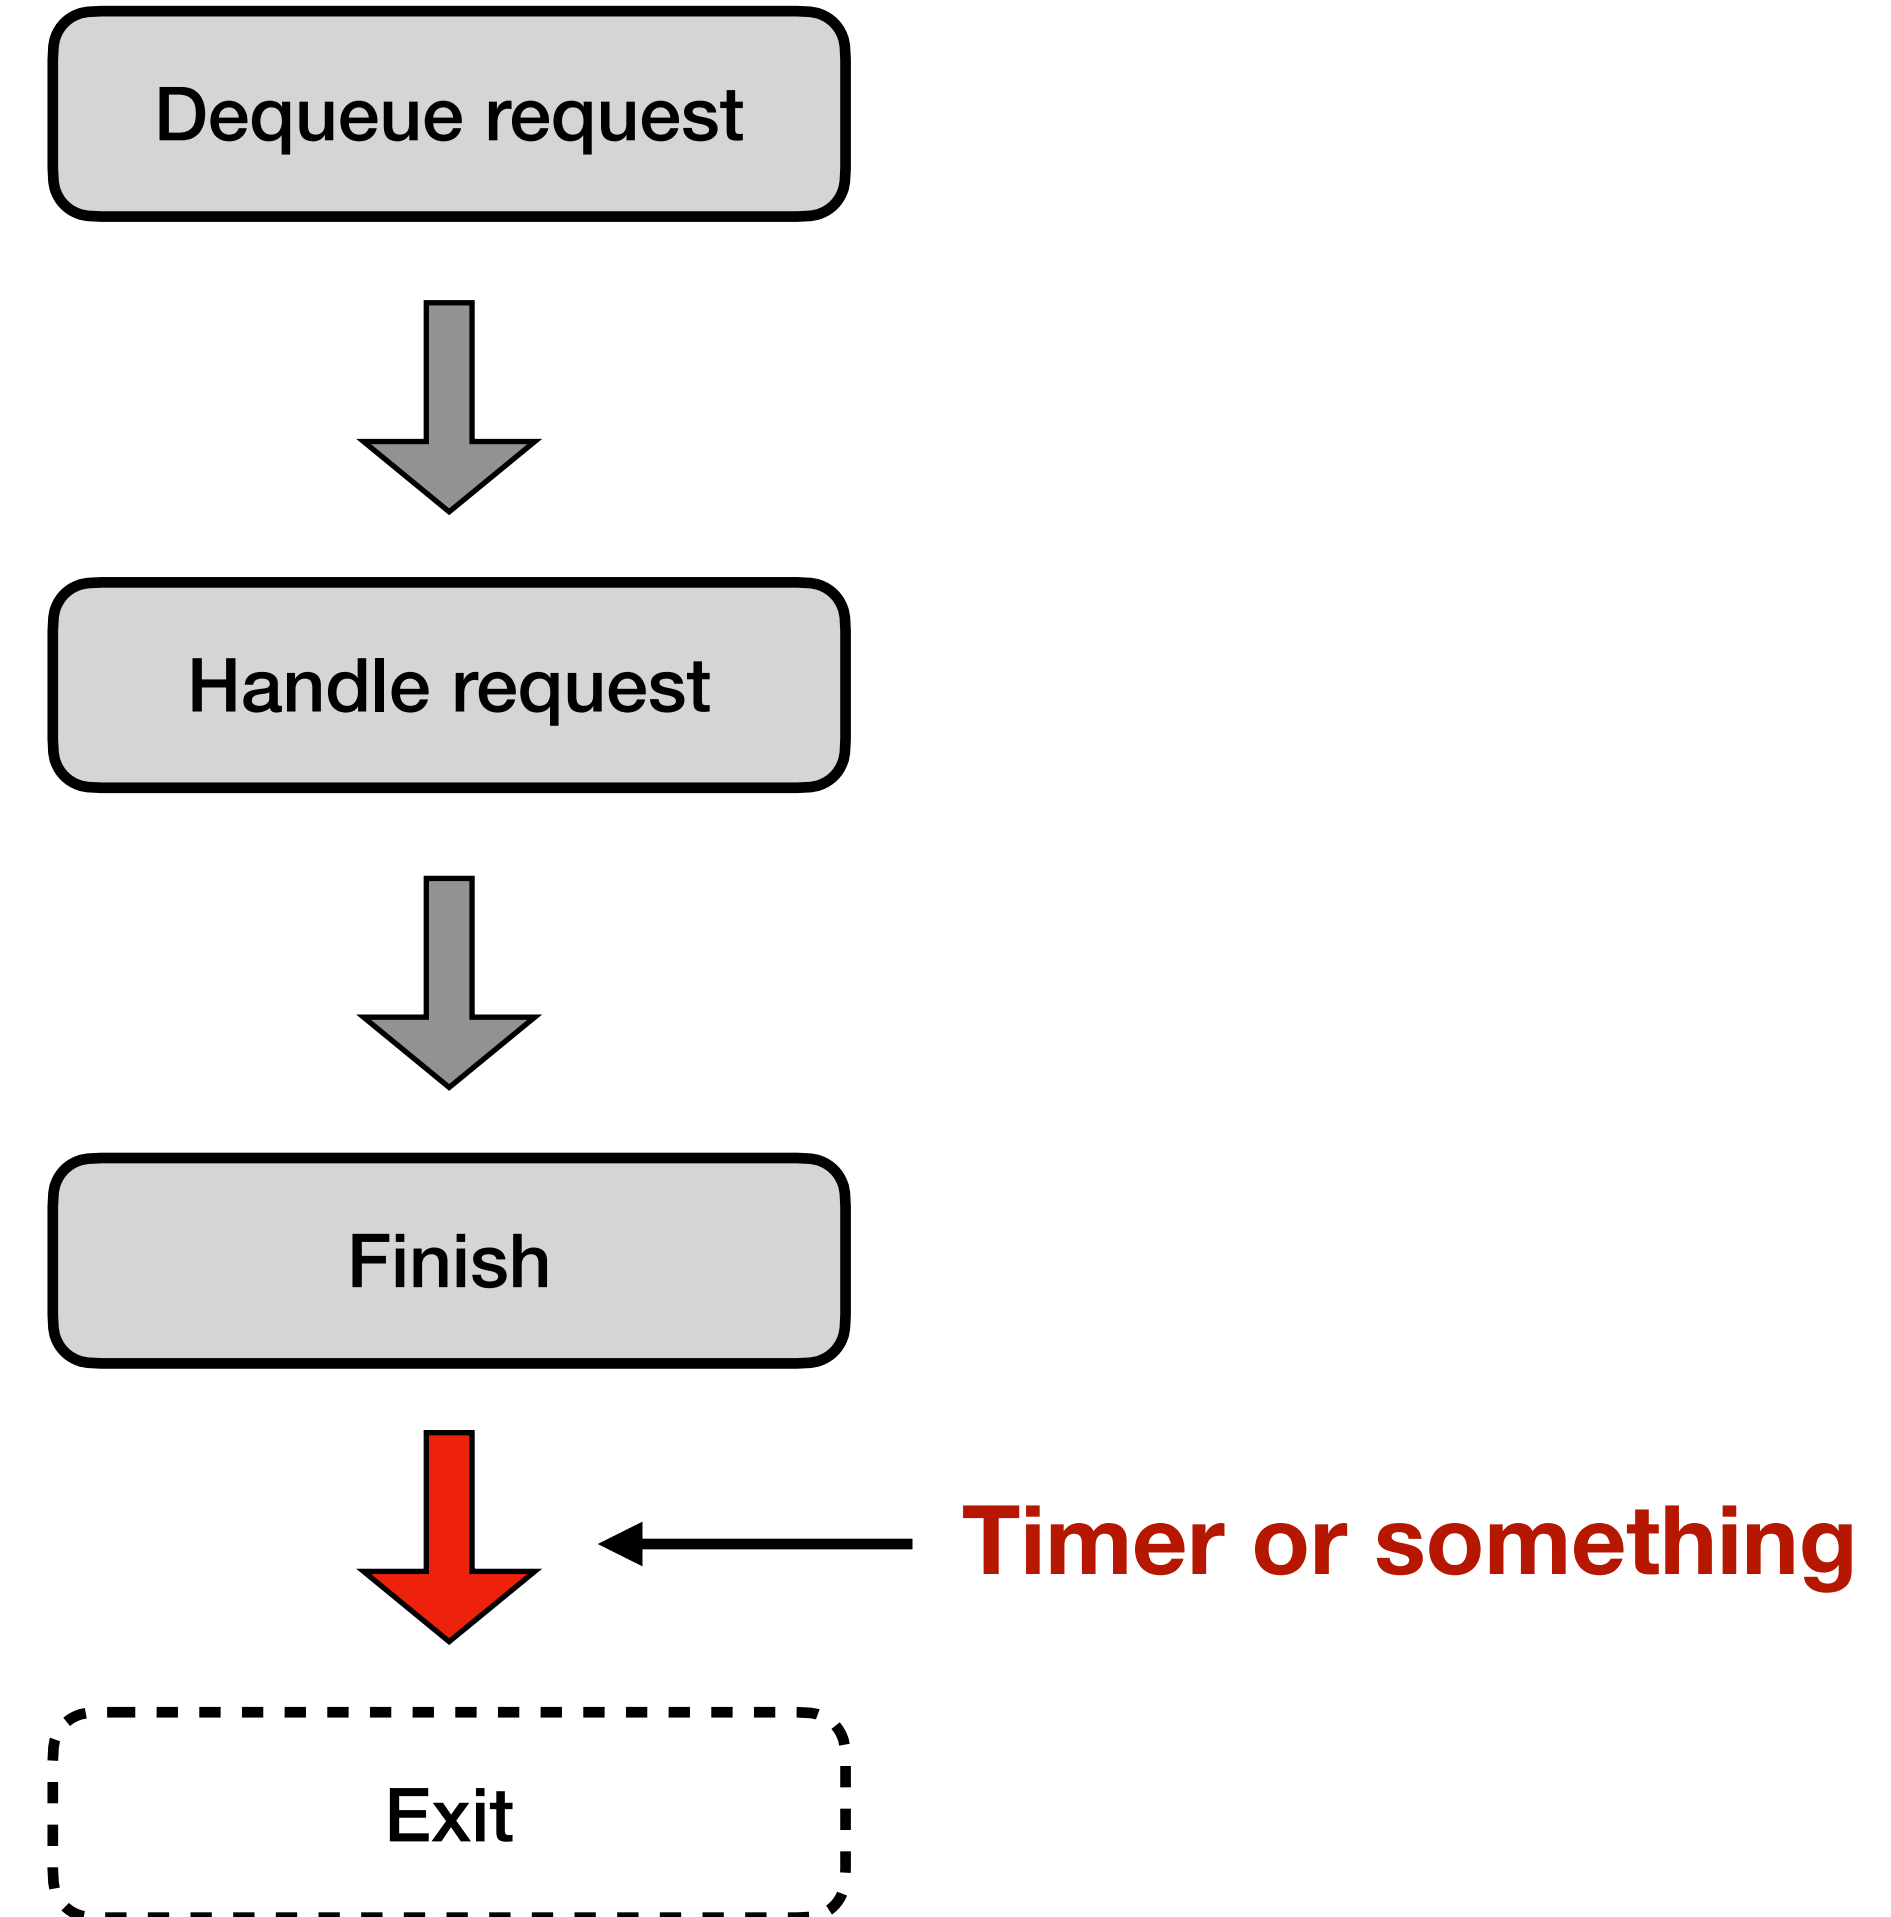
## Thread-B (cryptd\_queue\_worker)



## Thread-A

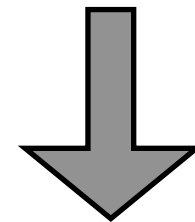


## Thread-B (cryptd\_queue\_worker)

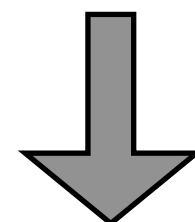


## Thread-A

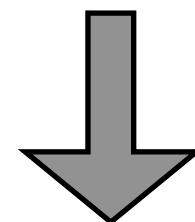
Dispatch packet



Waiting



Continue

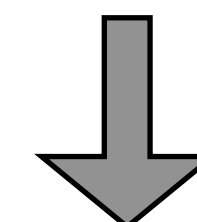


Exit

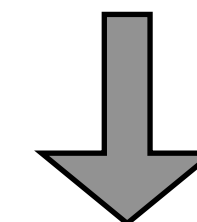
**Free TX/RX context ...**

## Thread-B (cryptd\_queue\_worker)

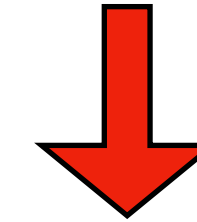
Dequeue request



Handle request



Finish

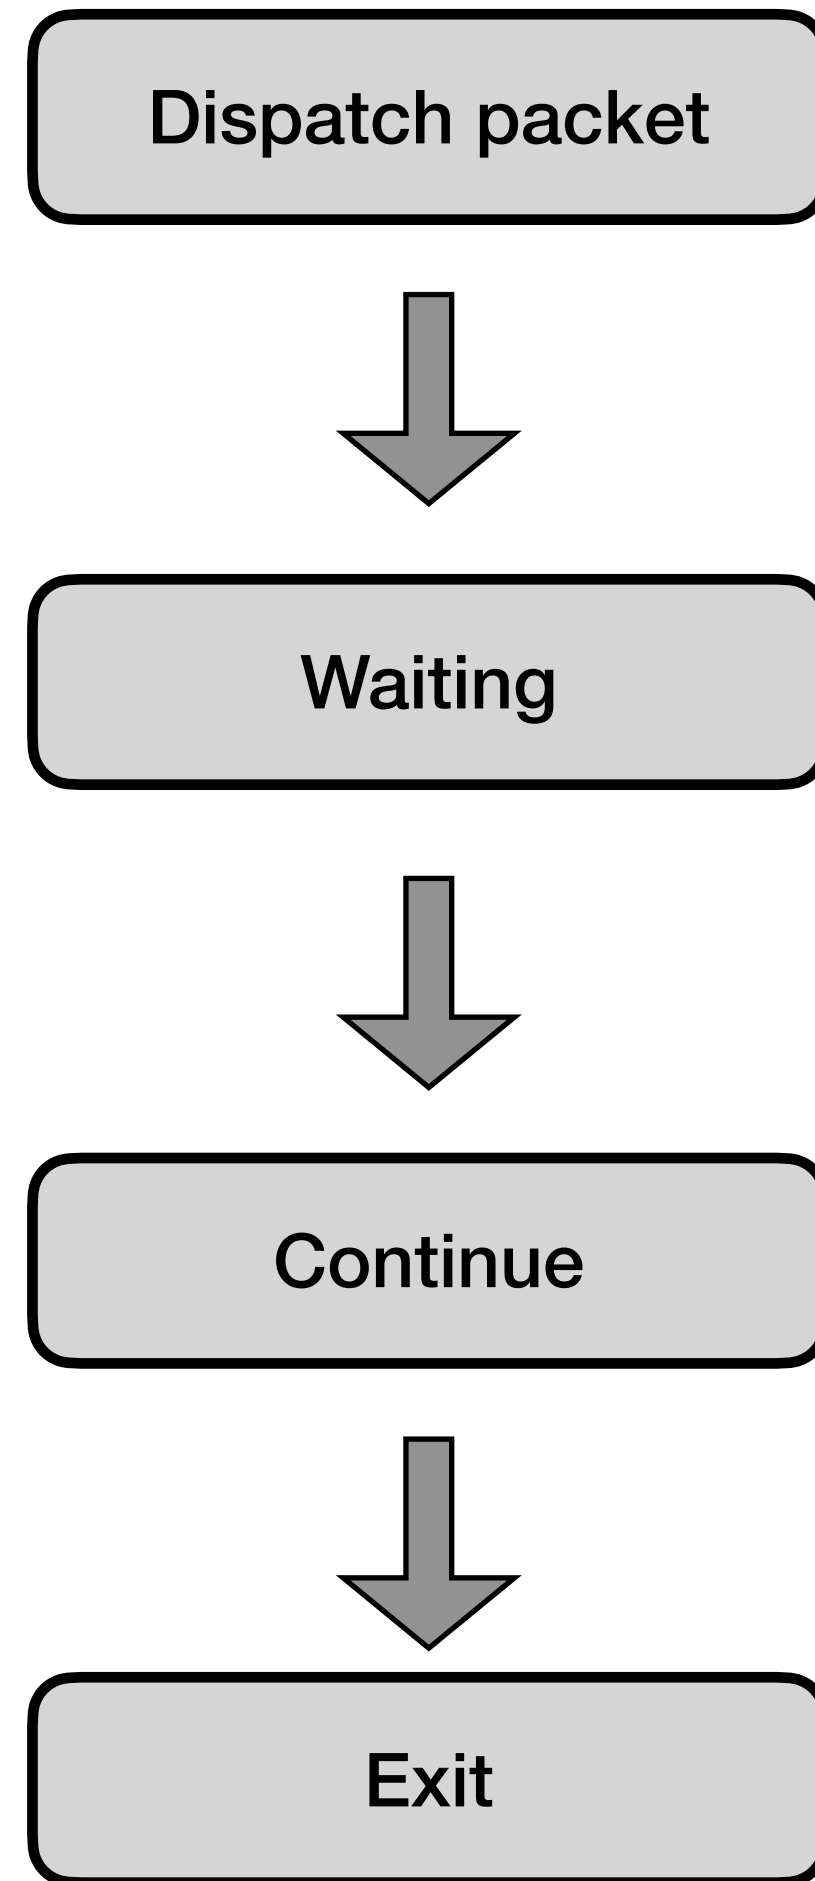


Exit

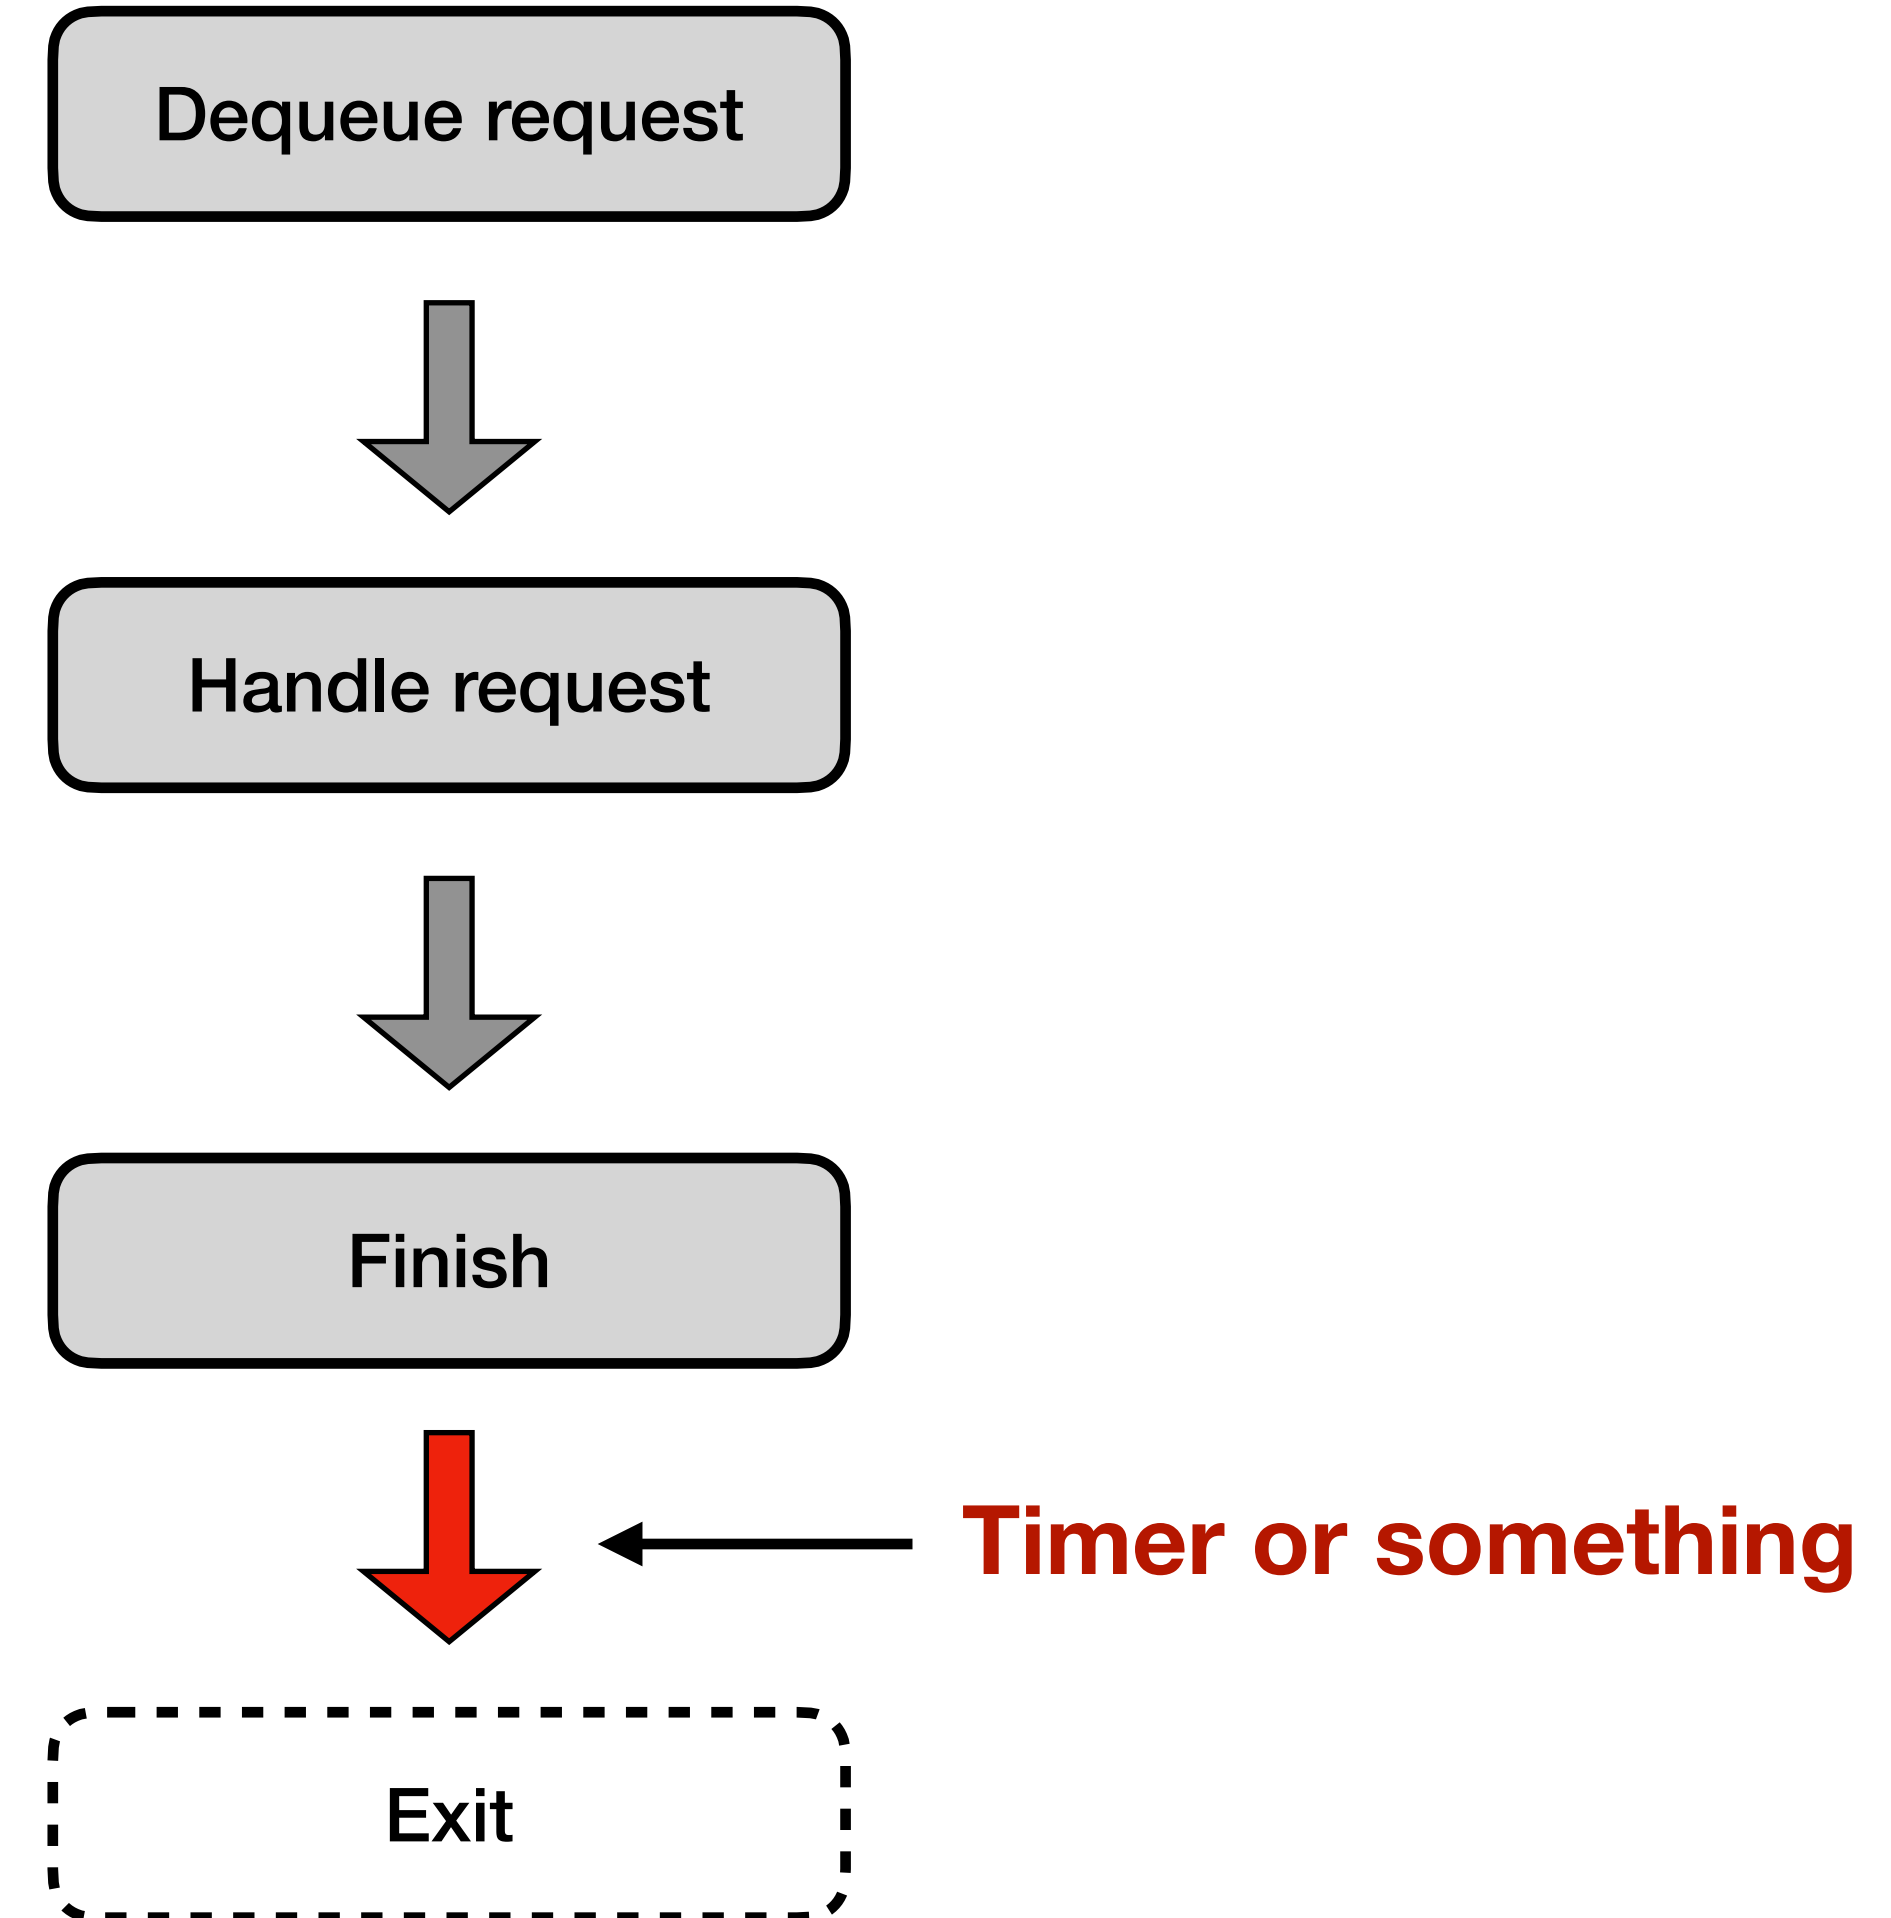
**← Timer or something**



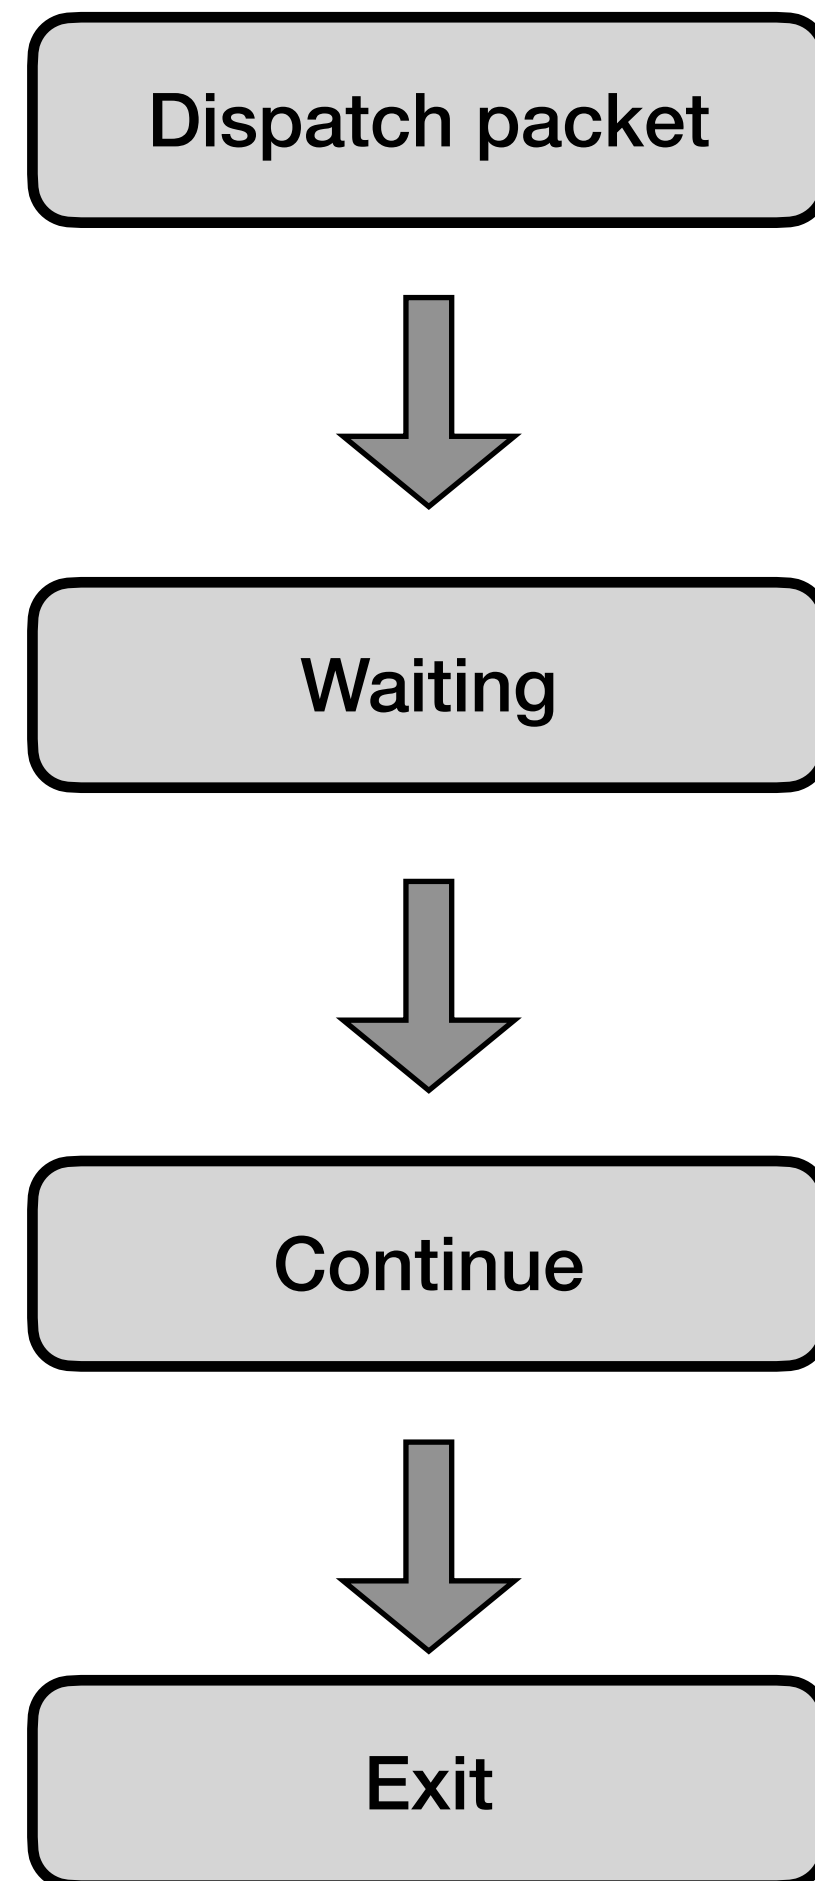
## Thread-A



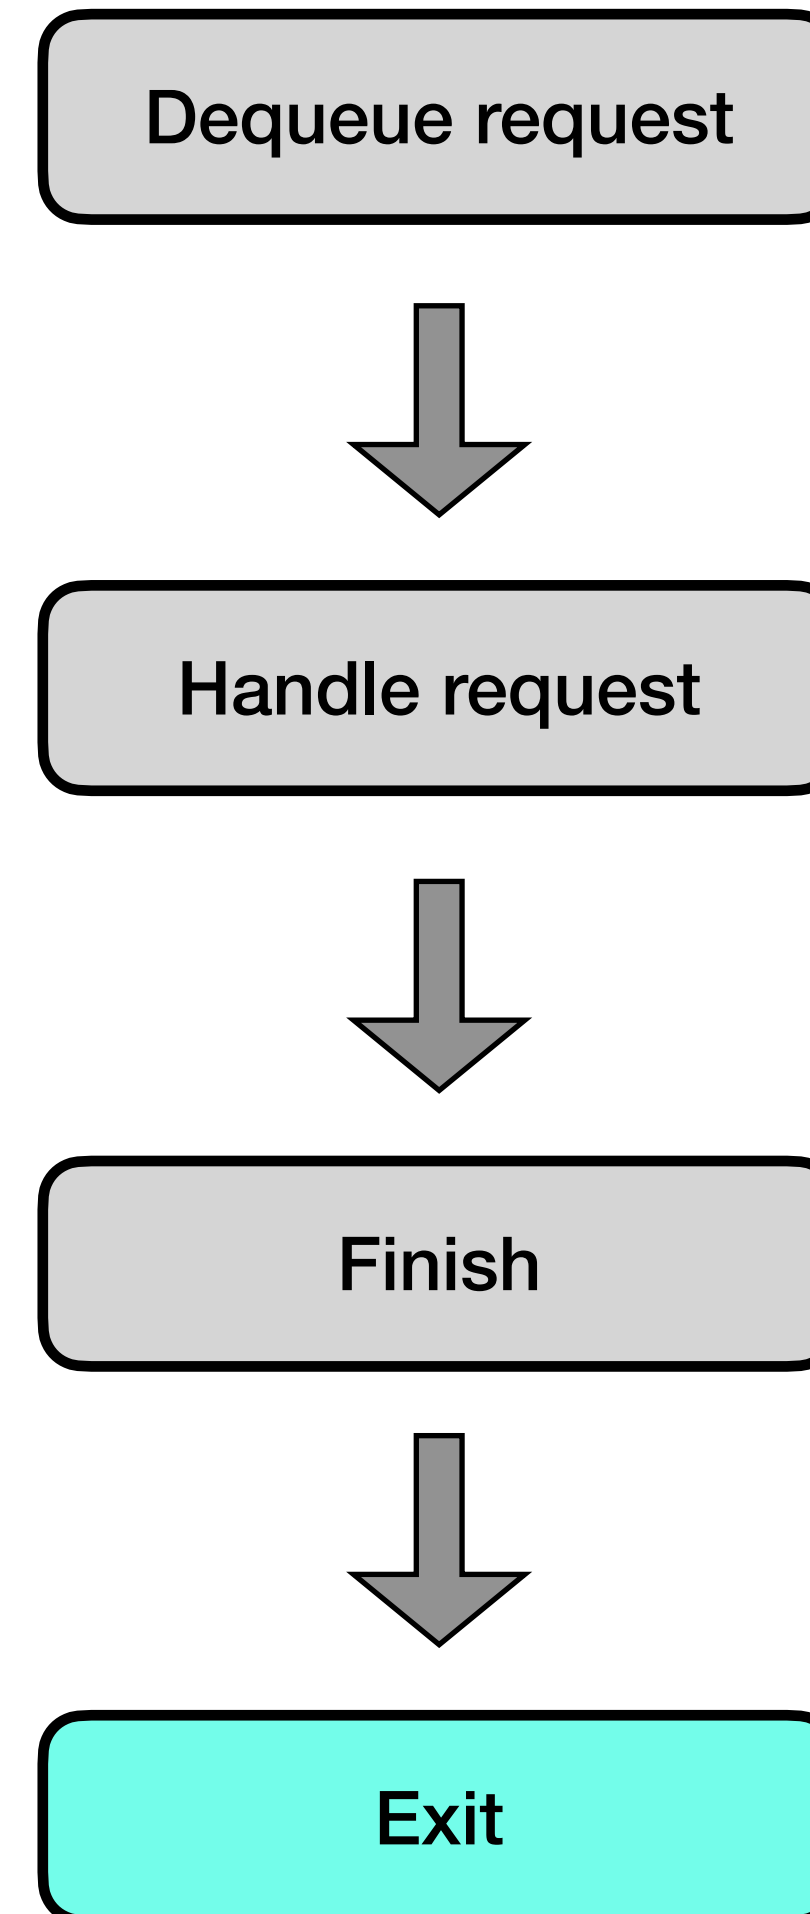
## Thread-B (cryptd\_queue\_worker)



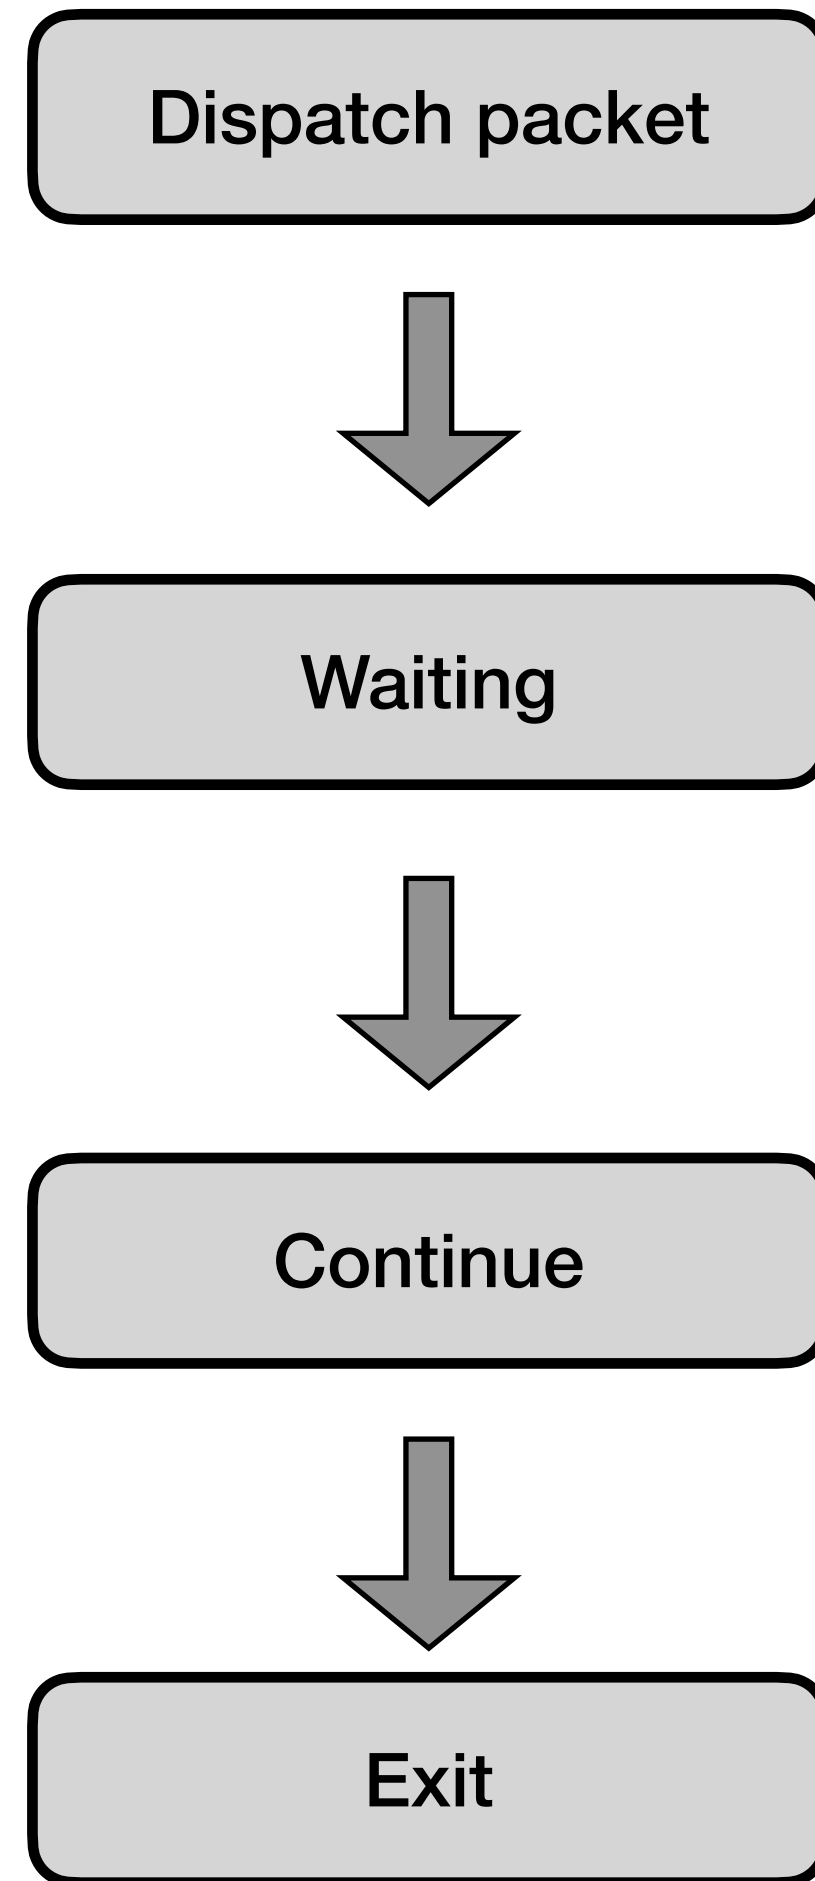
## Thread-A



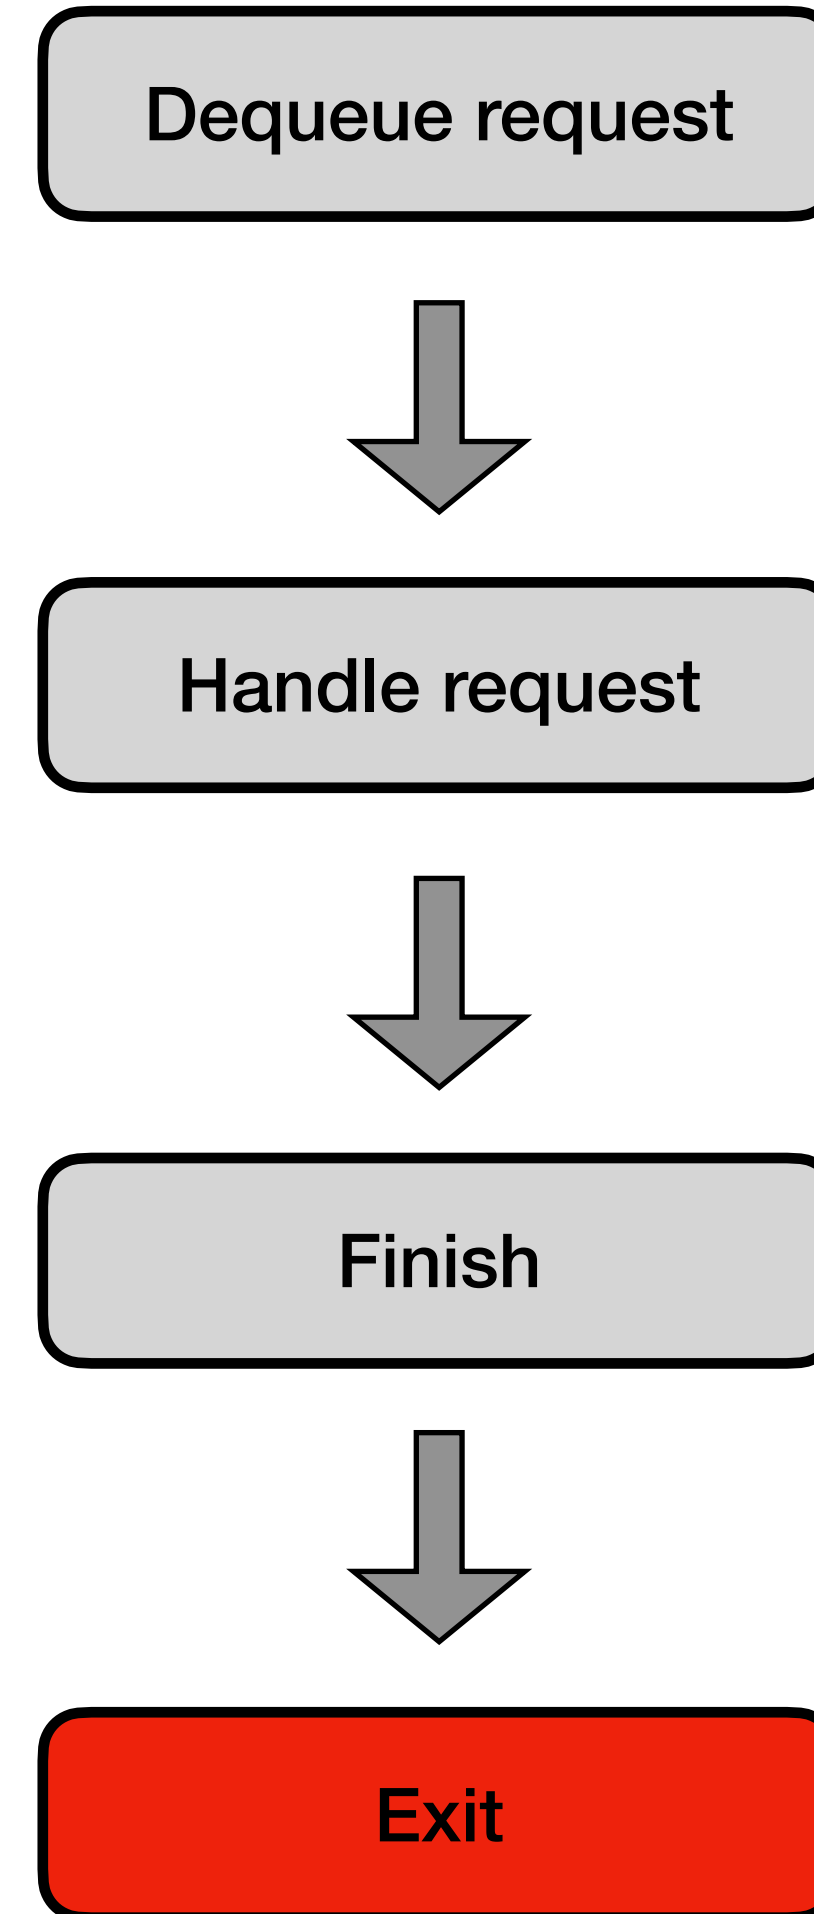
## Thread-B (cryptd\_queue\_worker)



## Thread-A



## Thread-B (cryptd\_queue\_worker)



**UAF when accessing TX/RX context object**