

Déjà Vu in Linux io_uring: Breaking Memory Sharing Again After Generations of Fixes

Pumpkin Chang (@u1f383)

October 10, 2025

DEV✓CORE

/HEXAICON/ 2025

\$ whoami

- Pumpkin 🎃 (@u1f383)
- Security researcher at DEVCORE
- Focus on Linux Kernel, Android and VM



\$ Outline

- Linux io_uring Introduction
 - Architecture & Memory Sharing
- Review of Three Previous Bugs
 - Memory Sharing Design Issues in io_uring
- CVE-2025-21836: Root Cause & Exploitation
- Takeaways

Introduction

\$ Overview

- io_uring
 - Introduced in Linux 5.1 in 2019
 - Provides high-performance async I/O
 - Reduces syscall overhead by asynchronous handling
 - **Shares objects** between user and kernel space for fast data exchange

\$ Overview

- Three core components:
 - I/O Command
 - Encapsulates each I/O request as a **command**
 - Prepares and validates it based on command type
 - Command Worker
 - Processes I/O requests in the **kernel I/O thread**
 - Memory Sharing
 - **Shares memory with user space** to avoid costly data copying

\$ Overview

- Three core components:
 - I/O Command
 - Encapsulates each I/O request as a **command**
 - Prepares and validates it based on command type
 - Command Worker
 - Processes I/O requests in the **kernel I/O thread**
 - Memory Sharing
 - Shares memory with user space to avoid costly data copying

\$ Overview

- Three core components:
 - I/O Command
 - Encapsulates each I/O request as a command
 - Prepares and validates it based on command type
 - Command Worker
 - Processes I/O requests in the kernel I/O thread
 - Memory Sharing
 - Shares memory with user space to avoid costly data copying

\$ Memory Sharing

- io_uring memory sharing:
 1. Submission Queue (SQ) and Completion Queue (CQ)
 - Used for submitting **commands** and fetching **results**
 2. Kernel buffers
 - Reallocated memory for zero-copy data transfer

\$ Memory Sharing

- io_uring memory sharing:
 1. Submission Queue (SQ) and Completion Queue (CQ)
 - Used for submitting commands and fetching results
 2. Kernel buffers
 - Store user data
 - **Reallocated** memory for zero-copy data transfer

\$ Memory Sharing

- io_uring memory sharing:
 1. Submission Queue (SQ) and Completion Queue (CQ)
 - Used for submitting commands and fetching results
 2. Kernel buffers
 - Store user data
 - **Reallocated** memory for zero-copy data transfer

What we focus on !

\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - Abstracts shared memory management

\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - Abstracts shared memory management
 - Process can **register** shared memory to io_uring context

\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - Abstracts shared memory management
 - Process can register shared memory to io_uring context
- Managed by I/O buffer object
 - Stored in an array within the context

\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - Abstracts shared memory management
 - Process can register shared memory to `io_uring` context
 - Managed by I/O buffer object
 - Stored in an array within the context
- Command can specify the **buffer index** to use it

User space

Process

Kernel space

1. Create an array to manage I/O buffer object

I/O buffer array

Index	I/O Buffer
0	
1	
...	

RAM

User space


Process

Kernel space

2. Assume there is an I/O buffer object with some backing pages

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process

READ

Flags: REQ_F_BUFFER_SELECT
Buffer index: 0

io_uring worker

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	

3. Process specifies the buffer index in I/O command

Page

Page

...

User space

Process

Kernel space

io_uring worker

READ

Index: 0

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	

4. The read data will be stored in I/O buffer object



\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - **Ring buffer**: a large, continuous memory region
 - [1] Memory can be **pre-allocated** by process
 - [2] or reserved by kernel
 - Provided buffer: multiple small, non-contiguous buffers
 - Use fragmented user-space memory to store data

User space

Kernel space

Process

1. Preallocate some anonymous pages

I/O buffer array


Index	I/O Buffer
0	
1	
...	




User space

Kernel space

Process



2. Register
- Flag: 0
- Index: 0



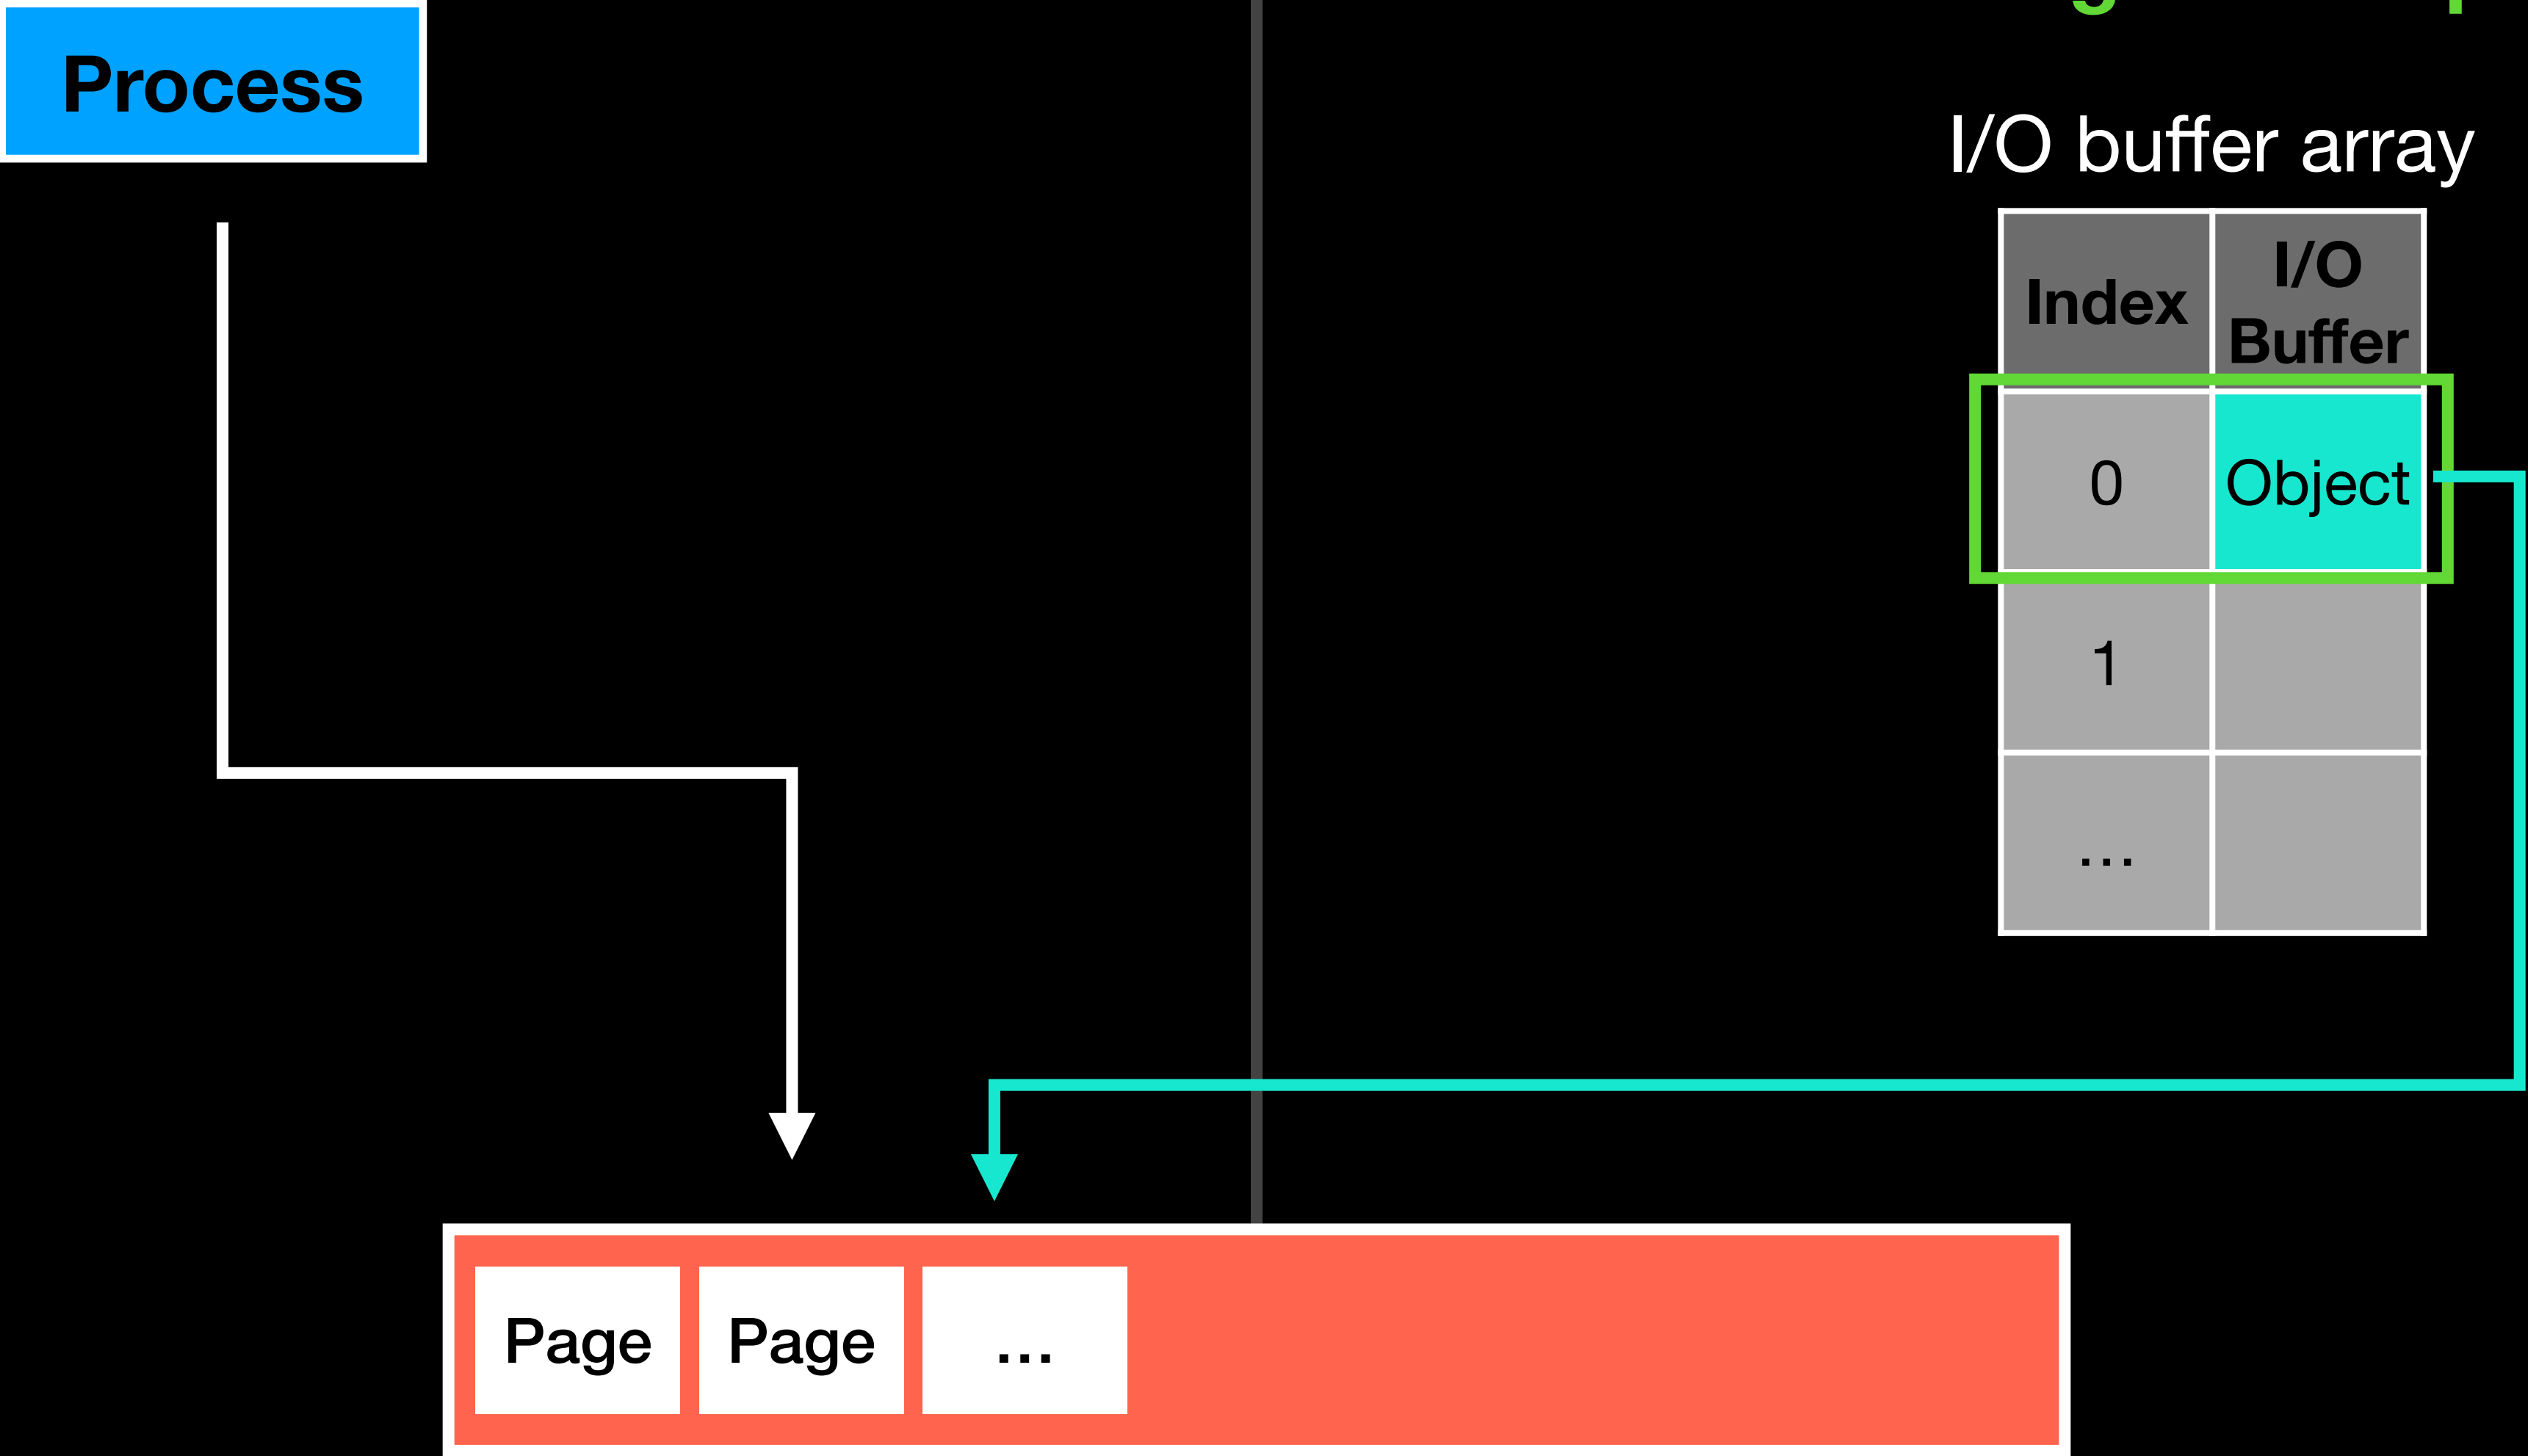
I/O buffer array

Index	I/O Buffer
0	
1	
...	

User space

Kernel space

3. Create an I/O buffer object to manage these pages



\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - **Ring buffer**: a large, continuous memory region
 - [1] Memory can be pre-allocated by process
 - [2] or **reserved** by kernel
 - Provided buffer: multiple small, non-contiguous buffers
 - Use fragmented user-space memory to store data

User space

Kernel space

Process



- 1. Register
 - Flag: IOU_PBUF_RING_MMAP
 - Index: 0

I/O buffer array

Index	I/O Buffer
0	
1	
...	



User space


Process

Kernel space

2. Allocate pages based on the request size

I/O buffer array


Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space


Process



3. Map the backing pages to user space

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	

4. The memory can now be accessed



\$ Memory Sharing

- I/O Buffer (Kernel buffer)
 - Ring buffer: a large, continuous memory region
 - [1] Memory can be pre-allocated by process
[2] or reserved by kernel
 - **Provided buffer**: multiple small, non-contiguous buffers
 - Use **fragmented** user-space memory to store data

User space

Kernel space

Process

PROVIDE_BUFFERS

Address: 0x7fff7c00000
nbuf: 3
len: 0x100
Buffer ID: 0

io_uring worker

I/O buffer array

Index	I/O Buffer
0	
1	
...	

1. Send PROVIDE_BUFFERS request

User space

Process

Kernel space

io_uring worker

PROV_BUF Buffer ID: 0

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	

2. Create an I/O buffer object

User space

Process

Kernel space

io_uring worker

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



Address: 0x7ffff7c00000
nbuf: 3
len: 0x100

3. The address, size and count are stored in the object

User space

Kernel space

Process

REMOVE_BUFFERS

nbuf: 1
Buffer ID: 0

io_uring worker

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	

Address: 0x7ffff7c00000
nbuf: 3
len: 0x100

4. Send REMOVE_BUFFERS request

User space

Process

Kernel space

io_uring worker

REMOV_BUF nbuf: 1
Buffer ID: 0

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	

5. Remove several sub-buffers

Address: 0x7fff7c00000
nbuf: 2
len: 0x100

\$ Memory Sharing

- Both types of buffer use the same `io_buffer_list` structure
- How to identify the buffer type?
 - Determined by flag fields: `is_mapped` and `is_mmap`
 - These fields are initialized during object creation

```
struct io_buffer_list {
```

```
/* If ->buf_nr_pages is set, then buf_pages/buf_ring are use
 * then these are classic provided buffers and ->buf_list is
 */
union {
    struct list_head buf_list;
    struct {
        struct page **buf_pages;
        struct io_uring_buf_ring *buf_ring;
    };
    struct rcu_head rcu;
};
__u16 bgid;

/* below is for ring provided buffers */
__u16 buf_nr_pages;
__u16 nr_entries;
__u16 head;
__u16 mask;

atomic_t refs;
```



is_mapped	is_mmap	Type
0	0	Provided buffer
0	1	SHOULD NOT HAPPEN
1	0	Ring buffer (with pre-allocated memory)
1	1	Ring buffer (with reserved memory)

```
/* ring mapped provided buffers */
__u8 is_mapped;
/* ring mapped provided buffers, but mmap'ed by application */
__u8 is_mmap;
```

The Evolution of Shared Memory

\$ RECON



Limiting io_uring

To protect our users, we decided to limit the usage of io_uring in Google products:

- **ChromeOS:** We [disabled](#) io_uring (while we explore new ways to sandbox it).
- **Android:** Our [seccomp-bpf filter](#) ensures that io_uring is unreachable to apps. Future Android releases will use SELinux to [limit io_uring access to a select few system processes](#).
- **GKE AutoPilot:** We are investigating disabling io_uring by default.
- It is disabled on production Google servers.

<https://security.googleblog.com/2023/06/learnings-from-kctf-vrps-42-linux.html>

\$ RECON

- Memory sharing
 - Easy to exploit
 - Often leads to **strong primitives**
- Hard to maintain
 - **State transitions** across alloc/update/release make bugs more likely

\$ Bug 1 - CVE-2024-0582

- io_uring/kbuf: Defer release of mapped buffer rings
 - Kernel reserves pages for the ring buffer
 - Mapped via `remap_pfn_range()` without incrementing refcount

\$ Bug 1 - CVE-2024-0582

- io_uring/kbuf: Defer release of mapped buffer rings
 - Kernel reserves pages for the ring buffer
 - Mapped via `remap_pfn_range()` without incrementing refcount
 - **Unregister** an I/O buffer will also release its pages


\$ Bug 1 - CVE-2024-0582

- io_uring/kbuf: Defer release of mapped buffer rings
 - Kernel reserves pages for the ring buffer
 - Mapped via `remap_pfn_range()` without incrementing refcount
 - Unregister an I/O buffer will also release its pages
 - However, the user mapping **remains accessible**, resulting in a page UAF

User space

Kernel space


Process



- 1. Register
 - Flag: IOU_PBUF_RING_MMAP
 - Index: 0

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space


Kernel space

Process

2. Map the these pages

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process

3. Unregister
- Index: 0

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process

4. The I/O buffer object is freed along with its backing pages

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process

5. These page are still accessible from user space

I/O buffer array

Index	I/O Buffer
0	
1	
...	



\$ Bug 1 - CVE-2024-0582

- How to fix?
 - Prevent unregistration of mapped ring buffer

\$ Bug 1 - CVE-2024-0582

- How to fix?
 - Prevent unregistration of mapped ring buffer
 - Update refcount of page or I/O buffer object via VMA op

\$ Bug 1 - CVE-2024-0582

- How to fix?
 - Prevent unregistration of mapped ring buffer
 - Update refcount of page or I/O buffer object via VMA op
 - Add counter field in object to track mappings

\$ Bug 1 - CVE-2024-0582

- How to fix?

- ❌ Prevent unregistration of mapped ring buffer

- ❌ Update refcount of page or I/O buffer object via VMA op

- ❌ Add counter field in object to track mappings

- ✅ Defer releasing these pages until io_uring context is closed



**Manage object
lifecycles pro-
perly**



**Introduce a new
feature as work-
around**

User space

Kernel space

Process

1. Add a new field

io_buf_list

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process



2. Unregistration frees I/O buffer, but pages are still alive

io_buf_list

I/O buffer array

Index	I/O Buffer
0	Object
1	
...	



User space

Kernel space

Process

io_buf_list

I/O buffer array

Index	I/O Buffer
0	
1	
...	

3. Safely access!



\$ Bug 1 - CVE-2024-0582

- Behind the patch
 - Introduce a new feature as a workaround
 - More code means more potential issues

\$ Bug 1 - CVE-2024-0582

- Behind the patch
 - Introduce a new feature as a workaround
 - More code means more potential issues
- Not using **reference counting** that precisely
 - Should hold a refcount either page or object

\$ Bug 1 - CVE-2024-0582

- Behind the patch
 - Introduce a new feature as a workaround
 - More code means more potential issues
 - Not using reference counting that precisely
 - Should hold a refcount either page or object
- Mixes functionality with **lifetime** management

\$ Bug 2 - Lockdep

- io_uring: free io_buffer_list entries via RCU
 - Uncovered when addressing CVE-2024-0582
 - **Deadlock** between the io_uring context lock and the memory management (mm) lock

Process 1

Register an I/O buffer

```
def register(uaddr):
    mutex_lock(context_lock)

    if is_not_present(uaddr):
        lock(mm)
        handle_page_fault(uaddr)
        unlock(mm)

    copy_from_user(uaddr, &req)
    obj = new_io_buffer_obj(&req)
    io_buffer_obj_arr[req.idx] = obj

    mutex_unlock(context_lock)
```

Process 2

Map an I/O buffer to user space

```
def mmap_v1(idx):
    lock(mm)
    mutex_lock(context_lock)

    obj = io_buffer_obj_arr[idx]
    if can_mmap(obj):
        handle_memory_mapping(obj)

    mutex_unlock(context_lock)
    unlock(mm)
```

Process 1

Register an I/O buffer

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)  
        unlock(mm)  
  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj  
  
    mutex_unlock(context_lock)
```

Process 2

Map an I/O buffer to user space

```
def mmap_v1(idx):  
    lock(mm)  
    mutex_lock(context_lock)  
  
    obj = io_buffer_obj_arr[idx]  
    if can_mmap(obj):  
        handle_memory_mapping(obj)  
  
    mutex_unlock(context_lock)  
    unlock(mm)
```

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
    lock(mm)  
    handle_page_fault(uaddr)  
    unlock(mm)  
  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj  
  
    mutex_unlock(context_lock)
```

Hold context lock

Process 2

```
def mmap_v1(idx):  
    lock(mm)  
    mutex_lock(context_lock)  
  
    obj = io_buffer_obj_arr[idx]  
    if can_mmap(obj):  
        handle_memory_mapping(obj)  
  
    mutex_unlock(context_lock)  
    unlock(mm)
```

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)  
        unlock(mm)  
  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj  
  
    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v1(idx):  
    lock(mm)  
  
    mutex_lock(context_lock)  
  
    obj = io_buffer_obj_arr[idx]  
    if can_mmap(obj):  
        handle_memory_mapping(obj)  
  
    mutex_unlock(context_lock)  
    unlock(mm)
```

Hold **mm** lock

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)  
  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj  
  
    mutex_unlock(context_lock)
```

Wait for **mm** lock

Process 2

```
def mmap_v1(idx):  
    lock(mm)  
    mutex_lock(context_lock)  
  
    obj = io_buffer_obj_arr[idx]  
    if can_mmap(obj):  
        handle_memory_mapping(obj)  
  
    mutex_unlock(context_lock)  
    unlock(mm)
```


Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)  
  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj  
  
    mutex_unlock(context_lock)
```

Wait for **mm** lock

Process 2

```
def mmap_v1(idx):  
    lock(mm)  
    mutex_lock(context_lock)  
  
    if can_mmap(obj):  
        handle_memory_mapping(obj)  
  
    mutex_unlock(context_lock)  
    unlock(mm)
```

Wait for **context** lock

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)  
  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj
```

Wait for **mm** lock

Process 2

```
def mmap_v1(idx):  
    lock(mm)  
    mutex_lock(context_lock)  
    [obj, req] = mmap_v1_obj[idx]  
    if can_mmap(obj):  
        handle_memory_mapping(obj)  
    mutex_unlock(context_lock)
```

Wait for **context** lock

DEADLOCK

\$ Bug 2 - Lockdep

- The key patch:
 - Replace the **mutex** lock with an **RCU lock**

```
@@ -3498,9 +3498,9 @@ static void *io_uring_validate_mmap_request(struct file *file,
                        unsigned int bgid;


                        bgid = (offset & ~IORING_OFF_MMAP_MASK) >> IORING_OFF_PBUF_SHIFT;
-                       mutex_lock(&ctx->uring_lock);
+                       rcu_read_lock();
                        ptr = io_pbuf_get_address(ctx, bgid);
-                       mutex_unlock(&ctx->uring_lock);
+                       rcu_read_unlock();
                        if (!ptr)
                                return ERR_PTR(-EINVAL);
                        break;
```

\$ Bug 2 - Lockdep

- The key patch:
- Replace the **mutex** lock with an **RCU lock**

```
@@ -3498,9 +3498,9 @@ static void *io_uring_validate_mmap_rec  
- mutex_lock(&ctx->uring_lock);  
+ rcu_read_lock();  
ptr = io_pbuf_get_address(ctx, bgid);  
- mutex_unlock(&ctx->uring_lock);  
+ rcu_read_unlock();
```

```
+ rcu_read_unlock();  
if (!ptr)  
    return ERR_PTR(-EINVAL);  
break;
```



\$ Bug 2 - Lockdep

- RCU (Read-Copy-Update)
 - **Lock-free** synchronization for read-mostly data
 - Old data is freed **after** readers finish
 - Writer API: `synchronize_rcu()`, `call_rcu()`, `kfree_rcu()`
 - Reader API: `rcu_read_lock()` & `rcu_read_unlock()`

\$ Bug 2 - Lockdep

- Mutex Lock
 - Ensures **only one** thread accesses a shared resource at a time
 - Lead to a deadlock :(
- RCU Lock
 - Lock-free, while also protecting reads to the I/O buffer
 - However, it allows concurrent access!

\$ Bug 2 - Lockdep

- Mutex Lock
 - Ensures only one thread accesses a shared resource at a time
 - Lead to a deadlock :(
- RCU Lock
 - Lock-free, while also protecting reads to the I/O buffer object
 - However, it allows **concurrent access!**

Process 1

```
def register(uaddr):
    mutex_lock(context_lock)

    if is_not_present(uaddr):
        lock(mm)
        handle_page_fault(uaddr)
        unlock(mm)

    copy_from_user(uaddr, &req)
    obj = new_io_buffer_obj(&req)
    io_buffer_obj_arr[req.idx] = obj
    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v2(idx):
    lock(mm)
    - mutex_lock(context_lock)
    + rcu_read_lock()

    obj = io_buffer_obj_arr[idx]
    if obj.is_mmap:
        addr = obj.buf_ring
        map_address_to_user_space(addr)

    - mutex_unlock(context_lock)
    + rcu_read_unlock()
```

The mutex lock is replaced with an **RCU lock**

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock) ✓  
  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)
```

So the register handler
is **not blocked** anymore


```
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj  
    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v2(idx):  
    lock(mm)  
    - mutex_lock(context_lock)  
    + rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
    if obj.is_mmap:  
        addr = obj.buf_ring  
        map_address_to_user_space(addr)  
  
    - mutex_unlock(context_lock)  
    + rcu_read_unlock()  
    unlock(mm)
```


Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
    if is_not_present(uaddr):  
        lock(mm)  
        handle_page_fault(uaddr)  
        unlock(mm)  
    copy_from_user(uaddr, &req)  
    obj = new_io_buffer_obj(&req)  
    io_buffer_obj_arr[req.idx] = obj
```



But It also means the registration handler may executes **concurrently**

Process 2

```
def mmap_v2(idx):  
    lock(mm)  
    - mutex_lock(context_lock)  
    + rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
    if obj.is_mmap:  
        addr = obj.buf_ring  
        map_address_to_user_space(addr)  
  
    - mutex_unlock(context_lock)  
    + rcu_read_unlock()  
    unlock(mm)
```

\$ Bug 3 - CVE-2024-35880

- io_uring/kbuf: hold io_buffer_list reference over mmap
 - After the previous patch, **concurrent access** to the I/O buffer is allowed
 - RCU prevents releasing I/O buffer objects, but not **updating** them

\$ Bug 3 - CVE-2024-35880

- **Unregistration** does not free I/O buffer object immediately:
 1. Clear `bl->is_mmap` flag
 2. Reset `bl->buf_list` with `INIT_LIST_HEAD()`
 3. Finally, call `kfree_rcu()`

\$ Bug 3 - CVE-2024-35880

- Unregistration does not free I/O buffer object immediately:
 1. Clear `bl->is_mmap` flag
 2. Reset `bl->buf_list` with `INIT_LIST_HEAD()`
 3. Finally, call `kfree_rcu()`
- Potential **race** between mmap handler and resource cleanup 😎

Process 1

Unregister an I/O buffer

```
def unregister(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    obj = io_buffer_obj_arr[req.idx]

    obj.refcount -= 1
    if obj.refcount == 0:
        obj.is_mmap = 0
        INIT_LIST_HEAD(&obj.buf_list)
        kfree_rcu(obj)

    mutex_unlock(context_lock)
```

Process 2

Map an I/O buffer to user space

```
def mmap_v2(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]
    if obj.is_mmap:
        addr = obj.buf_ring
        map_address_to_user_space(addr)

    rcu_read_unlock()
```

Process 1

Unregister an I/O buffer

```
def unregister(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    obj = io_buffer_obj_arr[req.idx]

    obj.refcount -= 1
    if obj.refcount == 0:
        obj.is_mmap = 0
        INIT_LIST_HEAD(&obj.buf_list)
        kfree_rcu(obj)

    mutex_unlock(context_lock)
```

Process 2

Map an I/O buffer to user space

```
def mmap_v2(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]
    if obj.is_mmap:
        addr = obj.buf_ring
        map_address_to_user_space(addr)

    rcu_read_unlock()
```

Process 1

```
def unregister(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    obj = io_buffer_obj_arr[req.idx]

    obj.refcount -= 1
    if obj.refcount == 0:
        obj.is_mmap = 0
        INIT_LIST_HEAD(&obj.buf_list)
        kfree_rcu(obj)

    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v2(idx):
    rcu_read_lock()


    obj = io_buffer_obj_arr[idx]
    if obj.is_mmap:
        addr = obj.buf_ring
        space(addr)

    rcu_read_unlock()
```

Pass the check


Process 1

```
def unregister(uaddr):  
    mutex_lock(context_lock)  
  
    copy_from_user(uaddr, &req)  
    obj = io_buffer_obj_arr[req.idx]  
  
    obj.refcount -= 1  
    if obj.refcount == 0:  
        obj.is_mmap = 0  
        INIT_LIST_HEAD(&obj.buf_list)  
        kfree(obj)  
    mutex_unlock(context_lock)
```

 **Reset the linked list buf_list**

Process 2

```
def mmap_v2(idx):  
    rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
    if obj.is_mmap:  
        addr = obj.buf_ring  
        map_address_to_user_space(addr)  
  
    rcu_read_unlock()
```




```
struct io_buffer_list {  
    /*  
     * If ->buf_nr_pages is set, then buf_pages/buf_ring are used. If not,  
     * then these are classic provided buffers and ->buf_list is used.  
     */  
    union {  
        struct list_head buf_list;  
        struct {  
            struct page **buf_pages;  
            struct io_uring_buf_ring *buf_ring;  
        };  
    };  
    __u16 ready;  
    __u16 mask;  
    atomic_t refs;  
    /* ring mapped provided buffers */  
    __u8 is_mapped;  
    /* ring mapped provided buffers, but mmap'ed by application */  
    __u8 is_mmap;  
};
```

```
union {  
    struct list_head buf_list;  
    struct {  
        struct page **buf_pages;  
        struct io_uring_buf_ring *buf_ring;  
    };  
};
```

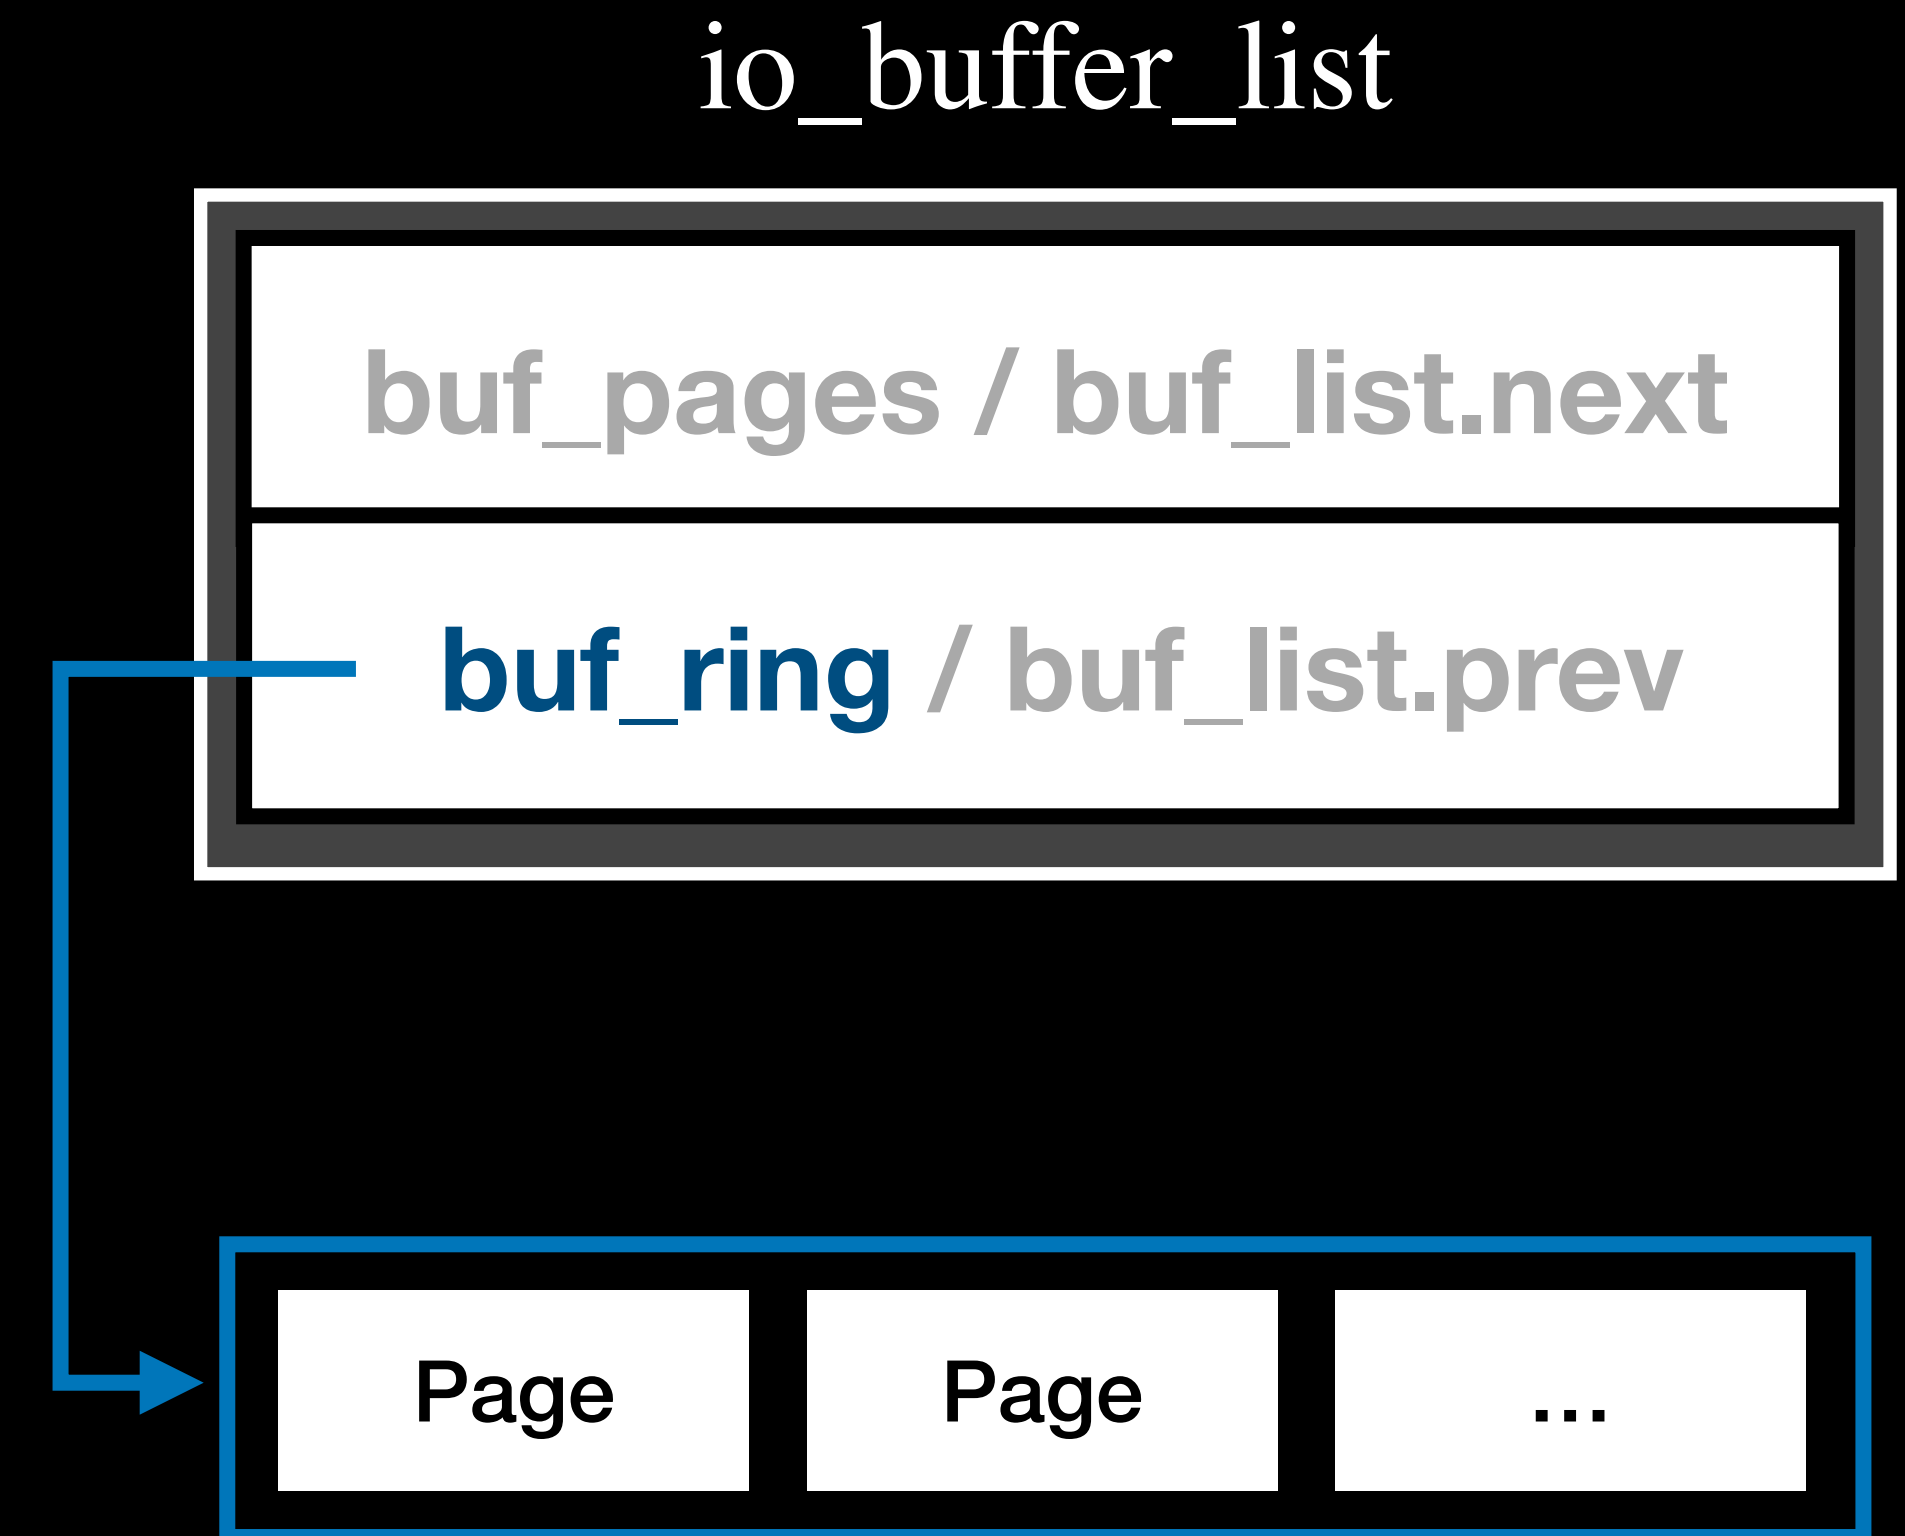
io_buffer_list

buf_pages / buf_list.next

buf_ring / buf_list.prev

Process 1

```
def unregister(uaddr):  
    mutex_lock(context_lock)  
  
    copy_from_user(uaddr, &req)  
    obj = io_buffer_obj_arr[req.idx]  
  
    obj.refcount -= 1  
    if obj.refcount == 0:  
        obj.is mmap = 0  
        → INIT_LIST_HEAD(&obj.buf_list)  
        kfree_rcu(obj)  
  
    mutex_unlock(context_lock)
```



Originally, buf_ring points to shared memory

Process 1

```
def unregister(uaddr):  
    mutex_lock(context_lock)  
  
    copy_from_user(uaddr, &req)  
    obj = io_buffer_obj_arr[req.idx]  
  
    obj.refcount -= 1  
    if obj.refcount == 0:  
        obj.is mmap = 0  
        → INIT_LIST_HEAD(&obj.buf_list)  
        kfree_rcu(obj)  
  
    mutex_unlock(context_lock)
```

io_buffer_list



Call **INIT_LIST_HEAD()** to update it to point to object itself

Process 1

```
def unregister(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    obj = io_buffer_obj_arr[req.idx]

    obj.refcount -= 1
    if obj.refcount == 0:
        obj.is_mmap = 0
        → INIT_LIST_HEAD(&obj.buf_list)
        kfree_rcu(obj)

    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v2(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]
    if obj.is_mmap:
        → addr = obj.buf_ring
        map_address_to_user_space(addr)

    rcu_read_unlock()
```

Get the address from **buf_ring**

Process 2

io_buffer_list

buf_pages / buf_list.next

buf_ring / buf_list.prev

addr

But the buf_ring now points to
the **starting address of object**

```
def mmap_v2(idx):  
    rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
    if obj.is_mmap:  
        addr = obj.buf_ring  
        map_address_to_user_space(addr)  
  
    rcu_read_unlock()
```

Process 2



```
def mmap_v2(idx):  
    rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
    if obj.is_mmap:  
        addr = obj.buf_ring  
        → map_address_to_user_space(addr)
```

Map I/O buffer object to user space

\$ Bug 3 - CVE-2024-35880

- RCU are introduced to prevent deadlock
 - Re-enable **concurrent access** to I/O buffer objects

\$ Bug 3 - CVE-2024-35880

- RCU are introduced to prevent deadlock
 - Re-enable concurrent access to I/O buffer objects
- Lead to a race condition in unregistration process
 - Concurrently retrieve the ring buffer from the reinitialized union field causes **incorrect memory mapping**

\$ Bug 3 - CVE-2024-35880

- Fixes:
 - **Update reference count** to prevent early unregistration
 - When the refcount reaches zero, the buffer can no longer be mapped

```
+     rcu_read_lock();  
+     bl = xa_load(&ctx->io_bl_xa, bgid);  
+     /* must be a mmap'able buffer ring and have pages */  
+     ret = false;  
+     if (bl && bl->is_mmap)  
+         ret = atomic_inc_not_zero(&bl->refs);  
+     rcu_read_unlock();
```

Process 1

```
def unregister(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    obj = io_buffer_obj_arr[req.idx]

    obj.refcount -= 1
    if obj.refcount == 0:
        obj.is_mmap = 0
        INIT_LIST_HEAD(&obj.buf_list)
        kfree_rcu(obj)

    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
+       if atomic_inc_not_zero(&obj.refcount) == 0:
+           early return

        addr = obj.buf_ring
        map_address_to_user_space(addr)

+       if atomic_dec_and_test(&obj.refcount) == 0:
+           kfree_rcu(obj)
```

Update **refcount** during memory mapping

Process 1

```
def unregister(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    obj = io_buffer_obj_arr[req.idx]

    obj.refcount -= 1
    if obj.refcount == 0:
        obj.is_mmap = 0
```

The ring_buf can no longer be destroyed **concurrently**

```
mutex_unlock(context_lock)
```

Process 2

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
        if atomic_inc_not_zero(&obj.refcount) == 0:
            early return

        addr = obj.buf_ring
        map_address_to_user_space(addr)

        if atomic_dec_and_test(&obj.refcount) == 0:
            kfree_rcu(obj)

    rcu_read_unlock()
```

\$ Mitigations

- By now:
 1. **Deferred page release** prevents the mapped memory from being freed
 2. **RCU protection** ensures the I/O buffer object is not freed too early
 3. **Correct refcount updates** prevent concurrent resets

\$ Mitigations

- By now:
 1. Deferred page release prevents the mapped memory from being freed
 2. RCU protection ensures the I/O buffer object is not freed too early
 3. Correct refcount updates prevent concurrent resets
- Unbreakable?

Breaking I/O Buffer Again:
CVE-2025-21836

\$ Insight

- According to the patch for CVE-2024-35880:
 - Concurrent access to the I/O buffer is still **allowed**

\$ Insight

- According to the patch for CVE-2024-35880:
 - Concurrent access to the I/O buffer is still allowed
- **Re-registering** an existing I/O buffer will modify object fields as well
 - Also called as “upgrading”
 - Reuse an empty provided buffer as ring buffer

\$ Insight

- According to the patch for CVE-2024-35880:
 - Concurrent access to the I/O buffer is still allowed
- Re-registering an existing I/O buffer will modify object fields as well
 - Also called as “upgrading”
 - Reuse an empty provided buffer as ring buffer
- Let’s revisit the **registration** process and see how it works!

```
def register(uaddr):
    mutex_lock(context_lock)

    copy_from_user(uaddr, &req)
    existing_obj = io_buffer_obj_arr[req.idx]

    if existing_obj and not is_empty_provided_buffer(obj):
        early return

    if existing_obj:
        obj = existing_obj
    else:
        obj = new_io_buffer_obj(&req)

    if reserved_memory_request(req):
        obj.is_mapped = 1
        obj.is_mmap = 1
        obj.buf_ring = alloc_pages()

    obj.refcount = 1
    io_buffer_obj_arr[req.idx] = obj

    mutex_unlock(context_lock)
```

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    copy_from_user(uaddr, &req)  
    existing_obj = io_buffer_obj_arr[req.idx]
```

```
if existing_obj and not is_empty_provided_buffer(obj):  
    early return
```

Only an **empty provided buffer** can be re-registered

```
if reserved_memory_request(req):  
    obj.is_mapped = 1  
    obj.is_mmap = 1  
    obj.buf_ring = alloc_pages()  
  
obj.refcount = 1  
io_buffer_obj_arr[req.idx] = obj  
  
mutex_unlock(context_lock)
```

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    copy_from_user(uaddr, &req)  
    existing_obj = io_buffer_obj_arr[req.idx]  
  
    if existing_obj and not is_empty_provided_buffer(obj):  
        early return
```

```
if reserved_memory_request(req):  
    obj.is_mapped = 1  
    obj.is_mmap = 1  
    obj.buf_ring = alloc_pages()  
  
obj.refcount = 1  
io_buffer_obj_arr[req.idx] = obj
```

Do the **same things** as in the first registration

\$ Root Cause

- At a glance, these seem safe operations:
 - Code reuse is common, so some dummy behavior is expected

\$ Root Cause

- At a glance, these seem safe operations:
 - Code reuse is common, so some dummy behavior is expected
 - A provided buffer is **not allowed to be mapped** (b1->is_mmap is false)
 - So memory mapping will early return :(

\$ Root Cause

- At a glance, these seem safe operations:
 - Code reuse is common, so some dummy behavior is expected
 - A provided buffer is **not allowed to be mapped** (b1->is_mmap is false)
 - So memory mapping will early return :(
 - ... won't it? 🤔

Process 1

Upgrade a provided buffer

```
def register(uaddr):
    mutex_lock(context_lock)

    # [...]

    obj = existing_obj

    if reserved_memory_request(req):
        obj.is_mapped = 1
        obj.is_mmap = 1
        obj.buf_ring = alloc_pages()

    obj.refcount = 1
    io_buffer_obj_arr[req.idx] = obj

    mutex_unlock(context_lock)
```

Process 2

Map an I/O buffer to user space

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
        if atomic_inc_not_zero(&obj.refcount) == 0:
            early return

        addr = obj.buf_ring
        map_address_to_user_space(addr)

        if atomic_dec_and_test(&obj.refcount) == 0:
            kfree_rcu(obj)

    rcu_read_unlock()
```


Process 1

```
def register(uaddr):
    mutex_lock(context_lock)

    # [...]

    obj = existing_obj

    if reserved_memory_request(req):
        obj.is_mapped = 1
        obj.is_mmap = 1
        obj.buf_ring = alloc_pages()

    obj.refcount = 1
    io_buffer_obj_arr[req.idx] = obj

    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
        if atomic_inc_not_zero(&obj.refcount) == 0:
            map_address_to_user_space(addr)

            if atomic_dec_and_test(&obj.refcount) == 0:
                kfree_rcu(obj)

    rcu_read_unlock()
```

A provided buffer **cannot be mapped since `is_mmap = 0`**

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    # [...]  
  
    obj = existing_obj  
  
    if reserved_memory_request(req):  
        obj.is_mapped = 1  
        obj.is_mmap = 1  
        obj.buf_ring = alloc_pages()  
  
    io_buffer_obj_arr[req.idx] = obj  
  
    mutex_unlock(context_lock)
```

But upgrading can make it **mappable**

Process 2

```
def mmap_v3(idx):  
    rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
  
    if obj.is_mmap: ✓  
        if atomic_inc_not_zero(&obj.refcount) == 0:  
            early return  
  
        addr = obj.buf_ring  
        map_address_to_user_space(addr)  
  
        if atomic_dec_and_test(&obj.refcount) == 0:  
            kfree_rcu(obj)  
  
    rcu_read_unlock()
```

Process 1

```
def register(uaddr):
    mutex_lock(context_lock)

    # [...]

    obj = existing_obj

    if reserved_memory_request(req):
        obj.is_mapped = 1
        → obj.is_mmap = 1
        obj.buf_ring = alloc_pages()

    obj.refcount = 1
    io_buffer_obj_arr[req.idx] = obj

    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
        → if atomic_inc_not_zero(&obj.refcount) == 0:
            early return

    if atomic_dec_and_test(&obj.refcount) == 0:
        kfree_rcu(obj)

    rcu_read_unlock()
```

Increase the refcount from 1 to 2

Process 1

```
def register(uaddr):
    mutex_lock(context_lock)

    # [...]

    obj = existing_obj

    if reserved_memory_request(req):
        obj.is_mapped = 1
        obj.is_mmap = 1
        obj.buf_ring = alloc_pages()
```


 `obj.refcount = 1`

The hardcoded assignment
sets refcount to **1**

Process 2

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
         if atomic_inc_not_zero(&obj.refcount) == 0:
            |     early return

    addr = obj.buf_ring
    map_address_to_user_space(addr)

    if atomic_dec_and_test(&obj.refcount) == 0:
        |     kfree_rcu(obj)

    rcu_read_unlock()
```

Process 1

```
def register(uaddr):
    mutex_lock(context_lock)

    # [...]

    obj = existing_obj

    if reserved_memory_request(req):
        obj.is_mapped = 1
        obj.is_mmap = 1
        obj.buf_ring = alloc_pages()

    obj.refcount = 1
    io_buffer_obj_arr[req.idx] = obj
    mutex_unlock(context_lock)
```

Process 2

```
def mmap_v3(idx):
    rcu_read_lock()

    obj = io_buffer_obj_arr[idx]

    if obj.is_mmap:
        if atomic_inc_not_zero(&obj.refcount) == 0:
            early return

        addr = obj.buf_ring
        map_address_to_user_space(addr)

        if atomic_dec_and_test(&obj.refcount) == 0:
            kfree_rcu(obj)
```

Decrease refcount from 1 to 0

Process 1

```
def register(uaddr):  
    mutex_lock(context_lock)  
  
    # [...]  
  
    obj = existing_obj  
  
    if reserved_memory_request(req):  
        obj.is_mapped = 1  
        obj.is_mmap = 1  
        obj.buf_ring = alloc_pages()  
  
    obj.refcount = 1  
    io_buffer_obj_arr[req.idx] = obj
```

Process 2

```
def mmap_v3(idx):  
    rcu_read_lock()  
  
    obj = io_buffer_obj_arr[idx]  
  
    if obj.is_mmap:  
        if atomic_inc_not_zero(&obj.refcount) == 0:  
            early return  
  
        addr = obj.buf_ring  
        map_address_to_user_space(addr)  
  
        if atomic_dec_and_test(&obj.refcount) == 0:  
            kfree_rcu(obj)
```

I/O buffer object UAF

\$ Root Cause

- This race is very hard to hit
 - Require memory mapping and upgrading in a specific order
- Once triggered, it allows access to the **freed I/O buffer object**
 - UAF on `io_buffer_list`
 - Full control over the **mapped address**

\$ Exploitation

1. Environment setup
2. Try to hit the race
3. Side channel via mmap return value to detect success
4. Wait 5 seconds for RCU drain
5. Reclaim the freed buffer object via spraying
6. Overwrite the kernel data `core_pattern[]`

\$ Exploitation

1. Environment setup
2. Try to hit the race
3. Side channel via mmap return value to detect success
4. Wait 5 seconds for RCU drain
5. Reclaim the freed buffer object via spraying
6. Overwrite the kernel data `core_pattern[]`

\$ Exploitation

- Target object: struct io_buffer_list
 - kmalloc-64-cg (provided buffer)
 - kmalloc-64 (ring buffer)
- Reclaim using **message queue** mechanism
 - Spray many struct msg_msgseg

\$ Exploitation

- `kfree_rcu()`
 - Tiny RCU: directly calls `call_rcu()`
 - Tree RCU:
 - Freed objects are batched
 - Batches are drained every 5 seconds (`KFREE_DRAIN_JIFFIES`)
- kernelCTF uses **Tree RCU**

\$ Exploitation

1. Environment setup
2. Try to hit the race
3. Side channel via mmap return value to detect success
4. Wait 5 seconds for RCU drain
5. Reclaim the freed buffer object via spraying
6. Overwrite the kernel data `core_pattern[]`

\$ Exploitation

- If race fails:
 - Does not lead to a kernel panic
 - The refcount of targeted `io_buffer_list` is **zero**

\$ Exploitation

- Side-channel race result by mmap
 - Race **fails**: the mapping handler behaves normally
 - Race succeeds: the mmap handler sees a zero refcount and returns an error

\$ Exploitation

- Side-channel race result by mmap
 - Race fails: the mapping handler behaves normally
 - Race **succeeds**: the mmap handler sees a zero refcount and **returns an error**
 - Thanks to `atomic_inc_not_zero()` 😊

\$ Exploitation

- KASLR bypass
 - Not provide an information leak primitive :(

\$ Exploitation

- KASLR bypass
 - Not provide an information leak primitive :(
 - **EntryBleed (CVE-2022-4543)**: time-based side-channel on x86_64
 - Measure prefetch timing of kernel addresses
 - Leak entry_SYSCALL_64 address
 - Technique by William

\$ Exploitation

- Read the flag
 1. Set corrupted bl->buf_ring to the kernel variable `core_pattern[]`
 2. Map the memory into user space
 3. Overwrite it with our executable path
 4. Trigger a SEGFAULT and get flag!

\$ Exploitation

1. Environment setup
2. Try to hit the race
3. Side channel via mmap return value to detect success
4. Wait 5 seconds for RCU drain
5. Reclaim the freed buffer object via spraying
6. Overwrite the kernel data `core_pattern[]`

\$ Exploitation

- Extending the race window with **timer interrupts**
 - Enqueue many timerfd waiters
 - The timer interrupt handler spends more time iterating the list
 - Proposed by Jann Horn
- In our case, we need **two timerfds** due to the narrow race window

Upgrading

Process 1

• Set bl->is_mmap to 1



mmap handling

Process 2



Upgrading

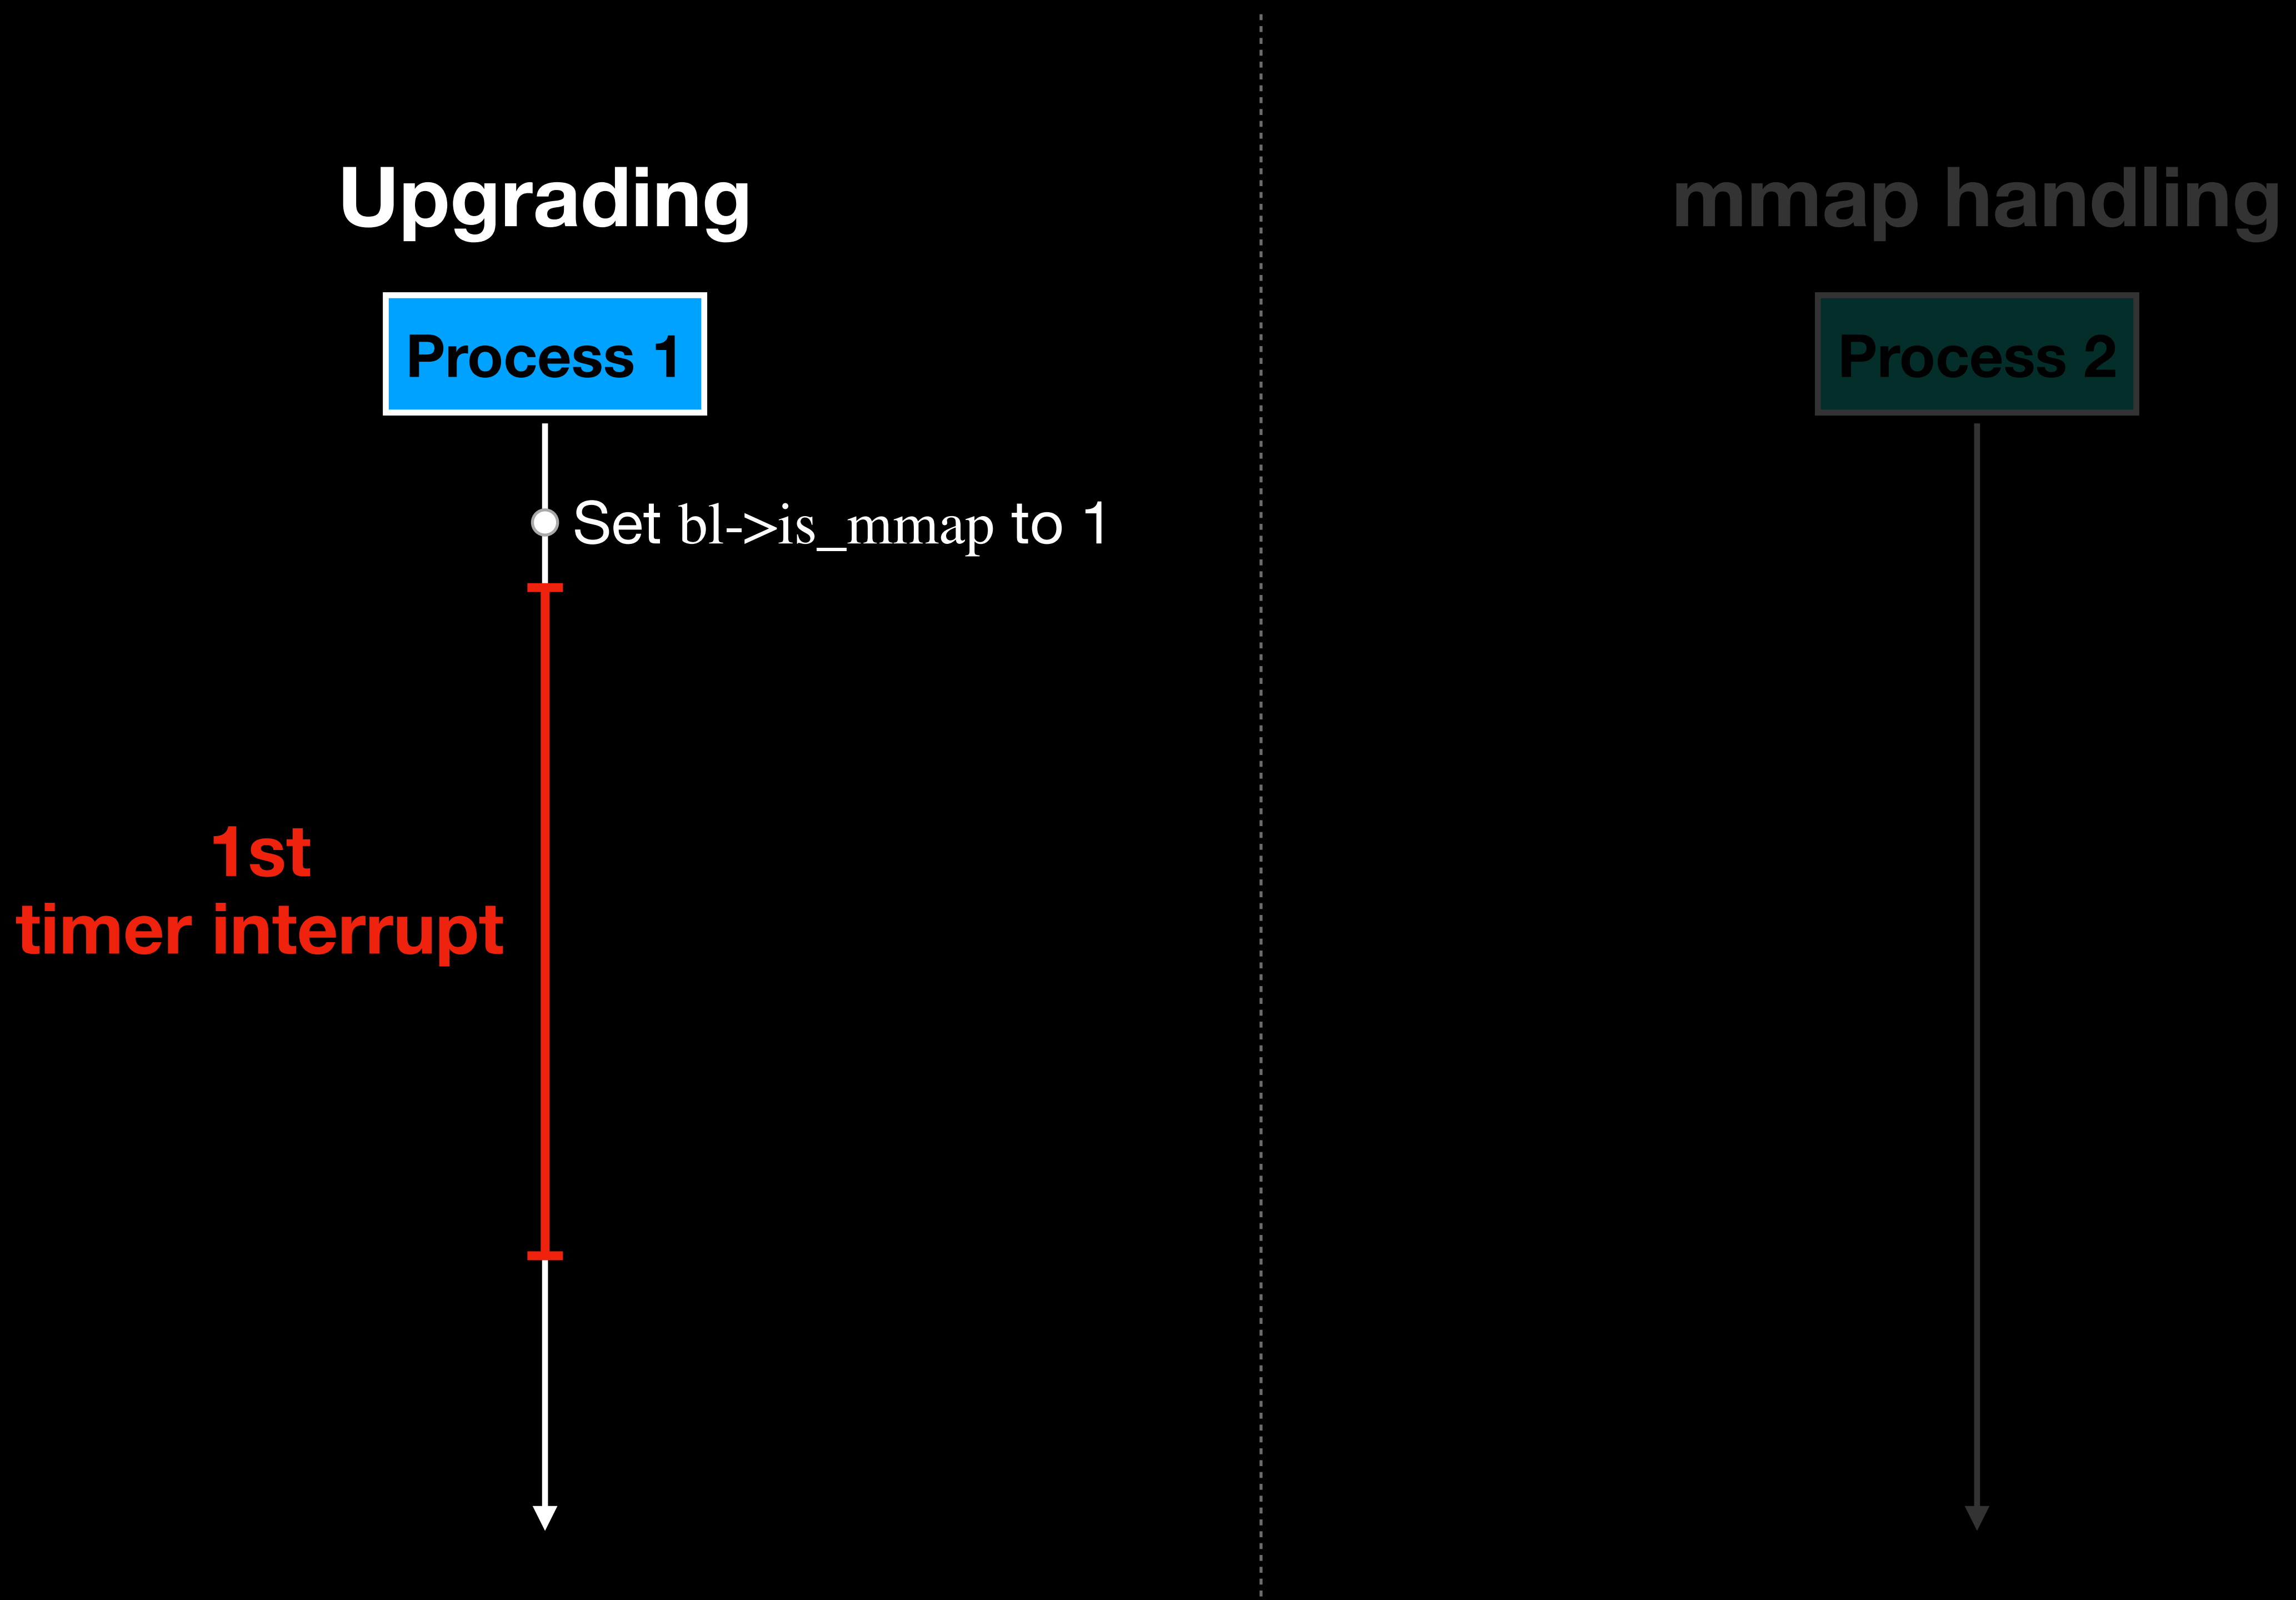
Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

mmap handling

Process 2



Upgrading

Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

mmap handling

Process 2

● Get I/O buffer

● Check bl->is_mmap

Upgrading

Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

mmap handling

Process 2

● Get I/O buffer

● Check bl->is_mmap ✓

Upgrading

Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

mmap handling

Process 2

- Get I/O buffer
- Check bl->is_mmap
- Inc bl->refs (1 -> 2)

Upgrading

Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

mmap handling

Process 2

● Get I/O buffer
● Check bl->is_mmap
● Inc bl->refs (1 -> 2)

**2nd
timer interrupt**

Upgrading

Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

● Set bl->refs to 1

mmap handling

Process 2

● Get I/O buffer

● Check bl->is_mmap

● Inc bl->refs (1 -> 2)

**2nd
timer interrupt**

Upgrading

Process 1

● Set bl->is_mmap to 1

**1st
timer interrupt**

● Set bl->refs to 1

mmap handling

Process 2

● Get I/O buffer

● Check bl->is_mmap

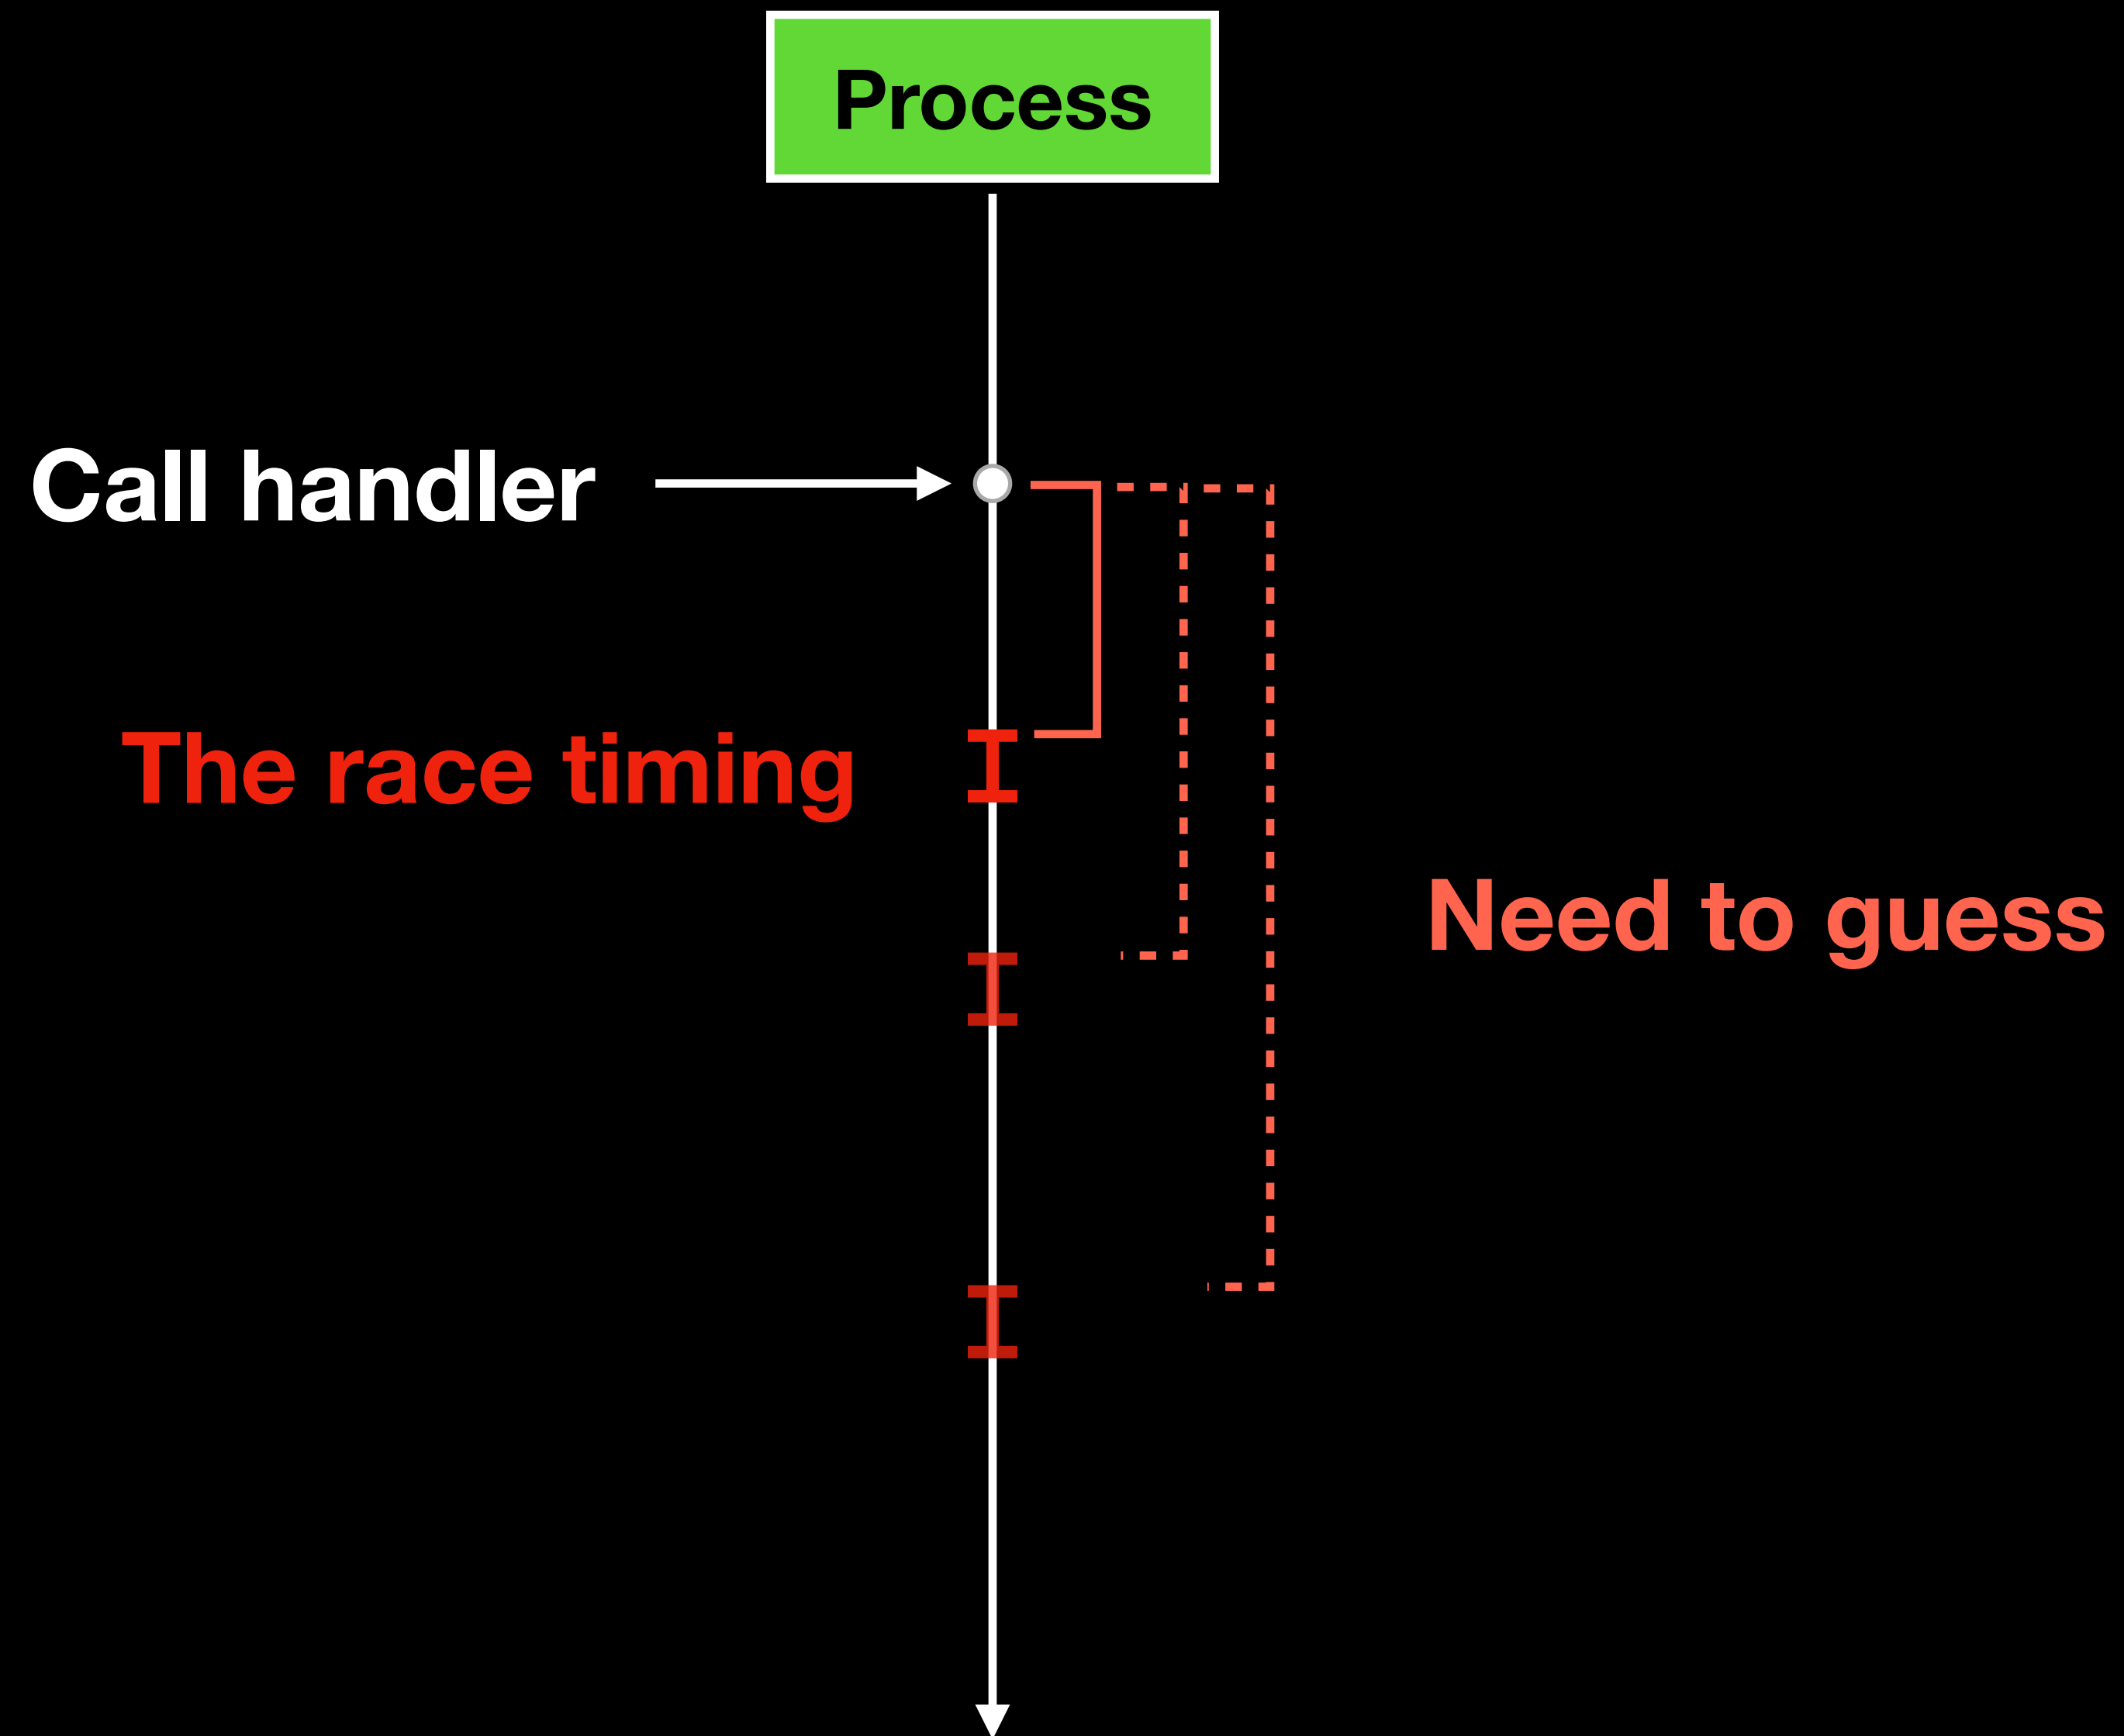
● Inc bl->refs (1 -> 2)

**2nd
timer interrupt**

● Dec bl->refs (1 -> 0)

\$ Exploitation

- **Unpredictable** timing gap between timer setup and target execution
 - Require guessing the correct race timing
 - Use a wide timeout range to increase chances
- Can we strategically bound the timeout range?



Process

The total time

I




Process



I

Syscall overhead




Demo time!

aaa@aaa:~/kernelctf/releases/lts-6.6.75/src\$


\$ Exploitation

- Success rate (per 30 min)
- kernelCTF: less than 1% 😓



\$ Exploitation

- Success rate (per 30 min)
 - kernelCTF: less than 1%
 - GitHub Actions: approximately 30% 🎉



\$ Patch

- io_uring/kbuf: reallocate buf lists on upgrade
- Upgrading now allocates a **new I/O buffer** instead of reusing the old one

```
@@ -642,12 +648,13 @@ int io_register_pbuf_ring(struct io_ring_ctx *ctx, void __user *arg)
    /* if mapped buffer ring OR classic exists, don't allow */
    if (bl->is_mapped || !list_empty(&bl->buf_list))
        return -EEXIST;
-   } else {
-       free_bl = bl = kzalloc(sizeof(*bl), GFP_KERNEL);
-       if (!bl)
-           return -ENOMEM;
+       io_destroy_bl(ctx, bl);
    }

+   free_bl = bl = kzalloc(sizeof(*bl), GFP_KERNEL);
+   if (!bl)
+       return -ENOMEM;
+
```

\$ Takeaways

- Memory sharing is common and requires careful handling
 - Complex ownership and lifetime management
 - Concurrency needed for performance increases risk
- Keep reference counts accurate to avoid UAF
- RCU prevents UAF, but not concurrent modifications

DEV✓CORE

Thanks!

Pumpkin 🎃 (@u1f383)
<https://u1f383.github.io/>